

Importing Libraries

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier

%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

Read csv file

In [2]:

```
bnk = pd.read_csv(r"C:\Users\SST190000\Downloads\Applied ML\archive\bank_marketing.csv",
sep = ';')
```

In [3]:

```
bnk
```

Out[3]:

	age	job	marital	education	default	housing	loan	contact	month	day_of_week	...	campaign
0	56	housemaid	married	basic.4y	no	no	no	telephone	may	mon	...	1
1	57	services	married	high.school	unknown		no	telephone	may	mon	...	1
2	37	services	married	high.school	no	yes	no	telephone	may	mon	...	1
3	40	admin.	married	basic.6y	no	no	no	telephone	may	mon	...	1
4	56	services	married	high.school	no	no	yes	telephone	may	mon	...	1
...
41183	73	retired	married	professional.course	no	yes	no	cellular	nov	fri	...	1
41184	46	blue-collar	married	professional.course	no	no	no	cellular	nov	fri	...	1
41185	56	retired	married	university.degree	no	yes	no	cellular	nov	fri	...	2
41186	44	technician	married	professional.course	no	no	no	cellular	nov	fri	...	1
41187	74	retired	married	professional.course	no	yes	no	cellular	nov	fri	...	3

41188 rows x 21 columns



Getting information about the dataset

In [4]:

```
bnk.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41188 entries, 0 to 41187
Data columns (total 21 columns):
#      Column              Non-Null Count  Dtype
#  ...  ...
```

```

0   age                41188 non-null int64
1   job                41188 non-null object
2   marital            41188 non-null object
3   education          41188 non-null object
4   default            41188 non-null object
5   housing            41188 non-null object
6   loan              41188 non-null object
7   contact            41188 non-null object
8   month              41188 non-null object
9   day_of_week        41188 non-null object
10  duration           41188 non-null int64
11  campaign           41188 non-null int64
12  pdays              41188 non-null int64
13  previous           41188 non-null int64
14  poutcome           41188 non-null object
15  emp.var.rate       41188 non-null float64
16  cons.price.idx     41188 non-null float64
17  cons.conf.idx      41188 non-null float64
18  euribor3m          41188 non-null float64
19  nr.employed        41188 non-null float64
20  y                  41188 non-null object
dtypes: float64(5), int64(5), object(11)
memory usage: 6.6+ MB

```

In [5]:

```
bnk.describe()
```

Out[5]:

	age	duration	campaign	pdays	previous	emp.var.rate	cons.price.idx	cons.conf.idx	e
count	41188.00000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000
mean	40.02406	258.285010	2.567593	962.475454	0.172963	0.081886	93.575664	-40.502600	-1.919640
std	10.42125	259.279249	2.770014	186.910907	0.494901	1.570960	0.578840	4.628198	1.454161
min	17.00000	0.000000	1.000000	0.000000	0.000000	-3.400000	92.201000	-50.800000	-1.919640
25%	32.00000	102.000000	1.000000	999.000000	0.000000	-1.800000	93.075000	-42.700000	-1.919640
50%	38.00000	180.000000	2.000000	999.000000	0.000000	1.100000	93.749000	-41.800000	-1.919640
75%	47.00000	319.000000	3.000000	999.000000	0.000000	1.400000	93.994000	-36.400000	-1.919640
max	98.00000	4918.000000	56.000000	999.000000	7.000000	1.400000	94.767000	-26.900000	-1.919640

In [6]:

```
bnk.columns
```

Out[6]:

```

Index(['age', 'job', 'marital', 'education', 'default', 'housing', 'loan',
       'contact', 'month', 'day_of_week', 'duration', 'campaign', 'pdays',
       'previous', 'poutcome', 'emp.var.rate', 'cons.price.idx',
       'cons.conf.idx', 'euribor3m', 'nr.employed', 'y'],
      dtype='object')

```

In [7]:

```
bnk.shape
```

Out[7]:

```
(41188, 21)
```

Getting information about the dataset

In [8]:

```
bnk.isnull().sum()
```

Out[8]:

```
age          0
job          0
marital      0
education    0
default      0
housing      0
loan         0
contact      0
month        0
day_of_week  0
duration     0
campaign     0
pdays       0
previous     0
poutcome     0
emp.var.rate 0
cons.price.idx 0
cons.conf.idx 0
euribor3m    0
nr.employed  0
Y            0
dtype: int64
```

Check for Duplicates and Null values

In [9]:

```
bnk.duplicated().sum()
```

Out[9]:

```
12
```

In [10]:

```
duplicates = bnk[bnk.duplicated()]
print("Duplicate rows : ")
duplicates
```

Duplicate rows :

Out[10]:

	age	job	marital	education	default	housing	loan	contact	month	day_of_week	...	campaign
1266	39	blue-collar	married	basic.6y	no	no	no	telephone	may	thu	...	1
12261	36	retired	married	unknown	no	no	no	telephone	jul	thu	...	1
14234	27	technician	single	professional.course	no	no	no	cellular	jul	mon	...	2
16956	47	technician	divorced	high.school	no	yes	no	cellular	jul	thu	...	3
18465	32	technician	single	professional.course	no	yes	no	cellular	jul	thu	...	1
20216	55	services	married	high.school	unknown	no	no	cellular	aug	mon	...	1
20534	41	technician	married	professional.course	no	yes	no	cellular	aug	tue	...	1
25217	39	admin.	married	university.degree	no	no	no	cellular	nov	tue	...	2
28477	24	services	single	high.school	no	yes	no	cellular	apr	tue	...	1
32516	35	admin.	married	university.degree	no	yes	no	cellular	may	fri	...	4
36951	45	admin.	married	university.degree	no	no	no	cellular	jul	thu	...	1
38281	71	retired	single	university.degree	no	no	no	telephone	oct	tue	...	1

Drop duplicates

In [11]:

```
bnk.drop_duplicates(keep=False, inplace=True)
bnk.duplicated().sum()
```

Out[11]:

0

In [12]:

```
bnk.isna().sum()
```

Out[12]:

```
age                0
job                0
marital            0
education          0
default            0
housing            0
loan              0
contact            0
month              0
day_of_week        0
duration           0
campaign           0
pdays            0
previous           0
poutcome           0
emp.var.rate       0
cons.price.idx     0
cons.conf.idx      0
euribor3m          0
nr.employed        0
y                  0
dtype: int64
```

Imputing 5%-10% null values

In [13]:

```
bnk['emp.var.rate'] = bnk['emp.var.rate'].mask(np.random.random(bnk['emp.var.rate'].shape) < .1)
bnk['duration'] = bnk['duration'].mask(np.random.random(bnk['duration'].shape) < .05)
bnk['campaign'] = bnk['campaign'].mask(np.random.random(bnk['campaign'].shape) < .1)
bnk['pdays'] = bnk['pdays'].mask(np.random.random(bnk['pdays'].shape) < .1)
bnk['previous'] = bnk['previous'].mask(np.random.random(bnk['previous'].shape) < .05)
```

In [14]:

```
bnk.isna().sum()
```

Out[14]:

```
age                0
job                0
marital            0
education          0
default            0
housing            0
loan              0
contact            0
month              0
day_of_week        0
duration           0
campaign           0
pdays            0
previous           0
poutcome           0
emp.var.rate       0
cons.price.idx     0
cons.conf.idx      0
euribor3m          0
nr.employed        0
y                  0
dtype: int64
```

```

day_of_week      0
duration         2138
campaign         4085
pdays          4088
previous        1980
poutcome         0
emp.var.rate    4192
cons.price.idx   0
cons.conf.idx    0
euribor3m        0
nr.employed      0
y                0
dtype: int64

```

In [15]:

```
bnk.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 41164 entries, 0 to 41187
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   age                   41164 non-null  int64
 1   job                   41164 non-null  object
 2   marital               41164 non-null  object
 3   education             41164 non-null  object
 4   default               41164 non-null  object
 5   housing               41164 non-null  object
 6   loan                  41164 non-null  object
 7   contact               41164 non-null  object
 8   month                 41164 non-null  object
 9   day_of_week           41164 non-null  object
10   duration              39026 non-null  float64
11   campaign              37079 non-null  float64
12   pdays                 37076 non-null  float64
13   previous              39184 non-null  float64
14   poutcome              41164 non-null  object
15   emp.var.rate          36972 non-null  float64
16   cons.price.idx         41164 non-null  float64
17   cons.conf.idx          41164 non-null  float64
18   euribor3m              41164 non-null  float64
19   nr.employed            41164 non-null  float64
20   y                     41164 non-null  object
dtypes: float64(9), int64(1), object(11)
memory usage: 6.9+ MB

```

Replace null values with mean

In [16]:

```

bnk['emp.var.rate'].fillna(bnk['emp.var.rate'].mean(),inplace=True)
bnk['duration'].fillna(bnk['duration'].mean(),inplace=True)
bnk['campaign'].fillna(bnk['campaign'].mean(),inplace=True)
bnk['pdays'].fillna(bnk['pdays'].mean(),inplace=True)
bnk['previous'].fillna(bnk['previous'].mean(),inplace=True)

```

In [17]:

```
bnk.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 41164 entries, 0 to 41187
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   age                   41164 non-null  int64
 1   job                   41164 non-null  object
 2   marital               41164 non-null  object
 3   education             41164 non-null  object
 4   default               41164 non-null  object

```

```

4 default 41164 non-null object
5 housing 41164 non-null object
6 loan 41164 non-null object
7 contact 41164 non-null object
8 month 41164 non-null object
9 day_of_week 41164 non-null object
10 duration 41164 non-null float64
11 campaign 41164 non-null float64
12 pdays 41164 non-null float64
13 previous 41164 non-null float64
14 poutcome 41164 non-null object
15 emp.var.rate 41164 non-null float64
16 cons.price.idx 41164 non-null float64
17 cons.conf.idx 41164 non-null float64
18 euribor3m 41164 non-null float64
19 nr.employed 41164 non-null float64
20 y 41164 non-null object
dtypes: float64(9), int64(1), object(11)
memory usage: 6.9+ MB

```

Exploratory Data Analysis

Exploring numerical variables in 'bnk'

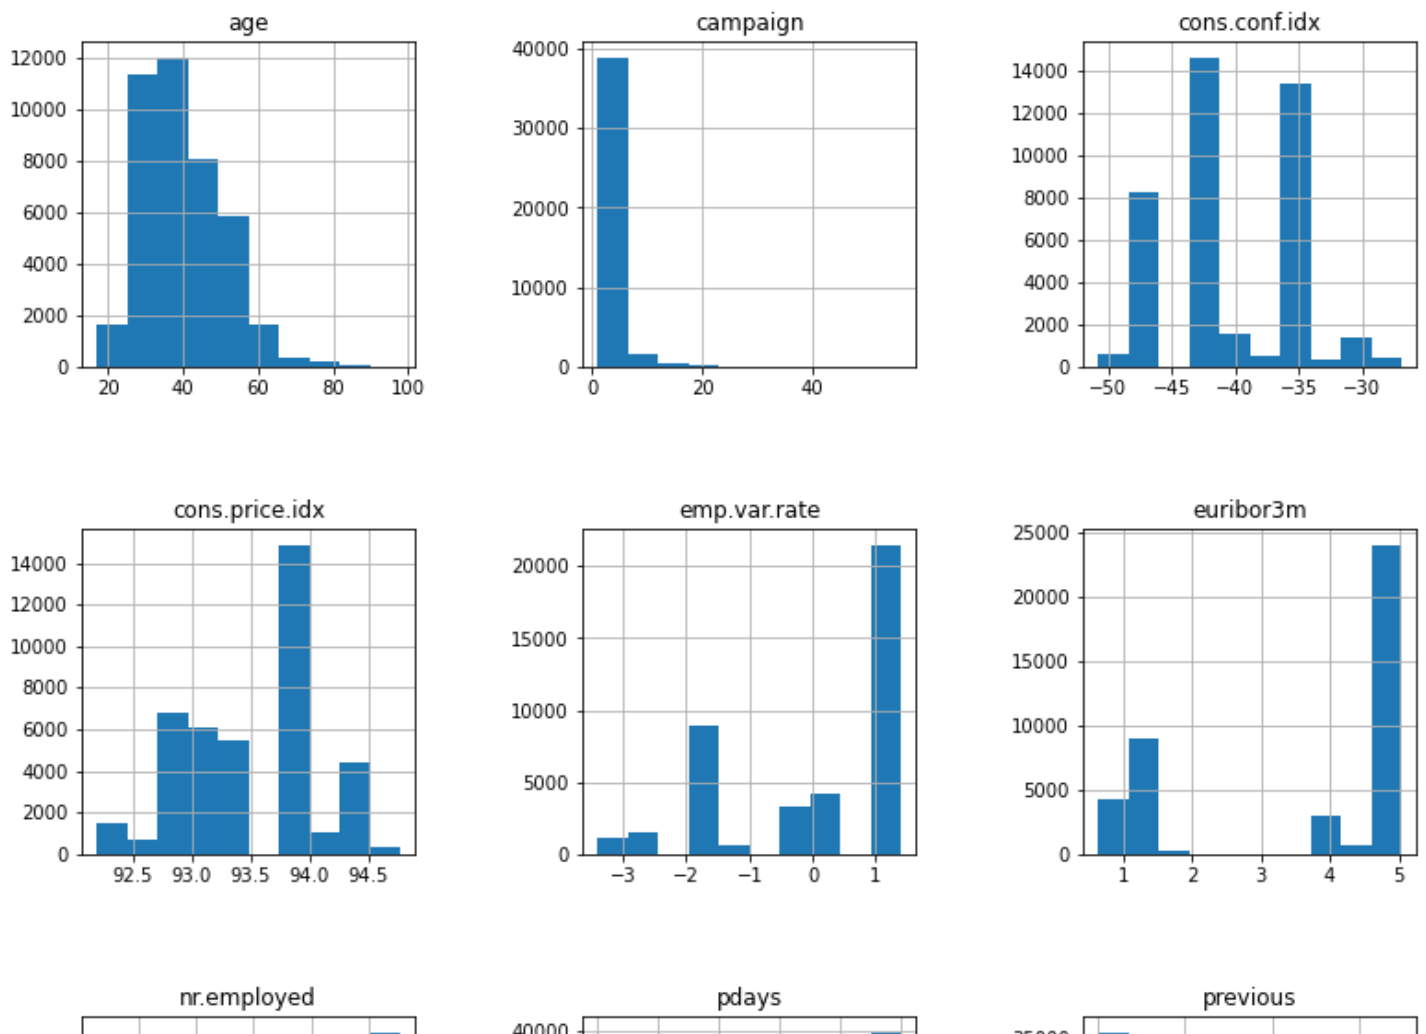
Histogram Subplots


In [18]:

```

col = ['age', 'campaign', 'pdays', 'previous', 'emp.var.rate', 'cons.price.idx', 'cons.conf.idx',
       'euribor3m', 'nr.employed']
bnk.hist(column=col, figsize=(13, 13))
plt.subplots_adjust(wspace = 0.5, hspace = 0.5)
plt.show()

```






Distribution of age variable

In [19]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
sns.distplot(bnk['age'], hist=True, kde=True,
              bins=int(150/10), color = 'purple',
              hist_kws={'edgecolor':'black'})
bca.set_title('Age Distribution', fontsize=15)
```

Out[19]:

Text(0.5, 1.0, 'Age Distribution')



Count of Duration

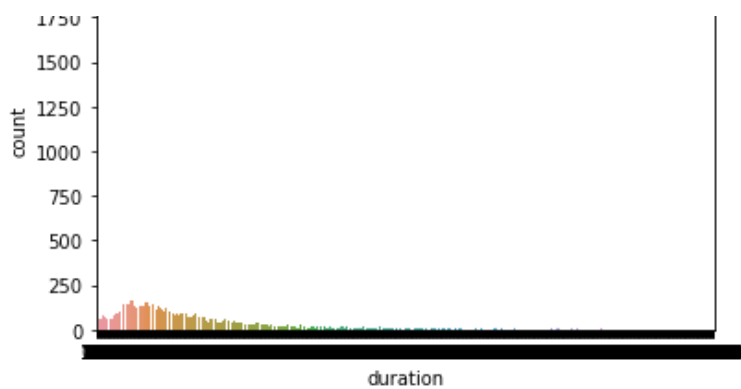
In [20]:

```
sns.countplot(x='duration', data=bnk)
bca.set_title('Count of Duration', fontsize=15)
```

Out[20]:

Text(0.5, 1.0, 'Count of Duration')





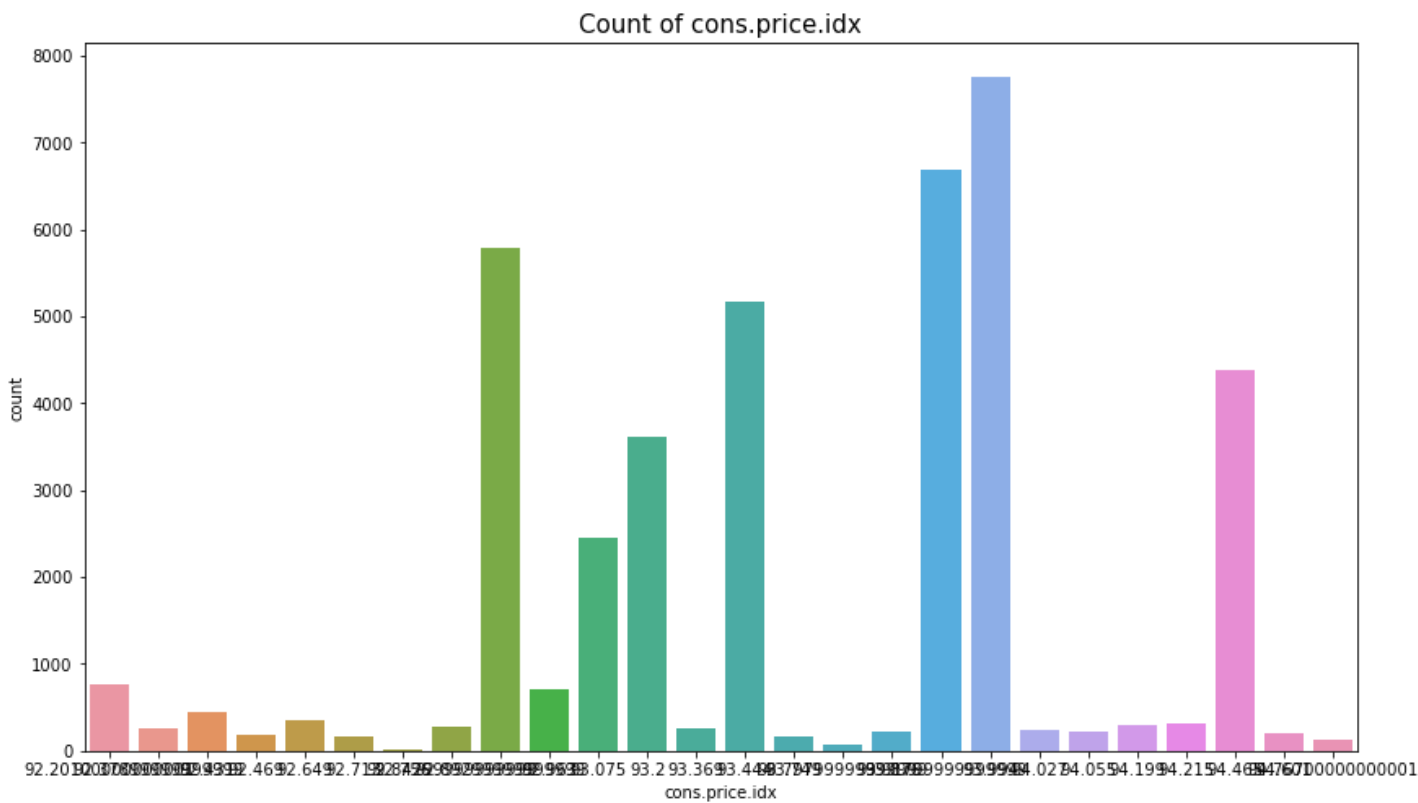
Count of cons.price.idx

In [21]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
sns.countplot(x='cons.price.idx', data=bnk)
bca.set_title('Count of cons.price.idx', fontsize=15)
```

Out[21]:

Text(0.5, 1.0, 'Count of cons.price.idx')



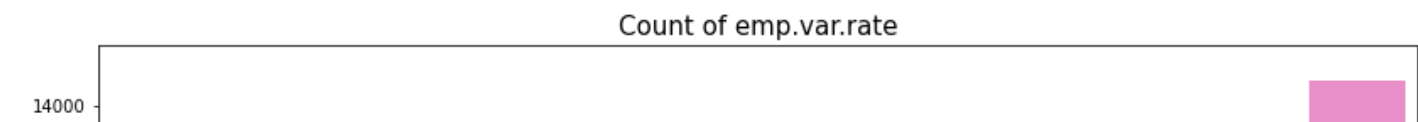
Count of emp.var.rate


In [22]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
sns.countplot(x='emp.var.rate', data=bnk)
bca.set_title('Count of emp.var.rate', fontsize=15)
```

Out[22]:

Text(0.5, 1.0, 'Count of emp.var.rate')





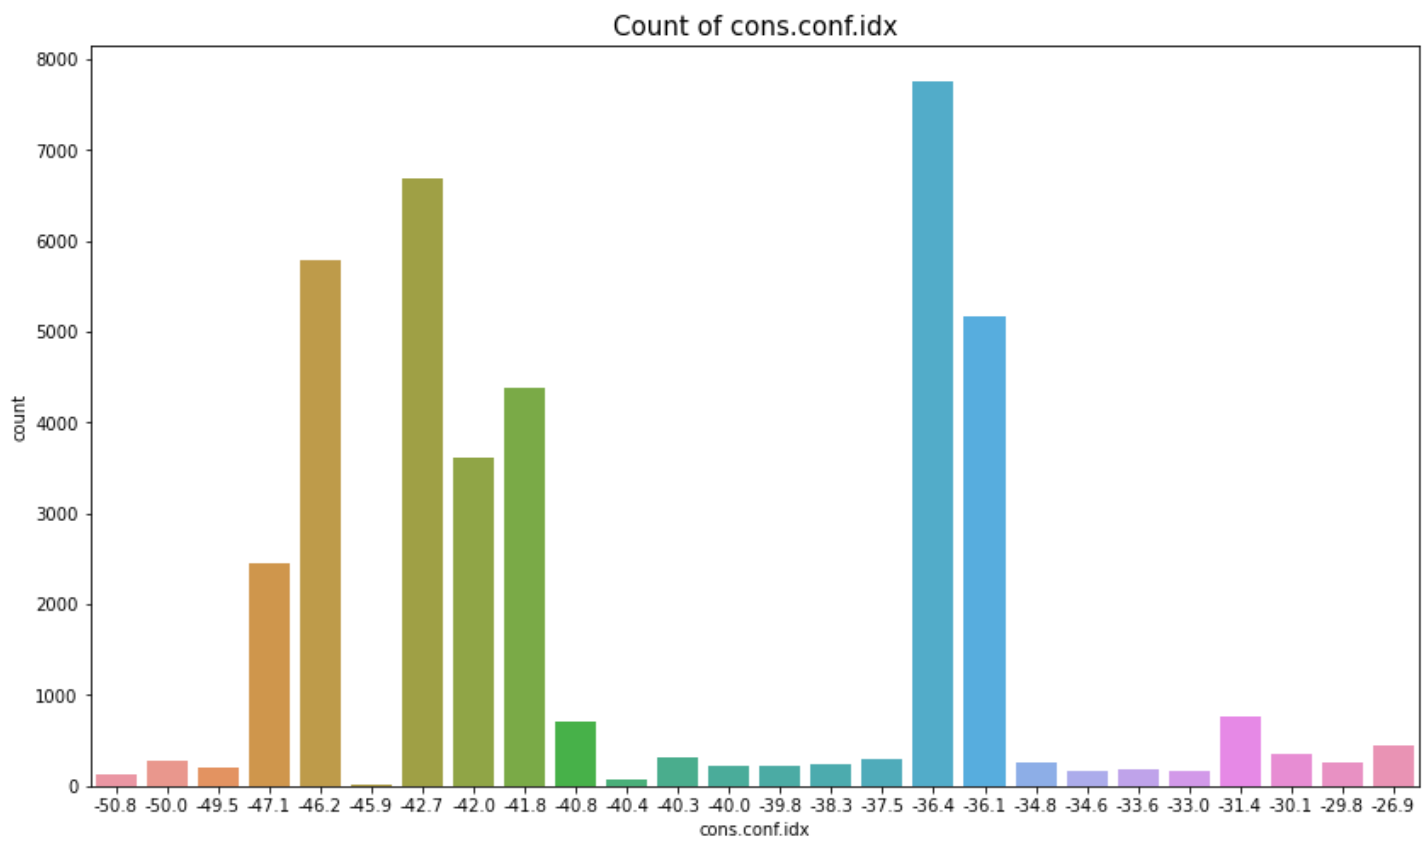
Count of cons.conf.idx

In [23]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
sns.countplot(x='cons.conf.idx', data=bnk)
bca.set_title('Count of cons.conf.idx', fontsize=15)
```

Out[23]:

Text(0.5, 1.0, 'Count of cons.conf.idx')




Count of euribor3m

In [24]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
sns.countplot(x='euribor3m', data=bnk)
bca.set_title('Count of euribor3m', fontsize=15)
```

Out[24]:

Text(0.5, 1.0, 'Count of euribor3m')

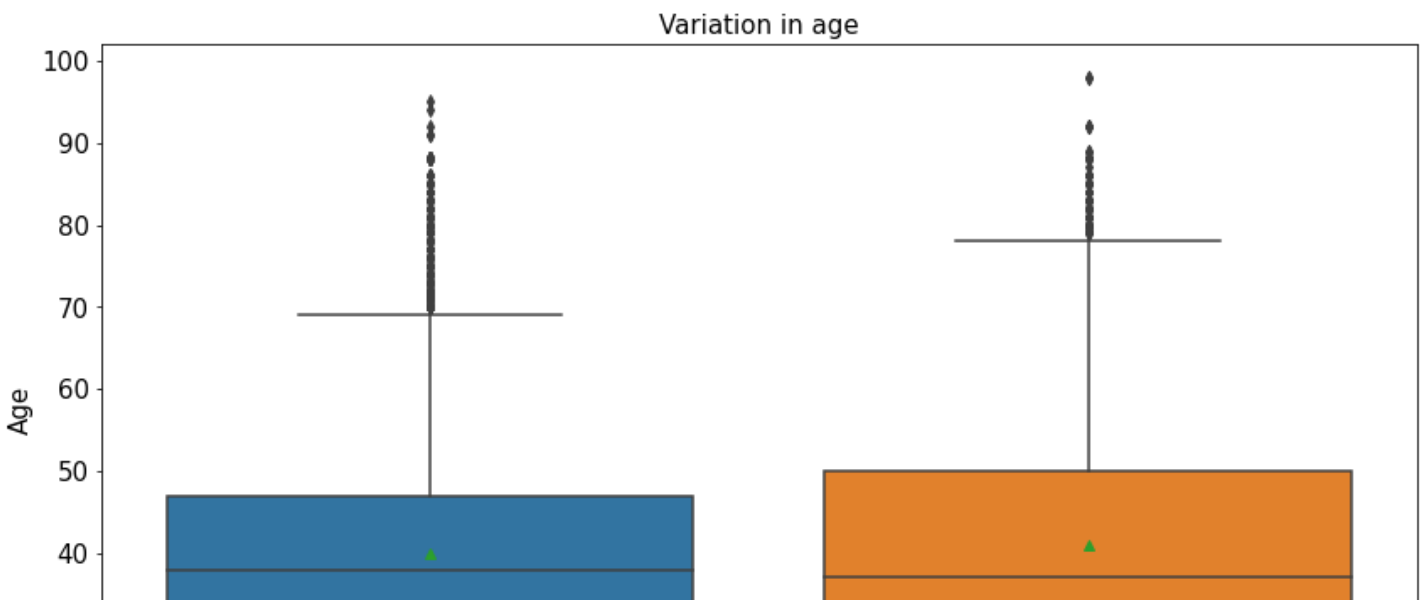



Exploring variation of numerical variables w.r.t target variable y

Variation in age

In [25]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot(x='y', y='age', data=bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Age', fontsize=15)
bca1.set_title('Variation in age', fontsize=15)
bca1.tick_params(labelsize=15)
```






Variation in duration

In [26]:


```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot( x='y', y= 'duration', data =bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Duration', fontsize=15)
bca1.set_title('Variation in duration', fontsize=15)
bca1.tick_params(labelsize=15)
```




Variation in Campaign

In [27]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot( x='y', y= 'campaign', data =bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Campaign', fontsize=15)
bca1.set_title('Variation in Campaign', fontsize=15)
bca1.tick_params(labelsize=15)
```

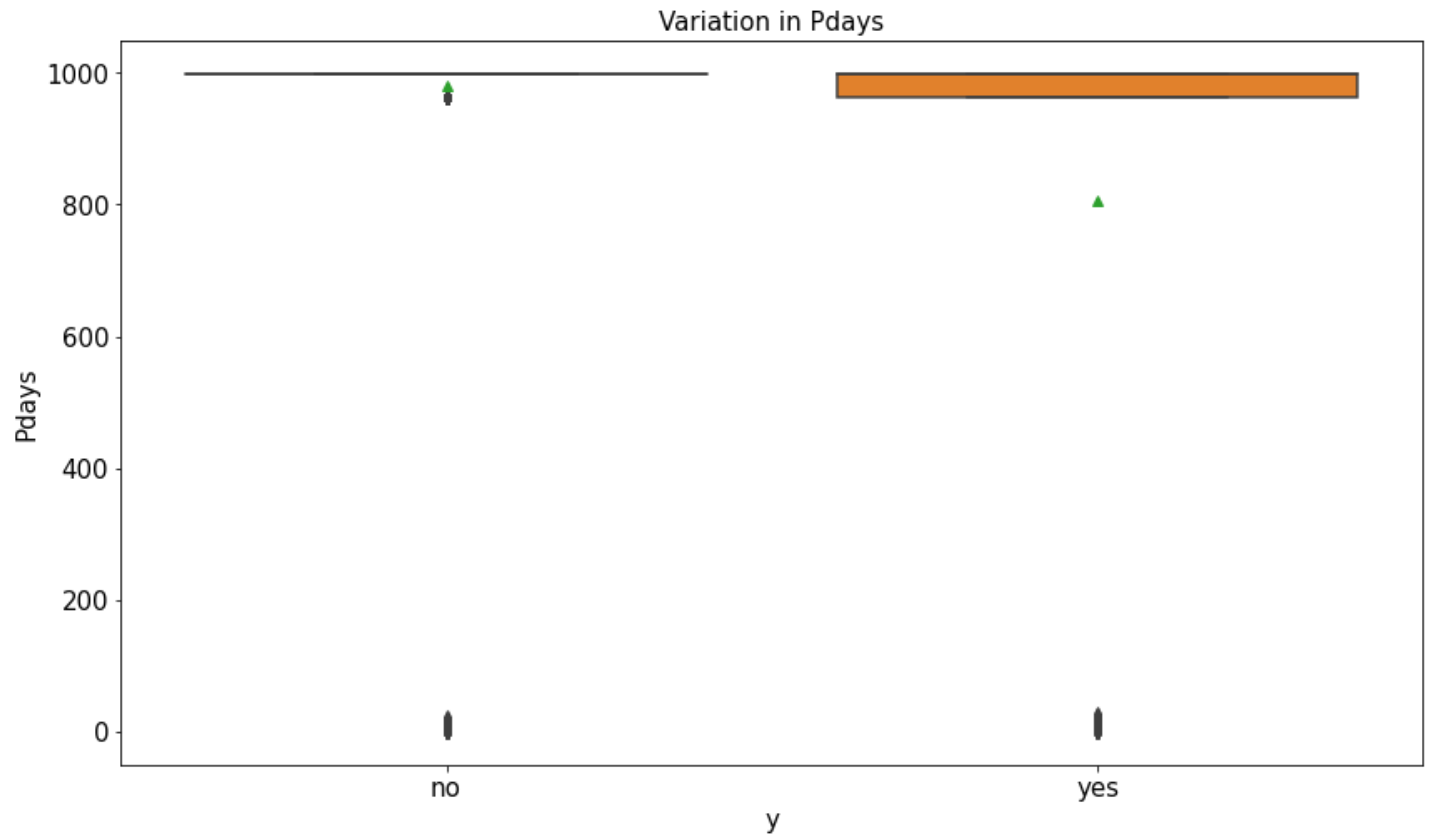




Variation in Pdays

In [28]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot( x='y', y= 'pdays', data =bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Pdays', fontsize=15)
bca1.set_title('Variation in Pdays', fontsize=15)
bca1.tick_params(labels=15)
```

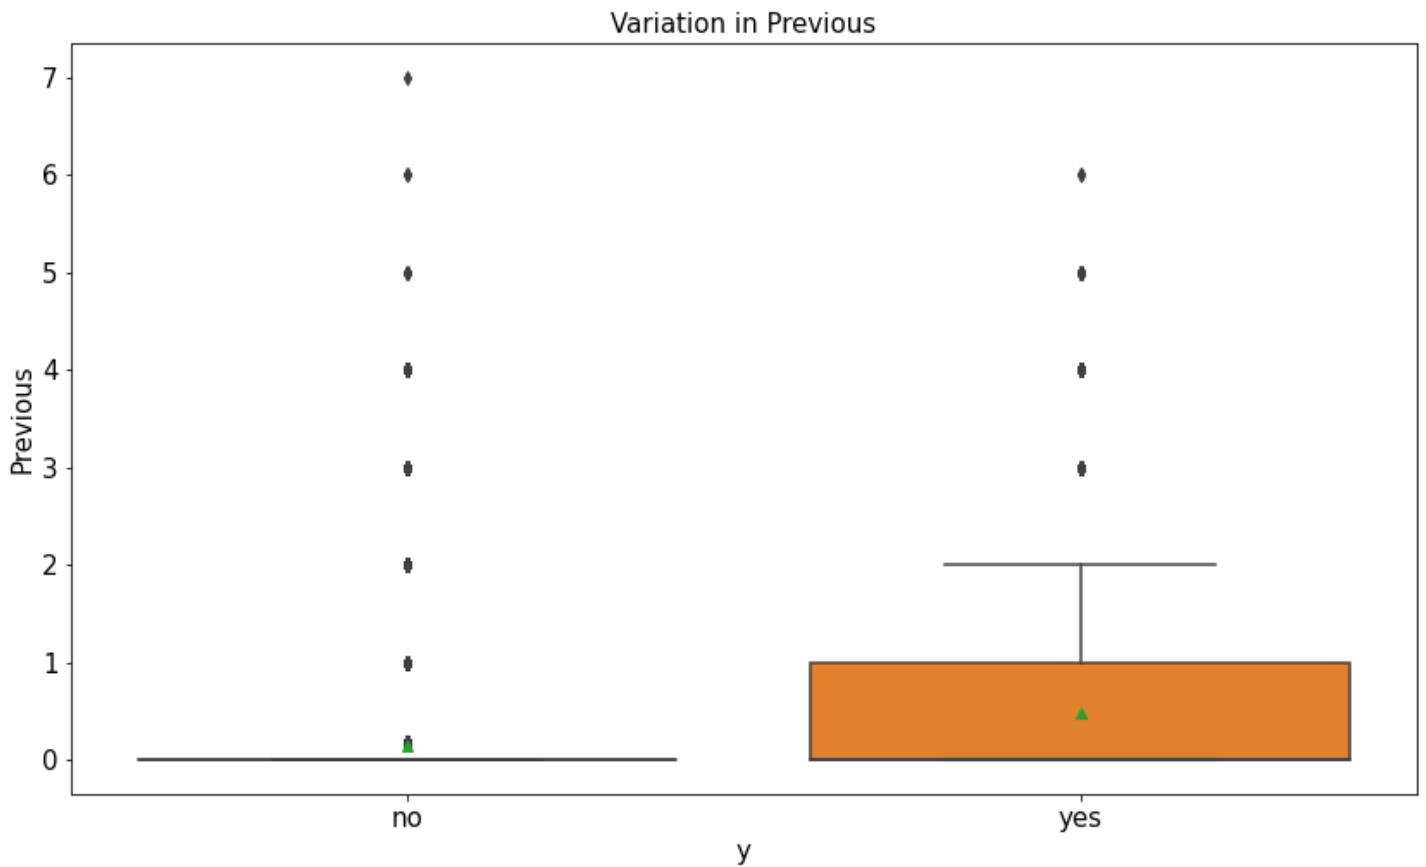


Variation in Previous

In [29]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot( x='y', y= 'previous', data =bnk, showmeans=True)
```

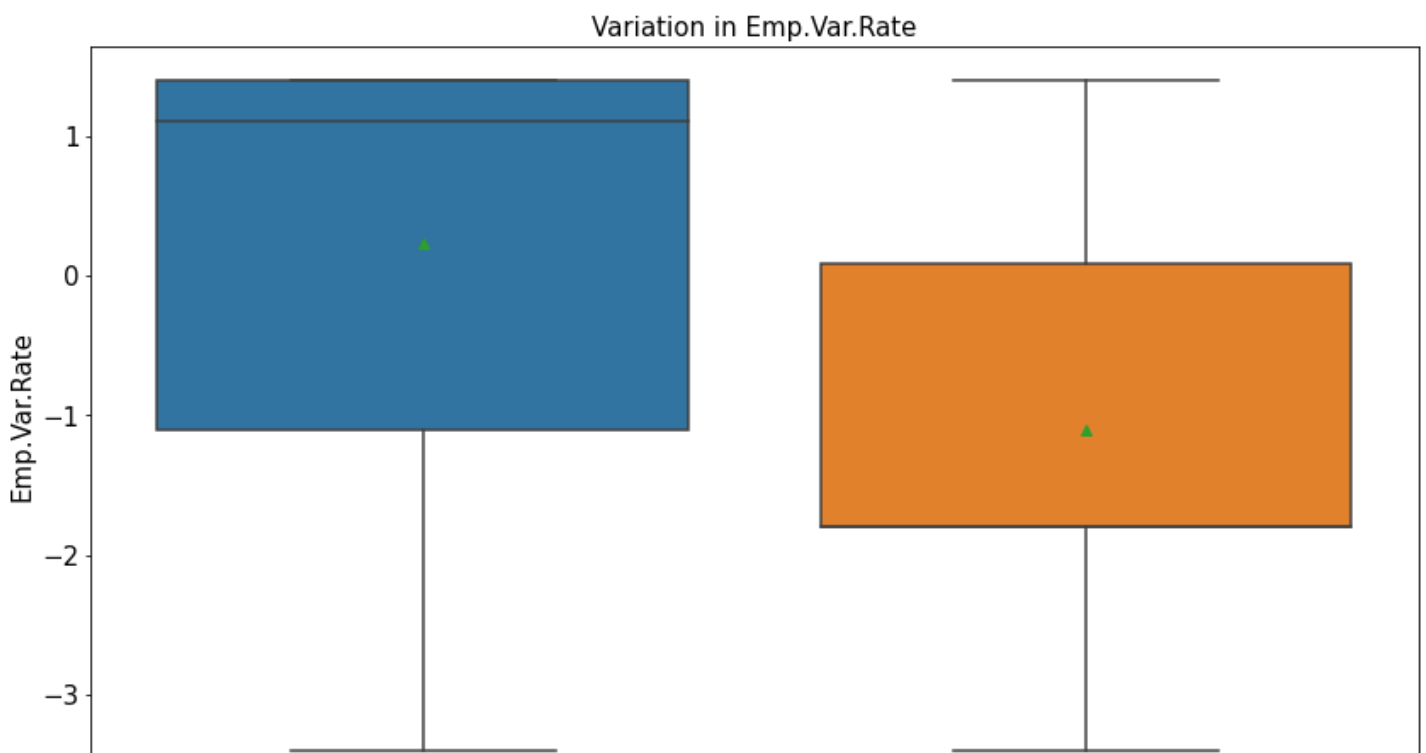
```
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Previous', fontsize=15)
bca1.set_title('Variation in Previous', fontsize=15)
bca1.tick_params(labelsize=15)
```



Variation in Emp.Var.Rate

In [30]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1=sns.boxplot( x='y', y= 'emp.var.rate', data =bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Emp.Var.Rate', fontsize=15)
bca1.set_title('Variation in Emp.Var.Rate', fontsize=15)
bca1.tick_params(labelsize=15)
```



no

yes


y

Variation in Cons.Price.Idx

In [31]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot( x='y', y= 'cons.price.idx', data =bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Cons.Price.Idx', fontsize=15)
bca1.set_title('Variation in Cons.Price.Idx', fontsize=15)
bca1.tick_params(labelsize=15)
```

Variation in Cons.Price.Idx





Variation in Cons.Conf.Idx

In [32]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot( x='y', y= 'cons.conf.idx', data =bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Cons.Conf.Idx', fontsize=15)
bca1.set_title('Variation in Cons.Conf.Idx', fontsize=15)
bca1.tick_params(labelsize=15)
```

Variation in Cons.Conf.Idx

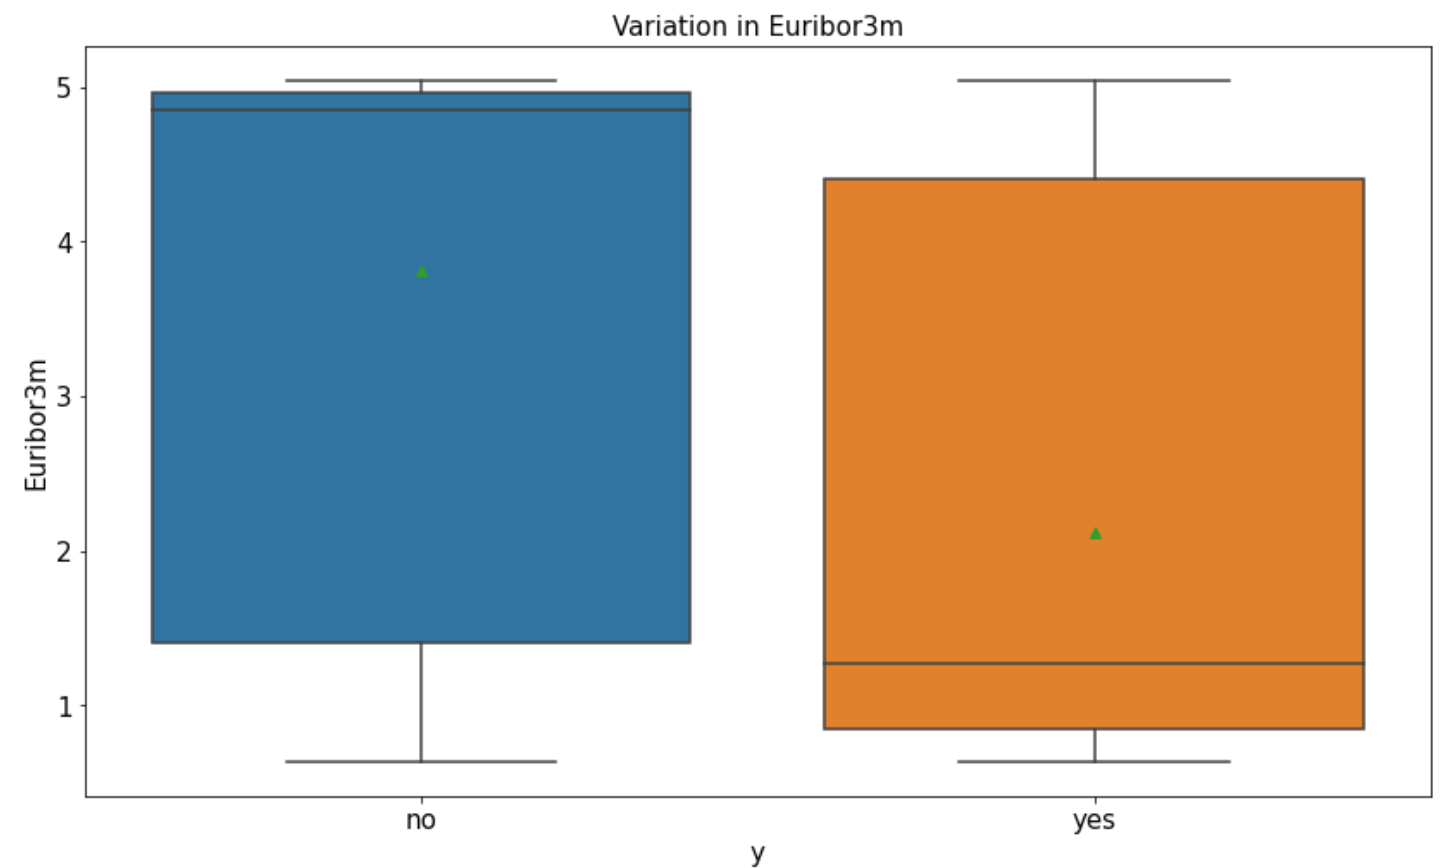




Variation in Euribor3m

In [33]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot( x='y', y= 'euribor3m', data =bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Euribor3m', fontsize=15)
bca1.set_title('Variation in Euribor3m', fontsize=15)
bca1.tick_params(labelsize=15)
```




Variation in Nr.Employed

In [34]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot( x='y', y= 'nr.employed', data =bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Nr.Employed', fontsize=15)
bca1.set_title('Variation in Nr.Employed', fontsize=15)
bca1.tick_params(labelsize=15)
```


Variation in Nr.Employed



Heatmap depicting correlation between all numerical variables

In [35]:

```
plt.subplots(figsize=(14,8))
sns.heatmap(bnk.corr(), annot=True)
plt.show()
```



Encoding and storing target variable 'y'

We perform one-hot-encoding on target variable 'y' in bnk dataframe as it is categorical data. We store the result in a new variable 'y'.

In [36]:

```
y = pd.get_dummies(bnk['y'], columns = ['y'], prefix = ['y'], drop_first = True)
bnk.head()
```

Out[36]:

	age	job	marital	education	default	housing	loan	contact	month	day_of_week	...	campaign	pdays	pre
0	56	housemaid	married	basic.4y	no	no	no	telephone	may	mon	...	2.570404	999.0	0.1
1	57	services	married	high.school	unknown	no	no	telephone	may	mon	...	1.000000	999.0	0.0
2	37	services	married	high.school	no	yes	no	telephone	may	mon	...	1.000000	999.0	0.0
3	40	admin.	married	basic.6y	no	no	no	telephone	may	mon	...	1.000000	999.0	0.0
4	56	services	married	high.school	no	no	yes	telephone	may	mon	...	1.000000	999.0	0.0

5 rows x 21 columns



Creating a new dataframe 'bank_client'

We are creating the bank_client dataset to store information of bank clients. The attributes included are namely - age, job, marital, education, default, housing, loan

In [37]:

```
bank_client = bnk.iloc[:, 0:7]
bank_client.head()
```

Out[37]:

	age	job	marital	education	default	housing	loan
0	56	housemaid	married	basic.4y	no	no	no
1	57	services	married	high.school	unknown	no	no
2	37	services	married	high.school	no	yes	no
3	40	admin.	married	basic.6y	no	no	no
4	56	services	married	high.school	no	no	yes

Exploring variables in bank_client

Age Count distribution


In [38]:

```
fig, bca = plt.subplots()
fig.set_size_inches(30, 15)
sns.countplot(x = 'age', data = bank_client)
bca.set_xlabel('Age', fontsize=25)
bca.set_ylabel('Count', fontsize=25)
bca.set_title('Age Count Distribution', fontsize=25)
```

Out[38]:

Text(0.5, 1.0, 'Age Count Distribution')


Age Count Distribution



Age Distribution

In [39]:

```
bca1 = sns.boxplot( y=bank_client["age"] )
bca1.set_xlabel('People Age', fontsize=15)
bca1.set_ylabel('Age', fontsize=15)
bca1.set_title('Age Distribution', fontsize=15)
bca1.tick_params(labels=15)
```




Jobs Count Distribution


In [40]:

```
fig, bca = plt.subplots()
fig.set_size_inches(30, 10)
sns.countplot(x = 'job', data = bank_client)
bca.set_xlabel('Jobs', fontsize=25)
bca.set_ylabel('Count', fontsize=25)
bca.set_title('Jobs Count Distribution', fontsize=25)
```

Out[40]:

Text(0.5, 1.0, 'Jobs Count Distribution')






Marital Status Count Distribution

In [41]:

```
fig, bca = plt.subplots()
fig.set_size_inches(10, 5)
sns.countplot(x = 'marital', data = bank_client)
bca.set_xlabel('Marital Status', fontsize=15)
bca.set_ylabel('Count', fontsize=15)
bca.set_title('Marital Status Count Distribution', fontsize=25)
```

Out[41]:

Text(0.5, 1.0, 'Marital Status Count Distribution')




Education Count Distribution


In [42]:

```
fig, bca = plt.subplots()
fig.set_size_inches(20, 5)
sns.countplot(x = 'education', data = bank_client)
bca.set_xlabel('Education', fontsize=15)
bca.set_ylabel('Count', fontsize=15)
bca.set_title('Education Count Distribution', fontsize=25)
```

Out[42]:

Text(0.5, 1.0, 'Education Count Distribution')






Housing Loan Count Distribution

In [43]:

```
fig, bca = plt.subplots()
fig.set_size_inches(10, 5)
sns.countplot(x = 'housing', data = bank_client)
bca.set_xlabel('Housing Loan', fontsize=15)
bca.set_ylabel('Count', fontsize=15)
bca.set_title('Housing Loan Count Distribution', fontsize=15)
```

Out[43]:

Text(0.5, 1.0, 'Housing Loan Count Distribution')




Personal Loan Count Distribution


In [44]:

```
fig, bca = plt.subplots()
fig.set_size_inches(10, 5)
sns.countplot(x = 'loan', data = bank_client)
bca.set_xlabel('Personal Loan', fontsize=15)
bca.set_ylabel('Count', fontsize=15)
bca.set_title('Personal Loan Count Distribution', fontsize=15)
```

Out[44]:

Text(0.5, 1.0, 'Personal Loan Count Distribution')






Default Credit Count Distribution

In [45]:

```
fig, bca = plt.subplots()
fig.set_size_inches(10, 5)
sns.countplot(x = 'default', data = bank_client)
bca.set_xlabel('Default Credit', fontsize=15)
bca.set_ylabel('Count', fontsize=15)
bca.set_title('Default Credit Count Distribution', fontsize=15)
```

Out[45]:

Text(0.5, 1.0, 'Default Credit Count Distribution')



Treating categorical variables

In [46]:

```
bank_client['job'].unique()
```

Out[46]:

```
array(['housemaid', 'services', 'admin.', 'blue-collar', 'technician',
       'retired', 'management', 'unemployed', 'self-employed', 'unknown',
       'entrepreneur', 'student'], dtype=object)
```

In [47]:

```
dummy = pd.get_dummies(bank_client['job'], prefix = 'Job_N_')
print(dummy)
```

	Job_N_admin.	Job_N_blue-collar	Job_N_entrepreneur	Job_N_housemaid	\
0	0	0	0	1	
1	0	0	0	0	
2	0	0	0	0	
3	1	0	0	0	
4	0	0	0	0	

```

...
41183      0      0      0      0
41184      0      1      0      0
41185      0      0      0      0
41186      0      0      0      0
41187      0      0      0      0

```

```

      Job_N_management  Job_N_retired  Job_N_self-employed  Job_N_services  \
0      0      0      0      0
1      0      0      0      1
2      0      0      0      1
3      0      0      0      0
4      0      0      0      1
...
41183      0      1      0      0
41184      0      0      0      0
41185      0      1      0      0
41186      0      0      0      0
41187      0      1      0      0

```

```

      Job_N_student  Job_N_technician  Job_N_unemployed  Job_N_unknown
0      0      0      0      0
1      0      0      0      0
2      0      0      0      0
3      0      0      0      0
4      0      0      0      0
...
41183      0      0      0      0
41184      0      0      0      0
41185      0      0      0      0
41186      0      1      0      0
41187      0      0      0      0

```

[41164 rows x 12 columns]

In [48]:

```

bank_client = bank_client.join(dummy)
bank_client

```

Out[48]:

	age	job	marital	education	default	housing	loan	Job_N_admin.	Job_N_blue-collar	Job_N_entrepreneur
0	56	housemaid	married	basic.4y	no	no	no	0	0	
1	57	services	married	high.school	unknown	no	no	0	0	
2	37	services	married	high.school	no	yes	no	0	0	
3	40	admin.	married	basic.6y	no	no	no	1	0	
4	56	services	married	high.school	no	no	yes	0	0	
...	
41183	73	retired	married	professional.course	no	yes	no	0	0	
41184	46	blue-collar	married	professional.course	no	no	no	0	1	
41185	56	retired	married	university.degree	no	yes	no	0	0	
41186	44	technician	married	professional.course	no	no	no	0	0	
41187	74	retired	married	professional.course	no	yes	no	0	0	

41164 rows x 19 columns

In [49]:

```

bank_client['marital'].unique()

```

Out[49]:

```
array(['married', 'single', 'divorced', 'unknown'], dtype=object)
```

In [50]:

```
lc=LabelEncoder()  
bank_client['Marital_N']=lc.fit_transform(bank_client['marital'])  
bank_client
```

Out[50]:

	age	job	marital	education	default	housing	loan	Job_N_admin.	Job_N_blue-collar	Job_N_entrepreneur
0	56	housemaid	married	basic.4y	no	no	no	0	0	
1	57	services	married	high.school	unknown	no	no	0	0	
2	37	services	married	high.school	no	yes	no	0	0	
3	40	admin.	married	basic.6y	no	no	no	1	0	
4	56	services	married	high.school	no	no	yes	0	0	
...
41183	73	retired	married	professional.course	no	yes	no	0	0	
41184	46	blue-collar	married	professional.course	no	no	no	0	1	
41185	56	retired	married	university.degree	no	yes	no	0	0	
41186	44	technician	married	professional.course	no	no	no	0	0	
41187	74	retired	married	professional.course	no	yes	no	0	0	

41164 rows x 20 columns



In [51]:

```
bank_client['education'].unique()
```

Out[51]:

```
array(['basic.4y', 'high.school', 'basic.6y', 'basic.9y',  
      'professional.course', 'unknown', 'university.degree',  
      'illiterate'], dtype=object)
```

In [52]:

```
bank_client=pd.concat((bank_client,pd.get_dummies(bank_client['education'])),axis=1)
```

In [53]:

```
bank_client['default'].unique()
```

Out[53]:

```
array(['no', 'unknown', 'yes'], dtype=object)
```

In [54]:

```
bank_client['housing'].unique()
```

Out[54]:

```
array(['no', 'yes', 'unknown'], dtype=object)
```

In [55]:

```
bank_client['loan'].unique()
```

Out[55]:

```
array(['no', 'yes', 'unknown'], dtype=object)
```

```
lc=LabelEncoder()  
bank_client['Default N']=lc.fit_transform(bank_client['default'])
```

```
lc=LabelEncoder()  
bank_client['Housing_N']=lc.fit_transform(bank_client['housing'])
```

```
lc=LabelEncoder()  
bank_client['Loan_N']=lc.fit_transform(bank_client['loan'])
```

```
bank_client.info()
```

In [60]:

```
bank_client = bank_client.drop(['job', 'marital', 'education', 'housing', 'default', 'loan'], axis = 1)
```

```
bank client.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 41164 entries, 0 to 41187
Data columns (total 25 columns):
 #   Column                Non-Null Count  Dtype

```



```

0   age                41164 non-null   int64
1   Job_N_admin.       41164 non-null   uint8
2   Job_N_blue-collar  41164 non-null   uint8
3   Job_N_entrepreneur 41164 non-null   uint8
4   Job_N_housemaid    41164 non-null   uint8
5   Job_N_management   41164 non-null   uint8
6   Job_N_retired      41164 non-null   uint8
7   Job_N_self-employed 41164 non-null   uint8
8   Job_N_services     41164 non-null   uint8
9   Job_N_student      41164 non-null   uint8
10  Job_N_technician   41164 non-null   uint8
11  Job_N_unemployed   41164 non-null   uint8
12  Job_N_unknown      41164 non-null   uint8
13  Marital_N          41164 non-null   int32
14  basic.4y           41164 non-null   uint8
15  basic.6y           41164 non-null   uint8
16  basic.9y           41164 non-null   uint8
17  high.school        41164 non-null   uint8
18  illiterate         41164 non-null   uint8
19  professional.course 41164 non-null   uint8
20  university.degree  41164 non-null   uint8
21  unknown            41164 non-null   uint8
22  Default_N          41164 non-null   int32
23  Housing_N          41164 non-null   int32
24  Loan_N             41164 non-null   int32

```

dtypes: int32(4), int64(1), uint8(20)

memory usage: 3.3 MB

In [62]:

```
bank_client['age'] = bank_client['age'].astype(int)
```

In [63]:

```
bank_client.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 41164 entries, 0 to 41187
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   age                41164 non-null   int32
1   Job_N_admin.       41164 non-null   uint8
2   Job_N_blue-collar  41164 non-null   uint8
3   Job_N_entrepreneur 41164 non-null   uint8
4   Job_N_housemaid    41164 non-null   uint8
5   Job_N_management   41164 non-null   uint8
6   Job_N_retired      41164 non-null   uint8
7   Job_N_self-employed 41164 non-null   uint8
8   Job_N_services     41164 non-null   uint8
9   Job_N_student      41164 non-null   uint8
10  Job_N_technician   41164 non-null   uint8
11  Job_N_unemployed   41164 non-null   uint8
12  Job_N_unknown      41164 non-null   uint8
13  Marital_N          41164 non-null   int32
14  basic.4y           41164 non-null   uint8
15  basic.6y           41164 non-null   uint8
16  basic.9y           41164 non-null   uint8
17  high.school        41164 non-null   uint8
18  illiterate         41164 non-null   uint8
19  professional.course 41164 non-null   uint8
20  university.degree  41164 non-null   uint8
21  unknown            41164 non-null   uint8
22  Default_N          41164 non-null   int32
23  Housing_N          41164 non-null   int32
24  Loan_N             41164 non-null   int32

```

dtypes: int32(5), uint8(20)

memory usage: 3.1 MB

In [64]:

```
def age(dataframe):
    dataframe.loc[dataframe['age'] <= 32, 'age'] = 1
    dataframe.loc[(dataframe['age'] > 32) & (dataframe['age'] <= 47), 'age'] = 2
    dataframe.loc[(dataframe['age'] > 47) & (dataframe['age'] <= 70), 'age'] = 3
    dataframe.loc[(dataframe['age'] > 70) & (dataframe['age'] <= 98), 'age'] = 4

    return dataframe

age(bank_client) ;
```

In [65]:

```
bank_client.head()
```

Out[65]:

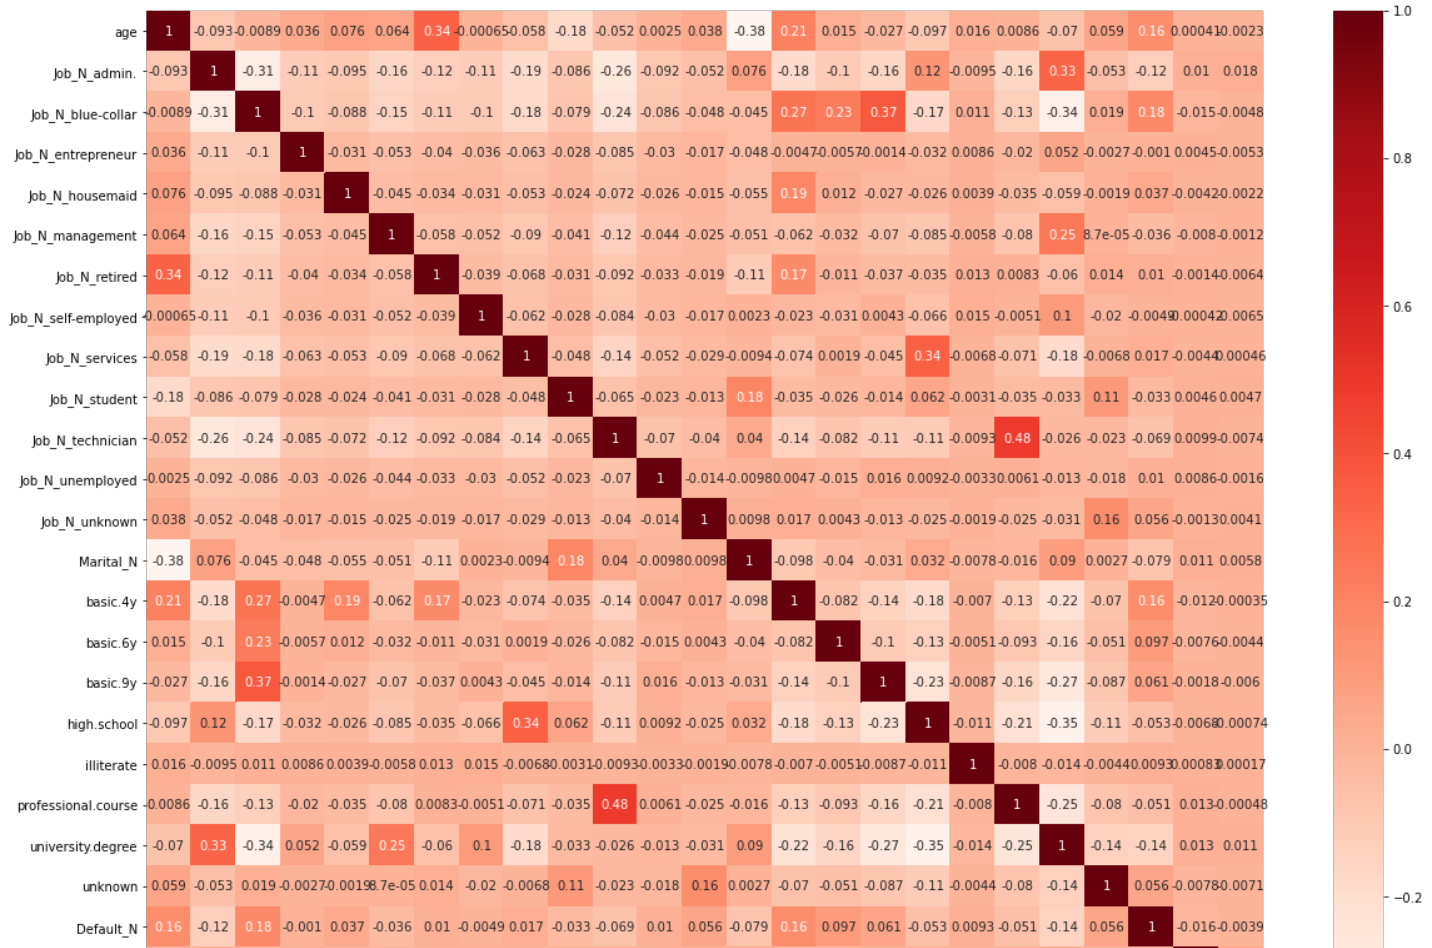
	age	Job_N_admin.	Job_N_blue-collar	Job_N_entrepreneur	Job_N_housemaid	Job_N_management	Job_N_retired	Job_N_self-employed
0	3	0	0	0	1	0	0	0
1	3	0	0	0	0	0	0	0
2	2	0	0	0	0	0	0	0
3	2	1	0	0	0	0	0	0
4	3	0	0	0	0	0	0	0

5 rows x 25 columns

Correlation between variables in bank_client dataset

In [66]:

```
plt.figure(figsize=(20,15))
cor = bank_client.corr()
sns.heatmap(cor, annot=True, cmap=plt.cm.Reds)
plt.show()
```



Creating new dataset 'other_attr'

We are now creating a dataset to store the attributes - contact, month, day_of_week and duration.

In [67]:

```
other_attr = bnk.iloc[:, 7:11]
other_attr.head()
```

Out[67]:

	contact	month	day_of_week	duration
0	telephone	may	mon	261.0
1	telephone	may	mon	149.0
2	telephone	may	mon	226.0
3	telephone	may	mon	151.0
4	telephone	may	mon	307.0

Check for null values

In [68]:

```
other_attr.isnull().sum()
```

Out[68]:


```
contact      0
month        0
day_of_week  0
duration     0
dtype: int64
```


Exploring the attributes in 'other_attr'

Calls duration

In [69]:

```
dur = sns.boxplot(x = 'duration', data = other_attr)
dur.set_xlabel('Calls', fontsize=15)
dur.set_ylabel('Duration', fontsize=15)
dur.set_title('Calls Distribution', fontsize=15)
dur.tick_params(labelsize=20)
```






Contacts Count

In [70]:

```
fig, bca = plt.subplots()
fig.set_size_inches(10, 5)
sns.countplot(x = 'contact', data = other_attr)
bca.set_xlabel('Contacts', fontsize=15)
bca.set_ylabel('Count', fontsize=15)
bca.set_title('Contacts Count', fontsize=15)
```

Out[70]:

Text(0.5, 1.0, 'Contacts Count')




Months Count


In [71]:

```
fig, bca = plt.subplots()
fig.set_size_inches(10, 5)
sns.countplot(x = 'month', data = other_attr, order = ['mar', 'apr', 'may', 'jun', 'jul',
, 'aug', 'sep', 'oct', 'nov', 'dec'])
bca.set_xlabel('Months', fontsize=15)
bca.set_ylabel('Count', fontsize=15)
bca.set_title('Months Count', fontsize=15)
```

Out[71]:

Text(0.5, 1.0, 'Months Count')






Days Of Week Count

In [72]:

```
fig, bca = plt.subplots()
fig.set_size_inches(10, 5)
sns.countplot(x = 'day_of_week', data = other_attr)
bca.set_xlabel('Days', fontsize=15)
bca.set_ylabel('Count', fontsize=15)
bca.set_title('DaysOfWeek Count', fontsize=15)
```

Out[72]:

Text(0.5, 1.0, 'DaysOfWeek Count')



Treating categorical variables

In [73]:

```
other_attr.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 41164 entries, 0 to 41187
Data columns (total 4 columns):
#   Column          Non-Null Count  Dtype
---  -
0   contact         41164 non-null  object
1   month           41164 non-null  object
2   day_of_week     41164 non-null  object
3   duration        41164 non-null  float64
dtypes: float64(1), object(3)
memory usage: 2.8+ MB
```

In [74]:

```
other_attr['contact'].unique()
```

```
Out[74]:
```

```
array(['telephone', 'cellular'], dtype=object)
```

```
In [75]:
```

```
other_attr['contact'] = other_attr['contact'].map({'telephone':1, 'cellular':2}).astype(int)
```

```
In [76]:
```

```
other_attr['month'].unique()
```

```
Out[76]:
```

```
array(['may', 'jun', 'jul', 'aug', 'oct', 'nov', 'dec', 'mar', 'apr',  
      'sep'], dtype=object)
```

```
In [77]:
```

```
other_attr["month"] = other_attr["month"].str.capitalize()
```

```
In [78]:
```

```
other_attr["month"] = pd.to_datetime(other_attr.month, format='%b', errors='coerce').dt.month  
other_attr = other_attr.sort_values(by="month")
```

```
In [79]:
```

```
other_attr['month'].unique()
```

```
Out[79]:
```

```
array([ 3,  4,  5,  6,  7,  8,  9, 10, 11, 12], dtype=int64)
```

```
In [80]:
```

```
other_attr['day_of_week'].unique()
```

```
Out[80]:
```

```
array(['tue', 'mon', 'thu', 'wed', 'fri'], dtype=object)
```

```
In [81]:
```

```
lc=LabelEncoder()  
other_attr['day_of_week']=lc.fit_transform(other_attr['day_of_week'])
```

```
In [82]:
```

```
other_attr['day_of_week'].unique()
```

```
Out[82]:
```

```
array([3, 1, 2, 4, 0])
```

```
In [83]:
```

```
other_attr['duration'] = other_attr['duration'].astype(int)
```

Creating new dataset 'cont_attr'

Here we are creating a new dataset for the social and economic context attributes, which are - emp.var.rate, cons.price.idx, cons.conf.idx, euribor3m, nr.employed.

```
In [84]:
```

```
cont_attr= bnk.loc[:, ['emp.var.rate', 'cons.price.idx', 'cons.conf.idx', 'euribor3m',
```

```
'nr.employed']]\ncont_attr.head()
```

Out[84]:

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	nr.employed
0	1.100000	93.994	-36.4	4.857	5191.0
1	0.079252	93.994	-36.4	4.857	5191.0
2	1.100000	93.994	-36.4	4.857	5191.0
3	1.100000	93.994	-36.4	4.857	5191.0
4	1.100000	93.994	-36.4	4.857	5191.0

Creating a dataset 'remain_attr'

This dataset contains all the remaining attributes (y excluded). These are - campaign, pdays, previous, poutcome

In [85]:

```
remain_attr = bnk.loc[:, ['campaign', 'pdays', 'previous', 'poutcome']]\nremain_attr.head()
```

Out[85]:

	campaign	pdays	previous	poutcome
0	2.570404	999.0	0.172596	nonexistent
1	1.000000	999.0	0.000000	nonexistent
2	1.000000	999.0	0.000000	nonexistent
3	1.000000	999.0	0.000000	nonexistent
4	1.000000	999.0	0.000000	nonexistent

Treating the categorical attributes

In [86]:

```
remain_attr['poutcome'].unique()
```

Out[86]:

```
array(['nonexistent', 'failure', 'success'], dtype=object)
```

In [87]:

```
remain_attr['poutcome'] = remain_attr['poutcome'].map({'nonexistent':1, 'failure':2, 'success':3}).astype(int)
```

Creating the 'final_bank' dataset

We will now merge/concat all the above datasets that we created and curated as per need into one final dataset for our analysis.

In [88]:

```
final_bank= pd.concat([bank_client, other_attr, cont_attr, remain_attr], axis = 1)
```

Exploring the dataset

```
final_bank.shape
```

```
final_bank.info()
```

```
final bank.isna().sum()
```

age	0
Job_N_admin.	0
Job_N_blue-collar	0
Job_N_entrepreneur	0
Job_N_housemaid	0
Job_N_management	0
Job_N_retired	0
Job_N_self-employed	0


```

Job_N_services      0
Job_N_student       0
Job_N_technician    0
Job_N_unemployed    0
Job_N_unknown       0
Marital_N           0
basic.4y            0
basic.6y            0
basic.9y            0
high.school         0
illiterate          0
professional.course 0
university.degree   0
unknown            0
Default_N           0
Housing_N           0
Loan_N              0
contact             0
month               0
day_of_week         0
duration            0
emp.var.rate        0
cons.price.idx      0
cons.conf.idx       0
euribor3m           0
nr.employed         0
campaign            0
pdays              0
previous            0
poutcome            0
dtype: int64

```

In [92]:

```
final_bank['campaign'].unique()
```

Out[92]:

```

array([ 2.57040373,  1.          ,  2.          ,  3.          ,  4.          ,
        5.          ,  6.          ,  7.          ,  8.          ,  9.          ,
       10.          , 11.          , 12.          , 13.          , 19.          ,
       18.          , 23.          , 14.          , 22.          , 25.          ,
       16.          , 17.          , 15.          , 20.          , 56.          ,
       42.          , 28.          , 26.          , 27.          , 32.          ,
       21.          , 24.          , 29.          , 31.          , 30.          ,
       35.          , 41.          , 37.          , 40.          , 33.          ,
       34.          , 43.          ])

```

In [93]:

```
final_bank['campaign'].fillna(final_bank['campaign'].mean(),inplace=True)
```

In [94]:

```
final_bank.isna().sum()
```

Out[94]:

```

age                0
Job_N_admin.       0
Job_N_blue-collar  0
Job_N_entrepreneur 0
Job_N_housemaid    0
Job_N_management   0
Job_N_retired       0
Job_N_self-employed 0
Job_N_services      0
Job_N_student       0
Job_N_technician    0
Job_N_unemployed    0
Job_N_unknown       0
Marital_N           0
basic.4y            0

```

```

basic.4y      0
basic.6y      0
basic.9y      0
high.school   0
illiterate    0
professional.course  0
university.degree  0
unknown       0
Default_N     0
Housing_N     0
Loan_N        0
contact       0
month         0
day_of_week   0
duration      0
emp.var.rate  0
cons.price.idx  0
cons.conf.idx  0
euribor3m     0
nr.employed   0
campaign      0
pdays        0
previous      0
poutcome      0
dtype: int64

```

In [95]:

```
final_bank.describe()
```

Out[95]:

	age	Job_N_admin.	Job_N_blue-collar	Job_N_entrepreneur	Job_N_housemaid	Job_N_management	Job_N_retired
count	41164.000000	41164.000000	41164.000000	41164.000000	41164.000000	41164.000000	41164.000000
mean	1.978598	0.253037	0.224759	0.035371	0.025751	0.071033	0.041687
std	0.735708	0.434757	0.417429	0.184717	0.158392	0.256883	0.199875
min	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	2.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
75%	2.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000
max	4.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

8 rows x 38 columns

Splitting the data

We already have our target variable stored in 'y' from the beginning. Also, we have separately curated our final_bank dataset. So, it does not contain our target variable y from our original dataset.

In [96]:

```

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(final_bank, y, test_size = 0.2, random_state = 0)

```

In [97]:

```

from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.metrics import confusion_matrix, accuracy_score
k_fold = KFold(n_splits=10, shuffle=True, random_state=0)

```

In [98]:

```
X_train.head()
```

Out[98]:

	age	Job_N_admin.	Job_N_blue-collar	Job_N_entrepreneur	Job_N_housemaid	Job_N_management	Job_N_retired	Job_N_s-emplo
20018	2	0	0	0	0	0	0	0
39695	1	1	0	0	0	0	0	0
17238	3	0	0	0	1	0	0	0
5924	3	0	1	0	0	0	0	0
34656	2	1	0	0	0	0	0	0

5 rows × 38 columns



Scaling the data

In our `final_bank` data, we can see that the minimum and maximum value ranges from quite high to quite low values. For this reason, we are scaling our data with `StandardScaler`. We do so to scale our features centred around the zero and have unit variance.

In [99]:

```
#train-test split
from sklearn.preprocessing import StandardScaler
sc_X = StandardScaler()
X_train = sc_X.fit_transform(X_train)
X_test = sc_X.transform(X_test)
```

Voting Classifier

In [100]:

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC

from sklearn.ensemble import VotingClassifier
```

Hard voting

In [101]:

```
log_reg_clf = LogisticRegression(random_state= 0, C = 100, max_iter = 1000)
log_reg_clf.fit(X_train, y_train)

dtree_clf = DecisionTreeClassifier(max_depth = 1, random_state = 0)
dtree_clf.fit(X_train, y_train)

svc_clf = SVC(C = 0.1, gamma = 0.01, probability = True, random_state= 0)
svc_clf.fit(X_train, y_train)

hard_voting_clf = VotingClassifier(estimators=[('lr', log_reg_clf), ('dt', dtree_clf), (
'svc', svc_clf)], voting='hard')
hard_voting_clf.fit(X_train, y_train)

from sklearn.metrics import accuracy_score
for clf in (log_reg_clf, dtree_clf, svc_clf, hard_voting_clf):
    clf.fit(X_train, y_train)
```

```
y_pred = clf.predict(X_test)
print(clf.__class__.__name__, '%.4f'%accuracy_score(y_test, y_pred))
```

```
LogisticRegression 0.9093
DecisionTreeClassifier 0.8876
SVC 0.8987
VotingClassifier 0.9014
```

Soft Voting

In [102]:

```
log_reg_clf = LogisticRegression(random_state= 0, C = 100, max_iter = 1000)
log_reg_clf.fit(X_train, y_train)

dtree_clf = DecisionTreeClassifier(max_depth = 1, random_state = 0)
dtree_clf.fit(X_train, y_train)

svc_clf = SVC(C = 0.1, gamma = 0.01, probability = True, random_state= 0)
svc_clf.fit(X_train, y_train)

soft_voting_clf = VotingClassifier(estimators=[('lr', log_reg_clf), ('dt', dtree_clf), (
'svc', svc_clf)], voting='soft')
soft_voting_clf.fit(X_train, y_train)

from sklearn.metrics import accuracy_score
for clf in (log_reg_clf, dtree_clf, svc_clf, soft_voting_clf):
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)
    print(clf.__class__.__name__, '%.4f'%accuracy_score(y_test, y_pred))
```

```
LogisticRegression 0.9093
DecisionTreeClassifier 0.8876
SVC 0.8987
VotingClassifier 0.9030
```

Bagging

Bagging for Decision Tree Classifier

In [104]:

```
from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier

dtree_clf = DecisionTreeClassifier(random_state=0)
dtree_bag_clf = BaggingClassifier(dtree_clf, n_estimators=500, max_samples=100, bootstrap
=True, n_jobs=-1, random_state=0)

dtree_bag_clf.fit(X_train, y_train)
y_pred = dtree_bag_clf.predict(X_test)
```

In [105]:

```
dtree_bag_clf.fit(X_train, y_train)

# train and test scores
print('Train score: %.2f'%dtree_bag_clf.score(X_train, y_train))
print('Test score: %.2f'%dtree_bag_clf.score(X_test, y_test))
```

```
Train score: 0.91
Test score: 0.91
```

In [106]:

```
print(confusion_matrix(y_test, dtree_bag_clf.predict(X_test) ))
from sklearn.metrics import classification_report
```

```
print(classification_report(y_train, dtree_bag_clf.predict(X_train)))
```

```
[[7128  180]
 [ 557  368]]
```

	precision	recall	f1-score	support
0	0.93	0.98	0.95	29218
1	0.67	0.39	0.49	3713
accuracy			0.91	32931
macro avg	0.80	0.68	0.72	32931
weighted avg	0.90	0.91	0.90	32931

Random Forest Classifier

GridSearch

In [107]:

```
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier

param_grid = {'n_estimators': [200, 300, 400, 500],
              'max_depth': np.arange(1, 10)}

rf_gridsearch = GridSearchCV(RandomForestClassifier(random_state=0), param_grid, cv=10,
                             return_train_score=True)
rf_gridsearch.fit(X_train, y_train)
print("Best parameters for RandomForest Clf: {}".format(rf_gridsearch.best_params_))
print("Best cross-validation score: {:.2f}".format(rf_gridsearch.best_score_))
```

Best parameters for RandomForest Clf: {'max_depth': 9, 'n_estimators': 500}
Best cross-validation score: 0.91

Random Forest Classifier

In [108]:

```
rf_clf = RandomForestClassifier(n_estimators=400, max_depth = 9, bootstrap=True, n_jobs=-1, random_state=0)
rf_clf.fit(X_train, y_train)

pred_rf = rf_clf.predict(X_test)

#train and test scores
print('Train score: {:.2f}'.format(rf_clf.score(X_train, y_train)))
print('Test score: {:.2f}'.format(rf_clf.score(X_test, y_test)))
```

Train score: 0.92
Test score: 0.91

In [109]:

```
print(confusion_matrix(y_test, rf_clf.predict(X_test)))
from sklearn.metrics import classification_report
print(classification_report(y_train, rf_clf.predict(X_train)))
```

```
[[7189  119]
 [ 641  284]]
```

	precision	recall	f1-score	support
0	0.93	0.99	0.96	29218
1	0.86	0.39	0.54	3713
accuracy			0.92	32931
macro avg	0.89	0.69	0.75	32931
weighted avg	0.92	0.92	0.91	32931

Pasting

Decision Tree Classifier

In [110]:

```
from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier

# pasting: bootstrap = False

dtree_clf = DecisionTreeClassifier(criterion = 'entropy', random_state=0)
dtree_bag_clf = BaggingClassifier(dtree_clf, n_estimators=500, max_samples=100, bootstrap=False, random_state=0)

dtree_bag_clf.fit(X_train, y_train)
y_pred = dtree_bag_clf.predict(X_test)

from sklearn.metrics import accuracy_score

# train and test scores
print('Train score: %.2f'%dtree_bag_clf.score(X_train, y_train))
print('Test score: %.2f'%dtree_bag_clf.score(X_test, y_test))
```

Train score: 0.91

Test score: 0.91

In [111]:

```
print(confusion_matrix(y_test, dtree_bag_clf.predict(X_test)))
from sklearn.metrics import classification_report
print(classification_report(y_train, dtree_bag_clf.predict(X_train)))
```

```
[[7172  136]
 [ 621  304]]
```

	precision	recall	f1-score	support
0	0.92	0.98	0.95	29218
1	0.70	0.32	0.44	3713
accuracy			0.91	32931
macro avg	0.81	0.65	0.69	32931
weighted avg	0.89	0.91	0.89	32931

SVC Classifier

In [112]:

```
from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier

# pasting: bootstrap = False

svc_clf = SVC(C = 0.1, gamma = 0.01, probability = True, random_state= 0)
svc_bag_clf = BaggingClassifier(svc_clf, n_estimators=500, max_samples=100, bootstrap=False, random_state=0)

svc_bag_clf.fit(X_train, y_train)
y_pred = svc_bag_clf.predict(X_test)

from sklearn.metrics import accuracy_score

#train and test scores
print('Train score: %.2f'%svc_bag_clf.score(X_train, y_train))
```

```
print('Test score: %.2f'%svc_bag_clf.score(X_test, y_test))
```

Train score: 0.89

Test score: 0.89

In [113]:

```
print(confusion_matrix(y_test, svc_bag_clf.predict(X_test)))
from sklearn.metrics import classification_report
print(classification_report(y_train, svc_bag_clf.predict(X_train)))
```

```
[[7289   19]
 [ 875   50]]
```

		precision	recall	f1-score	support
	0	0.89	1.00	0.94	29218
	1	0.86	0.06	0.11	3713
	accuracy			0.89	32931
	macro avg	0.87	0.53	0.53	32931
	weighted avg	0.89	0.89	0.85	32931

ADA Boost Classifier

Decision Tree

In [114]:

```
from sklearn.ensemble import AdaBoostClassifier

dtree_ada_clf = AdaBoostClassifier(DecisionTreeClassifier(max_depth=1), n_estimators=200,
                                   algorithm="SAMME.R", learning_rate=0.5, random_state=0)

dtree_ada_clf.fit(X_train, y_train)
predictions = dtree_ada_clf.predict(X_test)

#train and test scores
print('Train score: %.2f'%dtree_ada_clf.score(X_train, y_train))
print('Test score: %.2f'%dtree_ada_clf.score(X_test, y_test))
```

Train score: 0.91

Test score: 0.91

In [115]:

```
confusion_matrix(y_test, predictions)
```

Out[115]:

```
array([[7108, 200],
       [ 560, 365]], dtype=int64)
```

In [116]:

```
print(confusion_matrix(y_test, dtree_ada_clf.predict(X_test)))
from sklearn.metrics import classification_report
print(classification_report(y_train, dtree_ada_clf.predict(X_train)))
```

```
[[7108  200]
 [ 560  365]]
```

		precision	recall	f1-score	support
	0	0.93	0.97	0.95	29218
	1	0.66	0.40	0.50	3713
	accuracy			0.91	32931
	macro avg	0.79	0.69	0.72	32931

weighted avg 0.90 0.91 0.90 32931

Logistic Regression

In [117]:

```
from sklearn.ensemble import AdaBoostClassifier

log_reg_ada_clf = AdaBoostClassifier(LogisticRegression(solver='liblinear'), n_estimators
=500,

                                   algorithm="SAMME.R", learning_rate=0.1, random_state=0)

log_reg_ada_clf.fit(X_train, y_train)
predictions = log_reg_ada_clf.predict(X_test)

#train and test scores
print('Train score: %.2f'%log_reg_ada_clf.score(X_train, y_train))
print('Test score: %.2f'%log_reg_ada_clf.score(X_test, y_test))
```

Train score: 0.91

Test score: 0.91

In [118]:

```
print(confusion_matrix(y_test, log_reg_ada_clf.predict(X_test)))
from sklearn.metrics import classification_report
print(classification_report(y_train, log_reg_ada_clf.predict(X_train)))
```

```
[[7194  114]
 [ 663  262]]
```

		precision	recall	f1-score	support
	0	0.92	0.98	0.95	29218
	1	0.70	0.28	0.40	3713
accuracy				0.91	32931
macro avg		0.81	0.63	0.68	32931
weighted avg		0.89	0.91	0.89	32931

Gradient Boosting Classifier

In [119]:

```
from sklearn.ensemble import GradientBoostingClassifier

gbrt = GradientBoostingClassifier(random_state=0, max_depth=5, learning_rate=0.01)
gbrt.fit(X_train, y_train)

#train and test scores
print("Accuracy on training set: {:.3f}".format(gbrt.score(X_train, y_train)))
print("Accuracy on test set: {:.3f}".format(gbrt.score(X_test, y_test)))
```

Accuracy on training set: 0.904

Accuracy on test set: 0.901

In [120]:

```
print(confusion_matrix(y_test, gbrt.predict(X_test)))
from sklearn.metrics import classification_report
print(classification_report(y_train, gbrt.predict(X_train)))
```

```
[[7276   32]
 [ 779  146]]
```

		precision	recall	f1-score	support
	0	0.90	1.00	0.95	29218
	1	0.87	0.17	0.28	3713

accuracy			0.90	32931
macro avg	0.89	0.58	0.62	32931
weighted avg	0.90	0.90	0.87	32931

PCA

In [121]:

```
from sklearn.decomposition import PCA

pca = PCA()
pca.fit(X_train)

cumsum = np.cumsum(pca.explained_variance_ratio_)
d = np.argmax(cumsum >= 0.95) + 1
```

In [122]:

```
d
```

Out[122]:

```
29
```

In [123]:

```
pca = PCA(n_components=0.95)
X_reduced = pca.fit_transform(X_train)
```

In [124]:

```
pca.n_components_
```

Out[124]:

```
29
```

In [125]:

```
np.sum(pca.explained_variance_ratio_)
```

Out[125]:

```
0.9538315288981813
```

In [126]:

```
pca = PCA(n_components = 29)
X_reduced = pca.fit_transform(X_train)
X_recovered = pca.inverse_transform(X_reduced)
```

In [127]:

```
X_reduced_pca = X_reduced
```

In [128]:

```
from sklearn.decomposition import IncrementalPCA

n_batches = 100

inc_pca = IncrementalPCA(n_components=29)
for X_batch in np.array_split(X_train, n_batches):
    print(".", end="")
    inc_pca.partial_fit(X_batch)

X_train_reduced = inc_pca.transform(X_train)
```

.....
.....

In [129]:

```
from sklearn.decomposition import IncrementalPCA

n_batches = 100

inc_pca = IncrementalPCA(n_components=29)
for X_batch in np.array_split(X_test, n_batches):
    print(".", end="")
    inc_pca.partial_fit(X_batch)

X_test_reduced = inc_pca.transform(X_test)
```

.....
.....

In [130]:

```
X_train_reduced.shape
```

Out[130]:

```
(32931, 29)
```

In [131]:

```
X_test_reduced.shape
```

Out[131]:

```
(8233, 29)
```

Models on PCA data

Logistic Regression

In [132]:

```
from sklearn.linear_model import LogisticRegression
log_model = LogisticRegression()
log_model.fit(X_train_reduced, y_train)
log_pred = log_model.predict(X_test_reduced)
```

In [133]:

```
print(confusion_matrix(y_test, log_pred))
from sklearn.metrics import classification_report
print(classification_report(y_train, log_model.predict(X_train_reduced)))
```

```
[[7086  222]
 [ 711  214]]
```

	precision	recall	f1-score	support
0	0.92	0.98	0.95	29218
1	0.65	0.36	0.46	3713
accuracy			0.91	32931
macro avg	0.79	0.67	0.70	32931
weighted avg	0.89	0.91	0.89	32931

In [134]:

```
lr_score_train = log_model.score(X_train_reduced, y_train)
lr_score_train
```

Out[134]:

```
0.9056208435820352
```

In [135]:

```
lr_score_test = log_model.score(X_test_reduced,y_test)
lr_score_test
```

Out[135]:

```
0.886675573909875
```

KNN Classifier

In [136]:

```
from sklearn import model_selection
from sklearn.neighbors import KNeighborsClassifier
```

```
neighbors = np.arange(0,25)
```

```
cv_scores = []
```

In [137]:


```
# To determine best k-value
for k in neighbors:
    k_val = k+1
    knn_clf = KNeighborsClassifier(n_neighbors = k_val, weights='uniform', p=2, metric='
euclidean')
    k_fold = model_selection.KFold(n_splits=10, random_state=123)
    cross_val_scores = model_selection.cross_val_score(knn_clf, X_train_reduced, y_train
, cv=k_fold, scoring='accuracy')
    cv_scores.append(cross_val_scores.mean()*100)
    print("k=%d %0.2f (+/- %0.2f)" % (k_val, cross_val_scores.mean()*100, cross_val_scor
es.std()*100))
```

```
optimal_kval = neighbors[cv_scores.index(max(cv_scores))]
print ("The optimal number of neighbors is %d with %0.1f%%" % (optimal_kval, cv_scores[o
ptimal_kval]))
```

```
plt.plot(neighbors, cv_scores)
plt.xlabel('Number of Neighbors K')
plt.ylabel('Train Accuracy')
plt.show()
```

```
k=1 87.67 (+/- 0.61)
k=2 89.31 (+/- 0.62)
k=3 89.17 (+/- 0.43)
k=4 89.55 (+/- 0.53)
k=5 89.57 (+/- 0.44)
k=6 89.71 (+/- 0.47)
k=7 89.66 (+/- 0.35)
k=8 89.78 (+/- 0.45)
k=9 89.79 (+/- 0.43)
k=10 89.91 (+/- 0.37)
k=11 89.87 (+/- 0.35)
k=12 89.90 (+/- 0.45)
k=13 89.92 (+/- 0.45)
k=14 89.93 (+/- 0.46)
k=15 89.96 (+/- 0.40)
k=16 90.00 (+/- 0.46)
k=17 90.03 (+/- 0.39)
k=18 90.01 (+/- 0.46)
k=19 90.04 (+/- 0.46)
k=20 90.01 (+/- 0.53)
k=21 90.08 (+/- 0.48)
k=22 90.00 (+/- 0.50)
k=23 90.14 (+/- 0.46)
k=24 90.06 (+/- 0.44)
```

k=25 90.10 (+/- 0.47)
The optimal number of neighbors is 22 with 90.1%



In [138]:

```
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=24)
knn.fit(X_train_reduced, y_train)
knn_pred = knn.predict(X_test_reduced)
```

In [139]:

```
# for train data
from sklearn.metrics import classification_report
print(classification_report(y_train, knn.predict(X_train_reduced)))
```

	precision	recall	f1-score	support
0	0.91	0.99	0.95	29218
1	0.71	0.25	0.37	3713
accuracy			0.90	32931
macro avg	0.81	0.62	0.66	32931
weighted avg	0.89	0.90	0.88	32931

In [140]:

```
# for test data
confusion_matrix(y_test, knn.predict(X_test_reduced))
```

Out[140]:

```
array([[7249,  59],
       [ 875,  50]], dtype=int64)
```

In [141]:

```
print(classification_report(y_test, knn.predict(X_test_reduced)))
```

	precision	recall	f1-score	support
0	0.89	0.99	0.94	7308
1	0.46	0.05	0.10	925
accuracy			0.89	8233
macro avg	0.68	0.52	0.52	8233
weighted avg	0.84	0.89	0.84	8233

In [142]:

```
knn_score_test = knn.score(X_test_reduced, y_test)
knn_score_test
```

Out[142]:

```
0.88655411150249
```

Linear SVM

In [143]:

```
from sklearn.svm import LinearSVC

svm = LinearSVC()
svm.fit(X_train_reduced, y_train)
svc_pred = svm.predict(X_test_reduced)
```

In [144]:

```
# for train set
from sklearn.metrics import classification_report
print(classification_report(y_train, svm.predict(X_train_reduced)))
```

	precision	recall	f1-score	support
0	0.92	0.98	0.95	29218
1	0.67	0.31	0.42	3713
accuracy			0.90	32931
macro avg	0.79	0.64	0.68	32931
weighted avg	0.89	0.90	0.89	32931

In [145]:

```
#for test set
from sklearn.metrics import classification_report
print(classification_report(y_test, svm.predict(X_test_reduced)))
```

	precision	recall	f1-score	support
0	0.90	0.98	0.94	7308
1	0.48	0.17	0.25	925
accuracy			0.89	8233
macro avg	0.69	0.57	0.59	8233
weighted avg	0.86	0.89	0.86	8233

In [146]:

```
svm_score_train = svm.score(X_train_reduced, y_train)
svm_score_train
```

Out[146]:

```
0.904588381767939
```

In [147]:

```
svm_score_test = svm.score(X_test_reduced, y_test)
svm_score_test
```

Out[147]:

```
0.8860682618729503
```

Decision Tree Classifier

In [148]:

```
d_tree = DecisionTreeClassifier(criterion='entropy', random_state=0)
d_tree.fit(X_train_reduced, y_train)
```

```
d_tree_pred = d_tree.predict(X_test_reduced)
```

In [149]:

```
print(classification_report(y_test, d_tree.predict(X_test_reduced)))
```

	precision	recall	f1-score	support
0	0.91	0.88	0.89	7308
1	0.23	0.28	0.26	925
accuracy			0.81	8233
macro avg	0.57	0.58	0.57	8233
weighted avg	0.83	0.81	0.82	8233

In [150]:

```
print(classification_report(y_train, d_tree.predict(X_train_reduced)))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	29218
1	1.00	1.00	1.00	3713
accuracy			1.00	32931
macro avg	1.00	1.00	1.00	32931
weighted avg	1.00	1.00	1.00	32931

In [151]:

```
d_tree_score_train = d_tree.score(X_train_reduced,y_train)
d_tree_score_train
```

Out[151]:

0.9999696334760559

In [152]:

```
d_tree_score_test = d_tree.score(X_test_reduced,y_test)
d_tree_score_test
```

Out[152]:

0.814283979108466

Kernalized SVM(linear, rbf, poly)

In [153]:

```
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV

tuned_parameters = [{'kernel': ['rbf'], 'gamma': [0.1],
                      'C': [1]},
                    {'kernel': ['linear'], 'C': [1]},
                    {'kernel': ['poly'], 'degree':[3], 'C':[10] }]
```

In [154]:

```
clf = GridSearchCV(SVC(), tuned_parameters, cv=5, scoring='precision')
clf.fit(X_train, y_train)

print(clf.cv_results_)
```

```
{'mean_fit_time': array([91.10583048, 70.27198591, 80.00121274]), 'std_fit_time': array([
1.04248985, 7.38563372, 1.80327034]), 'mean_score_time': array([5.95953856, 2.66312752, 2
.90309567]), 'std_score_time': array([0.01897421, 0.03878049, 0.03692636]), 'param_C': ma
sked_array(data=[1, 1, 10],
```

```

        mask=[False, False, False],
        fill_value='?',
        dtype=object), 'param_gamma': masked_array(data=[0.1, --, --],
        mask=[False, True, True],
        fill_value='?',
        dtype=object), 'param_kernel': masked_array(data=['rbf', 'linear', 'poly'],
        mask=[False, False, False],
        fill_value='?',
        dtype=object), 'param_degree': masked_array(data=[--, --, 3],
        mask=[ True, True, False],
        fill_value='?',
        dtype=object), 'params': [{ 'C': 1, 'gamma': 0.1, 'kernel': 'rbf'}, { 'C': 1, '
kernel': 'linear'}, { 'C': 10, 'degree': 3, 'kernel': 'poly'}], 'split0_test_score': array
([0.67412141, 0.63099631, 0.60408163]), 'split1_test_score': array([0.63467492, 0.6293929
7, 0.53682171]), 'split2_test_score': array([0.62170088, 0.65116279, 0.56097561]), 'split
3_test_score': array([0.63276836, 0.64312268, 0.59578544]), 'split4_test_score': array([0
.59744409, 0.62666667, 0.5462963 ]), 'mean_test_score': array([0.63214193, 0.63626828, 0.
56879214]), 'std_test_score': array([0.02482757, 0.00933899, 0.02669553]), 'rank_test_sco
re': array([2, 1, 3])}

```

In [155]:

```

print('The best model is: ', clf.best_params_)
print('This model produces a mean cross-validated score (precision) of', clf.best_score_)

```

The best model is: { 'C': 1, 'kernel': 'linear'}

This model produces a mean cross-validated score (precision) of 0.6362682830306745

In [156]:

```

svm_ker_lin = SVC(kernel='linear', C=1)
svm_ker_rbf = SVC(kernel='rbf', gamma=0.1, C=1)
svm_ker_poly = SVC(kernel='poly', degree=3, C=10)

```

In [157]:

```

svm_ker_lin.fit(X_train_reduced, y_train)
svm_ker_rbf.fit(X_train_reduced, y_train)
svm_ker_poly.fit(X_train_reduced, y_train)

ker_lin_pred = svm.predict(X_test_reduced)
ker_rbf_pred = svm.predict(X_test_reduced)
ker_poly_pred = svm.predict(X_test_reduced)

```

In [158]:

```

print(classification_report(y_train, svm_ker_lin.predict(X_train_reduced)))

```

	precision	recall	f1-score	support
0	0.92	0.98	0.95	29218
1	0.65	0.29	0.41	3713
accuracy			0.90	32931
macro avg	0.79	0.64	0.68	32931
weighted avg	0.89	0.90	0.89	32931

In [159]:

```

print(classification_report(y_train, svm_ker_rbf.predict(X_train_reduced)))

```

	precision	recall	f1-score	support
0	0.94	0.99	0.97	29218
1	0.89	0.52	0.66	3713
accuracy			0.94	32931
macro avg	0.92	0.76	0.81	32931
weighted avg	0.94	0.94	0.93	32931

In [160]:

```
print(classification_report(y_train, svm_ker_poly.predict(X_train_reduced)))
```

	precision	recall	f1-score	support
0	0.94	0.98	0.96	29218
1	0.80	0.50	0.62	3713
accuracy			0.93	32931
macro avg	0.87	0.74	0.79	32931
weighted avg	0.92	0.93	0.92	32931

In [161]:

```
print(classification_report(y_test, svm_ker_poly.predict(X_test_reduced)))
```

	precision	recall	f1-score	support
0	0.91	0.91	0.91	7308
1	0.30	0.32	0.31	925
accuracy			0.84	8233
macro avg	0.61	0.61	0.61	8233
weighted avg	0.84	0.84	0.84	8233

In [162]:

```
print(classification_report(y_test, svm_ker_lin.predict(X_test_reduced)))
```

	precision	recall	f1-score	support
0	0.90	0.98	0.94	7308
1	0.53	0.15	0.24	925
accuracy			0.89	8233
macro avg	0.72	0.57	0.59	8233
weighted avg	0.86	0.89	0.86	8233

In [163]:

```
print(classification_report(y_test, svm_ker_rbf.predict(X_test_reduced)))
```

	precision	recall	f1-score	support
0	0.89	0.99	0.94	7308
1	0.31	0.02	0.04	925
accuracy			0.88	8233
macro avg	0.60	0.51	0.49	8233
weighted avg	0.82	0.88	0.84	8233

In [164]:

```
# Since linear kernel is our best model, we will consider it's train and test scores.
#test score
svm_ker_lin_score_test = svm_ker_lin.score(X_test_reduced,y_test)
svm_ker_lin_score_test
```

Out[164]:

0.8895906716871128

In [165]:

```
#train score
svm_ker_lin_score_train = svm_ker_lin.score(X_train_reduced,y_train)
```



```
svm_ker_lin_score_train
```

```
Out[165]:
```

```
0.9028574899031307
```

Results from Project 1 for all models :

Train scores for our models are as follows :

Logistic Regression : 0.9092951929792596

KNN Classification : 0.9031915216665148

Linear SVM : 0.9085056633567156

Kernalized SVM : 0.8997904709847864

Decision Tree Classifier : 0.9999696334760559

Test scores for our models are as follows :

Logistic Regression : 0.9086602696465444

KNN Classification : 0.9001579011296004

Linear SVM : 0.9069597959431556

Kernalized SVM : 0.9086602696465444

Decision Tree Classifier : 0.8933560063160452

As we can see from our train and test scores from Project 1 and our train and test scores when using our PCA reduced dataset, the scores have dropped. Hence, we can say that the dimensionally reduced dataset results in a poorer score compared to the original dataset. Although, PCA did help in saving computational time by reducing the features to 29 from 38.

```
In [166]:
```

```
import numpy
from keras.models import Sequential
from keras.layers import Dense
from keras.wrappers.scikit_learn import KerasClassifier

# create model
clf_model = Sequential()
clf_model.add(Dense(12, input_dim=45, activation='relu'))
clf_model.add(Dense(8, activation='relu'))
clf_model.add(Dense(1, activation='sigmoid'))

# Compile model
clf_model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
return model

# Fit the model
clf_model.fit(X_train, y_train, epochs=150, batch_size=10)

# evaluate the model
clf_model_scores = clf_model.evaluate(X_test, y_test)
print("\n%s: %.2f%%" % (clf_model.metrics_names[1], scores[1]*100))
```

ModuleNotFoundError

Traceback (most recent call last)

<ipython-input-166-2a2fa6886b37> in <module>

```
1 import numpy
----> 2 from keras.models import Sequential
3 from keras.layers import Dense
4 from keras.wrappers.scikit_learn import KerasClassifier
5
```

ModuleNotFoundError: No module named 'keras'

In []:

```
y_predict = clf_model.predict(X_test)
y_predict
```

In []: