

# AIRBNB Price Prediction in Newyork City

In this kernel we are focusing Price Prediction of New York City Airbnb Open Data Airbnb listings and metrics in NYC, NY, USA (2019) for linear regression.

Data Since 2008, guests and hosts have used Airbnb to expand on traveling possibilities and present more unique, personalized way of experiencing the world. This dataset describes the listing activity and metrics in NYC, NY for 2019. This data file includes all needed information to find out more about hosts, geographical availability, necessary metrics to make predictions and draw conclusions.

This data contains 16 columns, 4852 unique values(samples). Imported all necessary files and libraries, We removed unnecessary data from the dataset like last review, reviews per month and host name as they donot support the data required. We filled the null values with zero constant and did the visualization using seaborn, pyplot, matplotlib.

Variables id: listing ID name: name of the listing host\_id: host ID host\_name: name of the host neighbourhood\_group: location neighbourhood: area latitude: latitude coordinates longitude: longitude coordinates room\_type: listing space type price: price in dollars minimum\_nights: amount of nights minimum number\_of\_reviews: number of reviews last\_review: latest review reviews\_per\_month: number of reviews per month calculated\_host\_listings\_count: amount of listing per host availability\_365: number of days when listing is available for booking

We will perform Regression on this dataset.

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns

from scipy.stats import norm
from scipy import stats
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn import preprocessing
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
from sklearn import preprocessing
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from sklearn.svm import LinearSVR
from sklearn.svm import SVR
from sklearn import metrics

import warnings
warnings.filterwarnings('ignore')
```

In [2]:

```
my_data = pd.read_csv('NYC_AirBNB.csv')
```

In [3]:

```
my_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4852 entries, 0 to 4851
Data columns (total 30 columns):
Unnamed: 0      4852 non-null int64
id              4852 non-null int64
log_price      4852 non-null float64
property_type   4852 non-null object
room_type      4852 non-null object
amenities      4852 non-null object
accommodates   4852 non-null int64
bathrooms      4838 non-null float64
bed_type       4852 non-null object
```

```

cancellation_policy    4852 non-null object
cleaning_fee            4852 non-null object
city                   4852 non-null object
description             4852 non-null object
first_review           3833 non-null object
host_has_profile_pic    4814 non-null object
host_identity_verified  4814 non-null object
host_response_rate     3348 non-null object
host_since             4814 non-null object
instant_bookable       4852 non-null object
last_review            3837 non-null object
latitude               4852 non-null float64
longitude              4852 non-null float64
name                   4852 non-null object
neighbourhood          4852 non-null object
number_of_reviews      4852 non-null int64
review_scores_rating    3759 non-null float64
thumbnail_url          4494 non-null object
zipcode                4789 non-null object
bedrooms               4850 non-null float64
beds                   4840 non-null float64
dtypes: float64(7), int64(4), object(19)
memory usage: 1.1+ MB

```

In [4]:

```
my_data.head()
```

Out[4]:

Unnamed: 0	id	log_price	property_type	room_type	amenities	accommodates	bathrooms	bed_type	cancellation_
0	16553	44472	4.382027	Condominium	Private room	{TV,"Cable Internet","Wireless Internet","A...	1	1.0	Real Bed
1	12555	7007348	5.075174	Apartment	Entire home/apt	{TV,Internet,"Wireless Internet","Air conditio...	3	1.0	Real Bed
2	15012	10283574	4.852030	Apartment	Entire home/apt	{"Wireless Internet","Air conditioning",Kitch...	2	1.0	Real Bed
3	21502	1754527	5.010635	Apartment	Entire home/apt	{TV,Internet,"Wireless Internet","Air conditio...	4	1.0	Real Bed
4	13431	16823953	4.317488	Apartment	Private room	{TV,"Wireless Internet","Air conditioning",Kit...	2	1.0	Real Bed

5 rows × 30 columns

In [5]:

```

def calculate_nullvalue(df):
    """
        Calculating percentage of null value present in our data set.
    """

    missing_value = df.isnull().sum().to_frame()
    missing_value.columns = ['num_nullvalue']
    missing_value['percentge_missing'] = np.round(100 * (missing_value['num_nullvalue'] / df.shape[0]))
    missing_value.sort_values(by='num_nullvalue', ascending=False, inplace=True)

    return missing_value

```

In [6]:

```

num_nullvalue = calculate_nullvalue(my_data)
num_nullvalue

```

Out [6]:

	num_nullvalue	percentage_missing
host_response_rate	1504	31.0
review_scores_rating	1093	23.0
first_review	1019	21.0
last_review	1015	21.0
thumbnail_url	358	7.0
zipcode	63	1.0
host_identity_verified	38	1.0
host_since	38	1.0
host_has_profile_pic	38	1.0
bathrooms	14	0.0
beds	12	0.0
bedrooms	2	0.0
city	0	0.0
longitude	0	0.0
log_price	0	0.0
property_type	0	0.0
room_type	0	0.0
number_of_reviews	0	0.0
neighbourhood	0	0.0
name	0	0.0
latitude	0	0.0
description	0	0.0
amenities	0	0.0
instant_bookable	0	0.0
accommodates	0	0.0
bed_type	0	0.0
id	0	0.0
cancellation_policy	0	0.0
cleaning_fee	0	0.0
Unnamed: 0	0	0.0

In [7]:

```
#Finding the missing values in the dataframe
my_data.isnull().sum()
```

Out [7]:

```
Unnamed: 0      0
id              0
log_price       0
property_type   0
room_type       0
amenities       0
accommodates    0
bathrooms      14
bed_type        0
cancellation_policy  0
cleaning_fee    0
city           0
description     0
first_review   1019
host_has_profile_pic  38
host_identity_verified  38
```

```

host_response_rate    1504
host_since             38
instant_bookable       0
last_review           1015
latitude              0
longitude              0
name                  0
neighbourhood          0
number_of_reviews      0
review_scores_rating   1093
thumbnail_url         358
zipcode               63
bedrooms               2
beds                  12
dtype: int64

```

In [8]:

```

Conversion = {'Condo':['Timeshare','Loft','Guest suite','Condominium','Serviced apartment'],
              'Housing':['Vacation home','Townhouse','Casa particular','Villa','In-law'],
              'Hotel type 1':['Dorm','Guesthouse','Hostel'],
              'Hotel type 2':['Bed & Breakfast','Boutique hotel'],
              'Other':['Island','Yurt','Hut','Treehouse',
                      'Earth House','Tipi','Train','Parking Space','Lighthouse',
                      'Cabin','Camper/RV','Bungalow','Cave','Castle','Chalet','Boat','Tent']}
Conversion_real = {i : k for k, v in Conversion.items() for i in v}
my_data['property_type'].replace(Conversion_real,inplace =True)

```

In [9]:

```


plt.figure(figsize=(15,12))
sns.scatterplot(x='room_type', y='log_price', data=my_data)


plt.xlabel("Room Type", size=10)
plt.ylabel("Price", size=10)
plt.title("Room Type vs Price",size=15, weight='bold')

```

Out[9]:

Text(0.5, 1.0, 'Room Type vs Price')





In [10]:

```
top10_freq_neighbourhood=my_data.neighbourhood.value_counts().head(10)
print(top10_freq_neighbourhood)
```

```
Williamsburg      442
Bedford-Stuyvesant 296
Bushwick          248
Upper West Side   210
Harlem            208
Upper East Side   185
Crown Heights     184
Hell's Kitchen    178
Lower East Side   137
East Harlem       137
Name: neighbourhood, dtype: int64
```

In [11]:

```
top10_freq_neighbourhood_data=my_data[my_data['neighbourhood'].isin(['Williamsburg','Bedford-Stuyvesant','Harlem','Bushwick','Upper West Side','Hell\'s Kitchen','East Village','Upper East Side','Crown Heights','Midtown'])]
top10_freq_neighbourhood_data
```

Out[11]:

Unnamed: 0	id	log_price	property_type	room_type	amenities	accommodates	bathrooms	bed_type	
3	21502	1754527	5.010635	Apartment	Entire home/apt	{TV,Internet,"Wireless Internet","Air conditioning",Kitchen}	4	1.0	Real Bn
4	13431	16823953	4.317488	Apartment	Private room	{TV,"Wireless Internet","Air conditioning",Kitchen}	2	1.0	Real Bn
6	5389	2636988	4.262680	Apartment	Entire home/apt	{"Wireless Internet","Air conditioning",Heating,Essentials}	2	1.0	Real Bn
7	21857	20965965	4.248495	Apartment	Private room	{TV,Internet,"Wireless Internet","Air conditioning",Kitchen}	2	1.0	Real Bn
10	29298	2182851	3.806662	Apartment	Private room	{"Pets live on this property",Cat(s),"Smoke de...	2	1.0	Real Bn
...	...	...	...	...	...	...	...	...	
4842	24714	21023168	4.867534	Apartment	Entire home/apt	{"Wireless Internet","Air conditioning",Kitchen}	5	1.0	Real Bn
4843	8047	11145641	4.276666	Apartment	Private room	{TV,Internet,"Wireless Internet","Air conditioning",Kitchen}	2	1.0	Real Bn
4845	17804	14139651	3.688879	Apartment	Private room	{"Wireless Internet",Kitchen,Heating,Essentials}	2	1.0	Real Bn
4846	5608	12524258	4.094345	Apartment	Private room	{Internet,"Wireless Internet",Kitchen,Heating,Essentials}	4	1.0	Real Bn
...	...	...	...	...	...	...	...	...	

Unnamed: 0 20544 19120020 log\_price property\_type room\_type Private room amenities accommodates bathrooms bedrooms latitude longitude

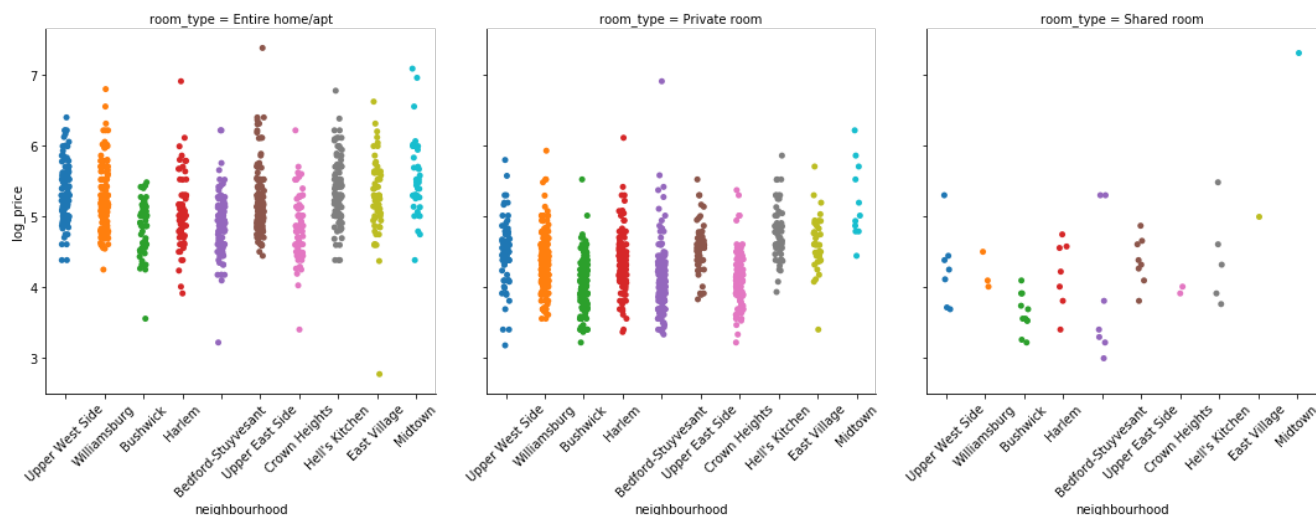
2111 rows x 30 columns

In [12]:

```
t=sns.catplot(x="neighbourhood", y="log_price", col="room_type", data=top10_freq_neighbourhood_data)
t.set_xticklabels(rotation=45)
```

Out[12]:

<seaborn.axisgrid.FacetGrid at 0x294aee9bb48>





In [13]:

```
plt.figure(figsize=(20,20))
sns.heatmap(my_data.corr(), annot=True, cmap="YlGnBu")
```

Out[13]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x294af20af48>



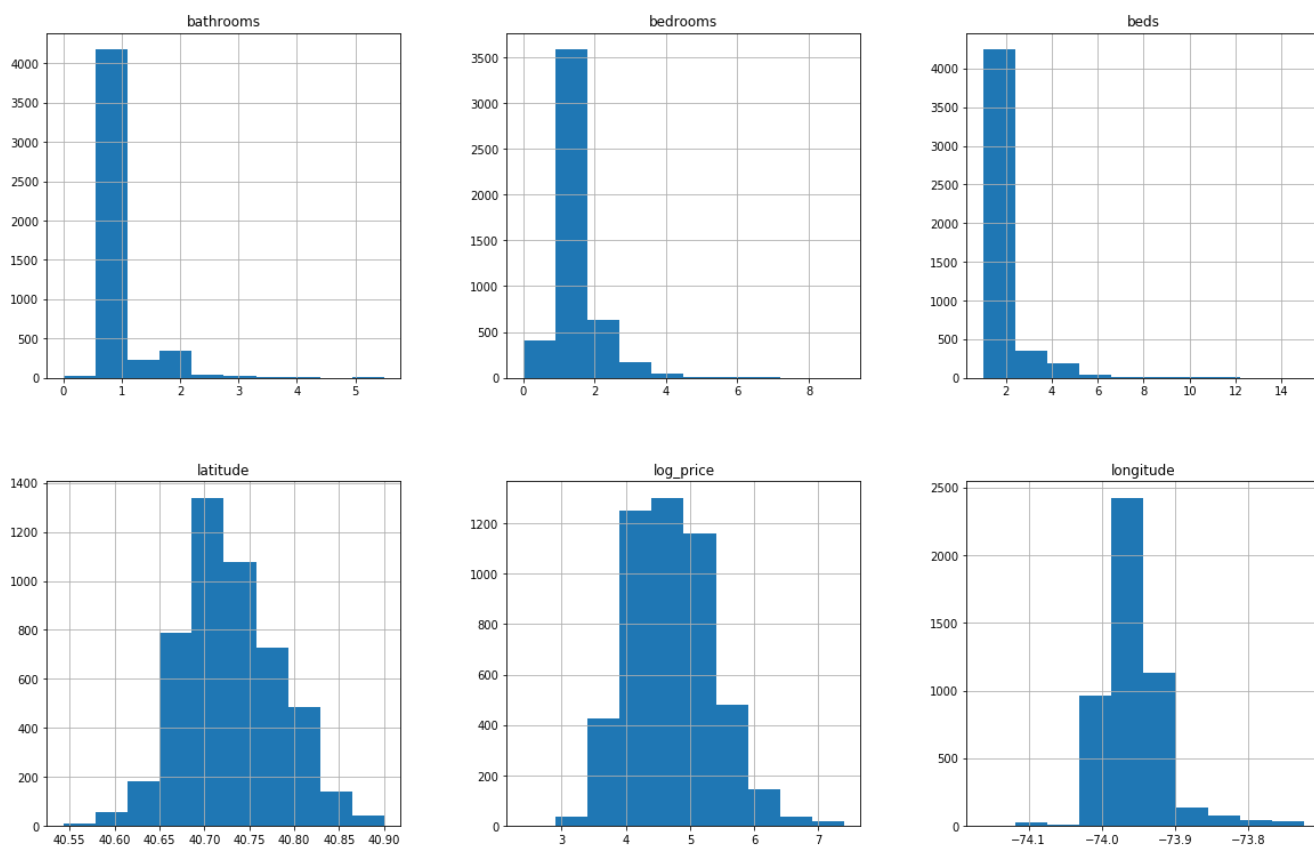





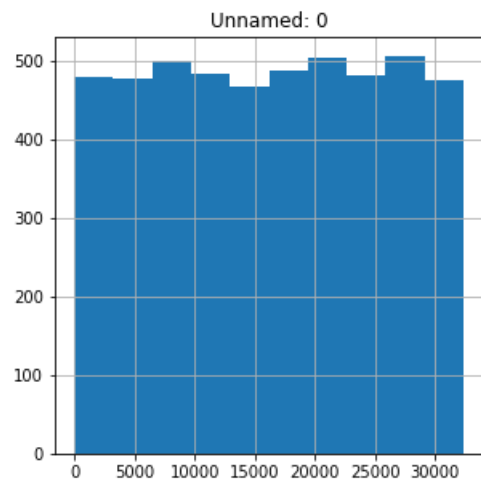
In [14]:

```
my_data[my_data.dtypes[(my_data.dtypes=="float")].index.values].hist(figsize=[20,20])
my_data[my_data.dtypes[(my_data.dtypes=="int64")].index.values].hist(figsize=[11,11])
```

Out[14]:

```
array([[<matplotlib.axes._subplots.AxesSubplot object at 0x00000294AF9ECCC8>,
        <matplotlib.axes._subplots.AxesSubplot object at 0x00000294AFA5C208>],
       [<matplotlib.axes._subplots.AxesSubplot object at 0x00000294AFA94188>,
        <matplotlib.axes._subplots.AxesSubplot object at 0x00000294AFACD288>]],
      dtype=object)
```





In [15]:

```
my_data.describe().transpose()
```

Out[15]:


	count	mean	std	min	25%	50%	75%	max
Unnamed: 0	4852.0	1.623937e+04	9.325224e+03	4.000000	8.123500e+03	1.633650e+04	2.442150e+04	3.234800e+04
id	4852.0	1.133776e+07	6.080574e+06	3152.000000	6.213140e+06	1.236683e+07	1.649103e+07	2.117444e+07
log_price	4852.0	4.708381e+00	6.594890e-01	2.397895	4.248495e+00	4.649176e+00	5.164786e+00	7.408531e+00
accommodates	4852.0	2.778648e+00	1.800746e+00	1.000000	2.000000e+00	2.000000e+00	4.000000e+00	1.600000e+01
bathrooms	4838.0	1.122261e+00	3.699294e-01	0.000000	1.000000e+00	1.000000e+00	1.000000e+00	5.500000e+00
latitude	4852.0	4.072880e+01	5.336419e-02	40.542680	4.069014e+01	4.072289e+01	4.076384e+01	4.090080e+01
longitude	4852.0	-7.395477e+01	4.202628e-02	-74.162537	-7.398367e+01	-7.395690e+01	-7.393956e+01	-7.372349e+01
number_of_reviews	4852.0	1.793157e+01	3.270683e+01	0.000000	1.000000e+00	5.000000e+00	2.000000e+01	4.740000e+02



review	scores	rating	count	mean	std	min	25%	50%	75%	max
			3759.0	9.349322e+01	8.581227e+00	20.000000	9.100000e+01	9.600000e+01	1.000000e+02	1.000000e+02
bedrooms			4850.0	1.157938e+00	6.998541e-01	0.000000	1.000000e+00	1.000000e+00	1.000000e+00	9.000000e+00
beds			4840.0	1.537397e+00	1.007436e+00	1.000000	1.000000e+00	1.000000e+00	2.000000e+00	1.500000e+01

In [16]:

```
plt.figure(figsize=(15,10))
sns.scatterplot(my_data.longitude,my_data.latitude,hue=my_data.neighbourhood)
plt.ioff()
```



- Sunnyside
- Mount Eden
- South Ozone Park
- Forest Hills
- Battery Park City
- Cobble Hill
- Inwood
- Downtown Brooklyn
- Woodside
- Brooklyn Heights
- Rego Park
- Riverdale
- Glendale
- St. George
- Bensonhurst
- Greenwich Village
- Flatlands
- Midland Beach
- East Elmhurst
- Borough Park
- Concourse Village
- Randall Manor
- Concord
- Highbridge
- Stapleton
- Midwood
- Ditmars / Steinway
- Brownsville
- Bay Ridge
- Corona
- Dyker Heights
- Kingsbridge Heights
- Eltingville
- Park Versailles
- Gravesend
- University Heights
- Norwood
- Co-op City
- Castle Hill
- Morningside Heights
- Longwood
- Melrose
- Jackson Heights
- DUMBO
- South Beach
- Claremont
- Roosevelt Island
- Elm Park
- Pelham Bay
- Brighton Beach
- Middle Village
- Parkchester
- Bergen Beach
- Kew Garden Hills
- Brooklyn Navy Yard
- Marble Hill
- Belmont
- Concourse
- Manhattan Beach
- Vinegar Hill
- Wakefield
- Fordham
- Morris Park
- Kingsbridge
- Throgs Neck
- Eastchester
- Whitestone
- Bedford Park
- Columbia Street Waterfront
- South Street Seaport
- Van Nest
- Brooklyn
- Edenwald
- Arrochar
- City Island
- Woodlawn
- Rosebank
- Sea Gate
- Utopia
- Queens
- Tompkinsville
- Manhattan
- Bronxdale

In [17]:

```
#imputing values with mean, median and mode
#mode is 1.0
my_data['bathrooms'].fillna(my_data['bathrooms'].mode()[0],inplace=True)
#mean is 93.50
my_data['review_scores_rating'].fillna(my_data['review_scores_rating'].mean(),inplace=True)
#mode is 1 bedroom
```

```
my_data['bedrooms'].fillna(my_data['bedrooms'].mode()[0],inplace=True)
#mode is 1 bed
my_data['beds'].fillna(my_data['beds'].mode()[0],inplace=True)
```

In [18]:

```
#Checking the number of null count
my_data.isnull().sum()
```

Out[18]:

```
Unnamed: 0          0
id                0
log_price         0
property_type     0
room_type        0
amenities        0
accommodates     0
bathrooms        0
bed_type         0
cancellation_policy 0
cleaning_fee     0
city             0
description       0
first_review     1019
host_has_profile_pic 38
host_identity_verified 38
host_response_rate 1504
host_since       38
instant_bookable  0
last_review     1015
latitude         0
longitude        0
name            0
neighbourhood    0
number_of_reviews 0
review_scores_rating 0
thumbnail_url    358
zipcode         63
bedrooms        0
beds            0
dtype: int64
```

In [19]:

```
#Removing the null values from is null
Remove_null=pd.DataFrame({"val":my_data['zipcode'].isnull()})
my_data=my_data[Remove_null['val']==False]
```

In [20]:

```
my_data=my_data.drop(['Unnamed: 0','amenities','bed_type','city','description','first_review',
'host_has_profile_pic','host_identity_verified','host_response_rate','host_since',
'last_review','name','neighbourhood','thumbnail_url','zipcode','id'],axis=1
)
```

In [21]:

```
my_data.isnull().sum()
```

Out[21]:

```
log_price          0
property_type      0
room_type          0
accommodates       0
bathrooms          0
cancellation_policy 0
cleaning_fee       0
instant_bookable   0
latitude           0
longitude          0
```

```
number_of_reviews    0
review_scores_rating  0
bedrooms             0
beds                 0
dtype: int64
```

In [22]:

```
my_data.describe()
```

Out[22]:

	log_price	accommodates	bathrooms	latitude	longitude	number_of_reviews	review_scores_rating	bedrooms
count	4789.000000	4789.000000	4789.000000	4789.000000	4789.000000	4789.000000	4789.000000	4789.000000
mean	4.705982	2.767592	1.121529	40.728855	-73.954704	17.957194	93.483237	1.156609
std	0.658785	1.785157	0.369942	0.053468	0.042065	32.749212	7.547036	0.697373
min	2.397895	1.000000	0.000000	40.542680	-74.162537	0.000000	20.000000	0.000000
25%	4.248495	2.000000	1.000000	40.690093	-73.983666	1.000000	93.000000	1.000000
50%	4.624973	2.000000	1.000000	40.722996	-73.956807	5.000000	93.493216	1.000000
75%	5.164786	4.000000	1.000000	40.763896	-73.939605	20.000000	98.000000	1.000000
max	7.408531	16.000000	5.500000	40.900803	-73.723488	474.000000	100.000000	9.000000

In [23]:

```
#categorical=['property_type','room_type','cancellation_policy','instant_bookable']
my_data=pd.concat((my_data,pd.get_dummies(my_data['property_type'])),axis=1)
my_data=pd.concat((my_data,pd.get_dummies(my_data['room_type'])),axis=1)
my_data=pd.concat((my_data,pd.get_dummies(my_data['cancellation_policy'])),axis=1)
my_data=pd.concat((my_data,pd.get_dummies(my_data['instant_bookable'])),axis=1)
my_data=pd.concat((my_data,pd.get_dummies(my_data['cleaning_fee'])),axis=1)
```

In [24]:

```
my_data=my_data.drop(['property_type','room_type','cancellation_policy','instant_bookable','cleaning_fee'],axis=1)
```

In [25]:

```
my_data.isnull().sum()
```

Out[25]:

```
log_price          0
accommodates       0
bathrooms          0
latitude           0
longitude           0
number_of_reviews  0
review_scores_rating 0
bedrooms           0
beds               0
Apartment          0
Condo              0
Hotel type 1       0
Hotel type 2       0
House              0
Housing            0
Other              0
Entire home/apt    0
Private room       0
Shared room        0
flexible           0
moderate           0
strict             0
Instant Booking    0
No Instant Booking 0
Cleaning Fee Req   0
```

```
No Cleaning Fee      0
dtype: int64
```

In [26]:

```
target = my_data['log_price']
target_df = pd.DataFrame(target)
target_df.head()
```

Out[26]:

	log_price
0	4.382027
1	5.075174
3	5.010635
4	4.317488
5	5.416100

In [27]:

```
features_df=my_data.drop(['log_price'],axis=1)
features_df.head()
```

Out[27]:

	accommodates	bathrooms	latitude	longitude	number_of_reviews	review_scores_rating	bedrooms	beds	Apartment	Condo	..
0	1	1.0	40.762239	73.981589	1	93.493216	1.0	1.0	0	1	..
1	3	1.0	40.677892	73.992054	17	98.000000	1.0	2.0	1	0	..
3	4	1.0	40.800331	73.965090	49	96.000000	1.0	2.0	1	0	..
4	2	1.0	40.711386	73.963529	0	93.493216	1.0	1.0	1	0	..
5	5	1.0	40.726874	73.979947	0	93.493216	2.0	2.0	1	0	..

5 rows × 25 columns



In [28]:

```
from sklearn.model_selection import train_test_split
X_train_old, X_test_old, y_train, y_test = train_test_split(features_df,target_df, test_size=0.25,
random_state = 0)
```

In [29]:

```
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
X_train = scaler.fit_transform(X_train_old)
X_test = scaler.transform(X_test_old)
```

In [30]:

```
X_train = pd.DataFrame(X_train, columns = X_train_old.columns)
print('Train dataset dimensionality:', X_train.shape)
print('Train dataset dimensionality:', y_train.shape)
```

```
Train dataset dimensionality: (3591, 25)
Train dataset dimensionality: (3591, 1)
```

In [31]:

```
X_test = pd.DataFrame(X_test, columns = X_test_old.columns)
print('Test dataset dimensionality:', X_test.shape)
print('Train dataset dimensionality:', y_test.shape)
```

Test dataset dimensionality: (1198, 25)

Train dataset dimensionality: (1198, 1)

## Bagging

**Bootstrapping** - resample method that repeatedly drawn sample from smaller data to form smaller data set.

**Bagging can be defined as Bootstrapping + Aggrgation and it is an ensemble method in which we first bootstrap our sample data and train them . After that , we aggregate them with equi weights**

## Model 1 Ridge Regressor with Bagging

In [32]:

```
# Using grid search to find the best parameter for bagging
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import BaggingRegressor
parameters_grid = {'n_estimators': [50,100,200,500],
                   'max_samples': [50,100,200,400,500]}

_grid_search_ = GridSearchCV(BaggingRegressor(), parameters_grid, cv=5, return_train_score=True)
_grid_search_.fit(X_train, y_train)
print("Best parameters: {}".format(_grid_search_.best_params_))
print("Best cross-validation score: {:.2f}".format(_grid_search_.best_score_))
```

Best parameters: {'max\_samples': 500, 'n\_estimators': 500}

Best cross-validation score: 0.67

In [33]:

```
# Using grid Search to find best parameter for model Ridge Regressor
from sklearn.linear_model import Ridge
parameters_grid = {'alpha':[0.001, 0.01, 0.1, 1, 10, 100, 1000]}
_grid_search_ = GridSearchCV(Ridge(), parameters_grid, cv=5, return_train_score=True)
_grid_search_.fit(X_train, y_train)
print("Best parameters: {}".format(_grid_search_.best_params_))
print("Best cross-validation score: {:.2f}".format(_grid_search_.best_score_))
```

Best parameters: {'alpha': 0.1}

Best cross-validation score: 0.58

In [34]:

```
from sklearn.ensemble import BaggingRegressor
from sklearn.linear_model import Ridge

ridge = Ridge(alpha=0.1)
bag_ridge_reg = BaggingRegressor(ridge, n_estimators=500, max_samples=500, bootstrap=True, n_jobs=-1, random_state=0)

bag_ridge_reg.fit(X_train, y_train)
y_pred = bag_ridge_reg.predict(X_test)

print('Score after applying Bagging on Ridge Regressor on Train data set: {:.2f}'.format(bag_ridge_reg.score(X_train, y_train)))
print('Score after applying Bagging on Ridge Regressor on Test data set: {:.2f}'.format(bag_ridge_reg.score(X_test, y_test)))
```

```
print('Score after applying Bagging on Ridge Regressor on Test data set: {:.2f}'.format(bag_ridge_reg.score(X_test, y_test)))
```

Score after applying Bagging on Ridge Regressor on Train data set: 0.59  
Score after applying Bagging on Ridge Regressor on Test data set: 0.56

## Model 2 Decision Tree with Bagging

In [35]:

```
# Grid search to find the best parameter for the model
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import GridSearchCV
parameters_grid = {'max_depth': np.arange(1, 10)}
_grid_search_ = GridSearchCV(DecisionTreeRegressor(random_state=0), parameters_grid, cv=5, return_train_score=True)
_grid_search_.fit(X_train, y_train)
print("Best parameters: {}".format(_grid_search_.best_params_))
print("Best cross-validation score: {:.2f}".format(_grid_search_.best_score_))
```

Best parameters: {'max\_depth': 6}  
Best cross-validation score: 0.61

In [36]:

```
# Grid search to find the best parameter for bagging
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import BaggingRegressor
parameters_grid = {'n_estimators': [50,100,200,500],
                  'max_samples': [50,100,200,400,500]}

_grid_search_ = GridSearchCV(BaggingRegressor(), parameters_grid, cv=5, return_train_score=True)
_grid_search_.fit(X_train, y_train)
print("Best parameters: {}".format(_grid_search_.best_params_))
print("Best cross-validation score: {:.2f}".format(_grid_search_.best_score_))
```

Best parameters: {'max\_samples': 500, 'n\_estimators': 500}  
Best cross-validation score: 0.67

In [37]:

```
# building the bagging model for Decision Tree Regressor using the best parameters
from sklearn.ensemble import BaggingRegressor
from sklearn.tree import DecisionTreeRegressor

dt_reg = DecisionTreeRegressor(max_depth = 6, random_state=0)
bag_dectree_reg = BaggingRegressor(dt_reg, n_estimators=500, max_samples=500, bootstrap=True, n_jobs=-1, random_state=0)

bag_dectree_reg.fit(X_train, y_train)
y_pred = bag_dectree_reg.predict(X_test)

print('Score after applying Bagging on Decision Tree Regressor on Train data Set: {:.2f}'.format(bag_dectree_reg.score(X_train, y_train)))
print('Score after applying Bagging on Decision Tree Regressor on Test data Set: {:.2f}'.format(bag_dectree_reg.score(X_test, y_test)))
```

Score after applying Bagging on Decision Tree Regressor on Train data Set: 0.70  
Score after applying Bagging on Decision Tree Regressor on Test data Set: 0.65

## Pasting

**In pasting sampling is done without replacement. Furthermore, bootstrap is set to false in pasting.**

## Model 1- Lasso Regressor with Pasting

In [38]:

```
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import BaggingRegressor
_param_grid = {'n_estimators': [50,100,200,500],
               'max_samples': [50,100,200,400,500]}

__grid__search = GridSearchCV(BaggingRegressor(), _param_grid, cv=5, return_train_score=True)
__grid__search.fit(X_train, y_train)
print("Best parameters: {}".format(__grid__search.best_params_))
print("Best cross-validation score: {:.3f}".format(__grid__search.best_score_))
```

Best parameters: {'max\_samples': 500, 'n\_estimators': 500}  
Best cross-validation score: 0.670

In [39]:

```
from sklearn.linear_model import Lasso
lasso_param_ = {'alpha':[0.001, 0.01, 0.1, 1, 10, 100, 1000]}
__grid__search = GridSearchCV(Lasso(), lasso_param_, cv=5, return_train_score=True)
__grid__search.fit(X_train, y_train)
print("Best parameters: {}".format(__grid__search.best_params_))
print("Best cross-validation score: {:.3f}".format(__grid__search.best_score_))
```

Best parameters: {'alpha': 0.001}  
Best cross-validation score: 0.581

In [40]:

```
lasso = Lasso(alpha=0.01)
pas_lasso_reg = BaggingRegressor(lasso, n_estimators=500, max_samples=500, bootstrap=False, n_jobs=-1, random_state=0)

pas_lasso_reg.fit(X_train, y_train)
y_pred = pas_lasso_reg.predict(X_test)

print('Score after aplying Pasting on Lasso Regressor on Train Set: {:.2f}'.format(pas_lasso_reg.score(X_train, y_train)))
print('Score after pasting Pasting on Lasso Regressor on Test Set: {:.2f}'.format(pas_lasso_reg.score(X_test, y_test)))
```

Score after aplying Pasting on Lasso Regressor on Train Set: 0.52  
Score after pasting Pasting on Lasso Regressor on Test Set: 0.49

## Model 2 - Knn Regressor with Pasting

In [41]:

```
from sklearn.neighbors import KNeighborsRegressor
from sklearn.model_selection import GridSearchCV
knn_param_grid = {'n_neighbors' : range(1,20), 'p': [1,2], 'weights': ['distance','uniform']}

grid_knn_rgr = GridSearchCV(KNeighborsRegressor(), param_grid = knn_param_grid, cv=10, return_train_score=True, n_jobs=-1)
grid_knn_rgr.fit(X_train, y_train)

print("Best parameters: {}".format(grid_knn_rgr.best_params_))
print("Best cross-validation score: {:.2f}".format(grid_knn_rgr.best_score_))
```

Best parameters: {'n\_neighbors': 16, 'p': 1, 'weights': 'distance'}  
Best cross-validation score: 0.59

In [42]:

```
from sklearn.ensemble import BaggingRegressor
```



```

from sklearn.neighbors import KNeighborsRegressor

bag_reg_knn2 = BaggingRegressor(KNeighborsRegressor(5, p=1, weights= 'distance'),max_features= 9, m
ax_samples=500, n_estimators= 200, random_state=0, bootstrap = False)
bag_reg_knn2.fit(X_train, y_train)
y_pred_knn2 = bag_reg_knn2.predict(X_test)

print('Train score after applying pasting in KNN Regressor:
{:.2f}%'.format(bag_reg_knn2.score(X_train, y_train)*100))
print('Test score after applying pasting in KNN Regresso:
{:.2f}%'.format(bag_reg_knn2.score(X_test, y_test)*100))

print()
print('MAE:', metrics.mean_absolute_error(y_test, y_pred_knn2))
print('MSE:', metrics.mean_squared_error(y_test, y_pred_knn2))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_pred_knn2)))
print("r2_Score:", r2_score(y_test, y_pred_knn2))

```

Train score after applying pasting in KNN Regressor: 66.03%  
Test score after applying pasting in KNN Regresso: 55.84%

MAE: 0.32189785453260594  
MSE: 0.19054611432352225  
RMSE: 0.43651588095225385  
r2\_Score: 0.5584223204113332

## Adaboosting

**Adaboost try to fit a sequence of weak learner on repeatedly modified data set**

## Model 1 - KNN Regressor with Adaboost

In [43]:

```

# Grid search to find the best adaboost
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import AdaBoostRegressor
parameters_grid = {'n_estimators': [50,100,200],
                   'learning_rate': [.01, .05, .1, 1]}

__grid_search__ = GridSearchCV(AdaBoostRegressor(random_state = 0), parameters_grid, cv=5, return_t
rain_score=True)
__grid_search__.fit(X_train, y_train)
print("Best parameters: {}".format(__grid_search__.best_params_))
print("Best cross-validation score: {:.2f}".format(__grid_search__.best_score_))

```

Best parameters: {'learning\_rate': 0.1, 'n\_estimators': 50}  
Best cross-validation score: 0.58

In [44]:

```

from sklearn.neighbors import KNeighborsRegressor
knn_reg = KNeighborsRegressor(n_neighbors=3)
knn_reg.fit(X_train,y_train)
y_pred = knn_reg.predict(X_test)
print('KNN regressor score on Train Set score: {:.2f}'.format(knn_reg.score(X_train, y_train)))
print('KNN regressor score on Test Set score: {:.2f}'.format(knn_reg.score(X_test, y_test)))

```

KNN regressor score on Train Set score: 0.77  
KNN regressor score on Test Set score: 0.50

In [45]:

```

from sklearn.ensemble import AdaBoostRegressor
knn_ada_reg = AdaBoostRegressor(KNeighborsRegressor(n_neighbors=3), n_estimators=100, learning_rat

```

```
e=0.05, random_state=0)
knn_ada_reg.fit(X_train, y_train)
y_pred = knn_ada_reg.predict(X_test)
print('KNN regressor score on Train Set after applying Adaboost Boosting:
{:.3f}'.format(knn_ada_reg.score(X_train, y_train)))
print('KNN regressor score on Test Set after applying Adaboost Boosting:
{:.3f}'.format(knn_ada_reg.score(X_test, y_test)))
```

KNN regressor score on Train Set after applying Adaboost Boosting: 0.869  
KNN regressor score on Test Set after applying Adaboost Boosting: 0.488

## Model 2 - Decision Tree Regressor

In [46]:

```
# Grid search to find the best adaboost
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import AdaBoostRegressor
_parameter_grid = {'n_estimators': [50,100,200,500],
                  'learning_rate': [0.01, .05, .1, 1]}

grid__search = GridSearchCV(AdaBoostRegressor(random_state = 0), _parameter_grid, cv=5, return_train_score=True)
grid__search.fit(X_train, y_train)
print("Best parameters: {}".format(grid__search.best_params_))
print("Best cross-validation score: {:.2f}".format(grid__search.best_score_))
```

Best parameters: {'learning\_rate': 0.01, 'n\_estimators': 500}  
Best cross-validation score: 0.58

In [47]:

```
from sklearn.tree import DecisionTreeRegressor
detr_reg = DecisionTreeRegressor(max_depth = 6, random_state=0)
detr_reg.fit(X_train, y_train)
y_pred = detr_reg.predict(X_test)

print('Decision tree regressor score on Train Set score: {:.2f}'.format(detr_reg.score(X_train, y_train)))
print('Decision tree regressor score on Test Set score: {:.2f}'.format(detr_reg.score(X_test, y_test)))
```

Decision tree regressor score on Train Set score: 0.70  
Decision tree regressor score on Test Set score: 0.60

In [48]:

```
from sklearn.ensemble import AdaBoostRegressor
ada_regr = AdaBoostRegressor(DecisionTreeRegressor(max_depth=6), n_estimators=100, learning_rate=0.05, random_state=0)
ada_regr.fit(X_train, y_train)

print('Decision tree regressor score on Train Set after Adaboost Boosting: {:.2f}'.format(ada_regr.score(X_train, y_train)))
print('Decision tree regressor score on Test Set after Adaboost Boosting: {:.2f}'.format(ada_regr.score(X_test, y_test)))
```

Decision tree regressor score on Train Set after Adaboost Boosting: 0.73  
Decision tree regressor score on Test Set after Adaboost Boosting: 0.64

## Gradient Boosting

In [49]:

```
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import GridSearchCV
gbrt = GradientBoostingRegressor(random_state=0)
```

```
param_grid={'n_estimators':[50,100,150], 'learning_rate':[0.5,1], 'max_depth': np.arange(1, 6)}
__grid_search__=GridSearchCV(gbrt,param_grid,cv=5,return_train_score=True)
__grid_search__.fit(X_train, y_train)
print("Best Parameters: {}".format(__grid_search__.best_params_))
print("Best cross-validation score: {:.2f}".format(__grid_search__.best_score_))
```

Best Parameters: {'learning\_rate': 0.5, 'max\_depth': 2, 'n\_estimators': 50}  
 Best cross-validation score: 0.66

In [50]:

```
# building the model with best parameters
gbrt = GradientBoostingRegressor(random_state=0, learning_rate=0.5, max_depth = 2, n_estimators=50)
gbrt.fit(X_train, y_train)
y_pred = gbrt.predict(X_test)
print("Score after applying Gradient Boosting on Train Set: {:.3f}".format(gbrt.score(X_train,
y_train)))
print("Score after applying Gradient Boosting on Test Set: {:.3f}".format(gbrt.score(X_test,
y_test)))
```

Score after applying Gradient Boosting on Train Set: 0.730  
 Score after applying Gradient Boosting on Test Set: 0.655

## Principal Component Analysis

**PCA technique is use to reduce the dimensionality of a data set consisting of many variables correlated to each other.**

In [51]:

```
from sklearn.decomposition import PCA

pca = PCA(n_components=0.95,random_state = 0)
pca.fit(X_train)
X_train_reduced = pca.transform(X_train)
X_test_reduced = pca.transform(X_test)
```

### Checking after PCA how many column get reduced

In [52]:

```
print('X_train shape',X_train.shape)
print('X_train_reduced shape',X_train_reduced.shape)
```

X\_train shape (3591, 25)  
 X\_train\_reduced shape (3591, 8)

## Scaling the data

In [53]:

```
mm = MinMaxScaler()
X_train_pca= mm.fit_transform(X_train_reduced)
X_test_pca = mm.transform(X_test_reduced)
```

## Dummy list will use if required.

In [54]:

```
train_score_pca=[]
test_score_pca=[]
```

```
models_pca =[]
```

## Model -1 KNN regressor after PCA Technique

In [55]:

```
np.random.seed(0)

x_range_1 = range(1,30,1)
tuned_parameters=dict(n_neighbors=x_range_1)

#Grid model
knn_reg_pca = KNeighborsRegressor()
grid_knn_pca=GridSearchCV(knn_reg_pca,tuned_parameters,cv=5,return_train_score=True)
grid_model_knn_pca=grid_knn_pca.fit(X_train_pca,y_train)

print(grid_model_knn_pca.best_params_)
print('validation score: {:.2f}'.format( grid_model_knn_pca.best_score_))
```

```
{'n_neighbors': 7}
validation score: 0.53
```

In [56]:

```
#General model
knn_pca=KNeighborsRegressor(n_neighbors=9)
knn_model_pca=knn_pca.fit(X_train_pca,y_train)
print('Train score: {}'.format(knn_model_pca.score(X_train_pca,y_train)))
print('Test score: {}'.format(knn_model_pca.score(X_test_pca,y_test)))
train_score_pca.append(knn_model_pca.score(X_train_pca,y_train))
test_score_pca.append(knn_model_pca.score(X_test_pca,y_test))
```

```
Train score: 0.6375807251169314
Test score: 0.5246936528801134
```

In [57]:

```
#calculating the accuracies
knn_accuracies_pca = cross_val_score(estimator = knn_model_pca, X = X_train_pca, y = y_train, cv =
10)
print("Accuracy: {:.2f} %".format(knn_accuracies_pca.mean()*100))
```

```
Accuracy: 53.17 %
```


In [58]:

```
#visualizing the train and test accuracy score for KNN Regressor
import matplotlib.pyplot as plt
views = pd.DataFrame(grid_knn_pca.cv_results_)
plt.plot(views['param_n_neighbors'],views['mean_test_score'],marker='o',c='r',label='Validation
Test score')
plt.plot(views['param_n_neighbors'],views['mean_train_score'],marker='*',c='g',label='Validation Tr
ain score')
plt.title('Number of Neighbors Vs. Mean Train/Validation Accuracy')
plt.xlabel('Number of Neighbors in KNN regressor')
plt.ylabel('Accuracy')
plt.legend()
```

Out[58]:

```
<matplotlib.legend.Legend at 0x294b1c71188>
```





In [59]:

```
#train and test accuracy score for KNN Regressor after running PCA
knn = KNeighborsRegressor(n_neighbors=7)
knn.fit(X_train_pca, y_train)
print('Train score on best parameters for KNN Regressor
{}'.format(knn.score(X_train_pca,y_train)))
print('Test score on best parameters for KNN Regressor {}'.format(knn.score(X_test_pca,y_test)))
```

Train score on best parameters for KNN Regressor 0.6570193928545962  
 Test score on best parameters for KNN Regressor 0.5201252133003174

## Model 2 -Linear Regression after PCA

In [60]:

```
#train and test accuracy score for linear regression after running PCA
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
lreg = LinearRegression()
#scaled train test
lreg.fit(X_train_pca, y_train)
print('Training score for regression model: {}'.format(lreg.score(X_train_pca, y_train)))
print('Testing score for regression model: {}'.format(lreg.score(X_test_pca, y_test)))
print('R-squared score (training): {:.2f}\n'.format(lreg.score(X_train_pca, y_train)))
print('R-squared score (test): {:.2f}'.format(lreg.score(X_test_pca, y_test)))
train_score_pca.append(lreg.score(X_train_pca, y_train))
test_score_pca.append(lreg.score(X_test_pca, y_test))
```

Training score for regression model: 0.4594824287389822  
 Testing score for regression model: 0.43657755698958867  
 R-squared score (training): 0.46

R-squared score (test): 0.44

In [61]:

```
#calculating the accuracies
lnr_accuracies_pca = cross_val_score(estimator = lreg, X = X_train_pca, y = y_train, cv = 10)
print("Accuracy: {:.2f} %".format(lnr_accuracies_pca.mean()*100))
```

Accuracy: 45.56 %

## Model 3 - Ridge regressor after PCA

In [62]:

```
np.random.seed(0)
x_range_2 = [0.01, 0.1, 1, 10, 100]
tuned_parameters = [{'alpha':x_range_2}]

#Grid model
ridge_pca = Ridge(max_iter=1000,tol=0.1,random_state=0)
grid_ridge_pca=GridSearchCV(ridge_pca,tuned_parameters,cv=5, return_train_score= True, iid = False)
grid_model_ridge_pca=grid_ridge_pca.fit(X_train_pca,y_train)

print("Best parameters: {}".format(grid_model_ridge_pca.best_params_))
```

```
print('Best parameters: {}'.format(grid_model_ridge_pca.best_params_))
print('Cross validation score: {:.2f}'.format(grid_model_ridge_pca.best_score_))
```

Best parameters: {'alpha': 1}  
Cross validation score: 0.45

In [63]:

```
#General model
ridge_l_pca=Ridge(alpha=0.1)
ridge_model_pca=ridge_l_pca.fit(X_train_pca,y_train)
print('Training score for Ridge regression model:
{}'.format(ridge_model_pca.score(X_train_pca,y_train)))
print('Testing score for Ridge regression model:
{}'.format(ridge_model_pca.score(X_test_pca,y_test)))
train_score_pca.append(ridge_model_pca.score(X_train_pca,y_train))
test_score_pca.append(ridge_model_pca.score(X_test_pca,y_test))
```

Training score for Ridge regression model: 0.45948226719811064  
Testing score for Ridge regression model: 0.43659616180333627

In [64]:

```
#calculating the accuracies
ridr_accuracies_pca = cross_val_score(estimator = ridge_model_pca, X = X_train_pca, y = y_train, cv
= 10)
print("Accuracy: {:.2f} %".format(ridr_accuracies_pca.mean()*100))
```


Accuracy: 45.56 %

In [65]:

```
#visualizing the train and test accuracy score for Ridge regression
import matplotlib.pyplot as plt
vector = pd.DataFrame(grid_model_ridge_pca.cv_results_)
plt.plot(vector['param_alpha'],vector['mean_test_score'],marker='o',c='r',label='Validation Test score')
plt.plot(vector['param_alpha'],vector['mean_train_score'],marker='*',c='g',label='Validation Train score')
plt.title('Alpha Vs. Mean Train/Validation Accuracy')
plt.xlabel(r'$\alpha$')
plt.ylabel('Accuracy')
plt.legend()
```

Out[65]:

<matplotlib.legend.Legend at 0x294b1ce8908>



## Model 4 Lasso Regression after PCA

In [66]:

```
#list the best parameter value for lasso regressor
from sklearn.linear_model import Lasso
np.random.seed(0)
x_range_3 = [0.01, 0.1, 1, 10, 100]
tuned_parameters = [{'alpha':x_range_3}]

#Grid model
lasso_pca = Lasso(max_iter=1000,tol=0.1,random_state=0)
grid_lasso_pca=GridSearchCV(lasso_pca,tuned_parameters,cv=5, return_train_score= True, iid = False)
grid_model_lasso_pca=grid_lasso_pca.fit(X_train_pca,y_train)

print("Best parameters: {}".format(grid_model_lasso_pca.best_params_))
print('Best Crossvalidation score: {:.2f}'.format( grid_model_lasso_pca.best_score_))
```

Best parameters: {'alpha': 0.01}  
Best Crossvalidation score: 0.42

In [67]:

```
#General model
lasso_1_pca=Lasso(alpha=0.01, tol=0.1)
lasso_model_pca=lasso_1_pca.fit(X_train_pca,y_train)
print('Train score on best parameters for Lasso Regressor:
{}'.format(lasso_model_pca.score(X_train_pca,y_train)))
print('Test score on best parameters for Lasso Regressor:
{}'.format(lasso_model_pca.score(X_test_pca,y_test)))
train_score_pca.append(lasso_model_pca.score(X_train_pca,y_train))
test_score_pca.append(lasso_model_pca.score(X_test_pca,y_test))
```

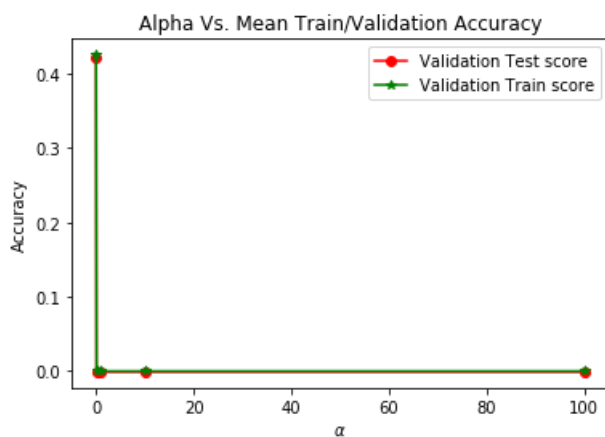
Train score on best parameters for Lasso Regressor: 0.42537317346395176  
Test score on best parameters for Lasso Regressor: 0.40669138836971286

In [68]:

```
#visualizing the train and test accuracy score for lasso regression
import matplotlib.pyplot as plt
vis_results = pd.DataFrame(grid_lasso_pca.cv_results_)
plt.plot(vis_results['param_alpha'],vis_results['mean_test_score'],marker='o',c='r',label='Validati
on Test score')
plt.plot(vis_results['param_alpha'],vis_results['mean_train_score'],marker='*',c='g',label='Validat
ion Train score')
plt.title('Alpha Vs. Mean Train/Validation Accuracy')
plt.xlabel(r'$\alpha$')
plt.ylabel('Accuracy')
plt.legend()
```

Out[68]:

<matplotlib.legend.Legend at 0x294b1d89f48>



## Model 5 -Polynomial Regerssion after PCA

In [69]:

```
from sklearn.preprocessing import PolynomialFeatures
train_score_list = []
test_score_list = []
regressor = LinearRegression()
for n in range(1,4):
    poly = PolynomialFeatures(n)
    X_train_poly_pca = poly.fit_transform(X_train_pca)
    X_test_poly_pca = poly.transform(X_test_pca)
    regressor.fit(X_train_poly_pca, y_train)
    train_score_list.append(regressor.score(X_train_poly_pca, y_train))
    test_score_list.append(regressor.score(X_test_poly_pca, y_test))

train = [sum(train_score_list)/len(train_score_list)]
test = [sum(test_score_list)/len(test_score_list)]
print(train)
print(test)
```

```
[0.5509409687032404]
[0.5029318211353652]
```


## Plotting Accuracy graph for Polynomial Regression after

In [70]:

```
x_axis = range(1,4)
plt.plot(x_axis, train_score_list, c = 'g', label = 'Train Score')
plt.plot(x_axis, test_score_list, c = 'b', label = 'Test Score')
plt.xlabel('degree')
plt.ylabel('accuracy')
plt.legend()
```

Out[70]:

<matplotlib.legend.Legend at 0x294b1d89a08>



## Model 5 - Support Vector Machine after PCA

In [71]:

```
from sklearn.svm import SVR
from sklearn.svm import LinearSVR
_paramtr_grid = {'C': [0.001, 0.01, 0.1, 1, 10, 100],
                  'gamma': [0.001, 0.01, 0.1, 1, 10, 100]}

grid_search = GridSearchCV(SVR(), _paramtr_grid, cv=5, return_train_score=True)
grid_search.fit(X_train_pca, y_train)
print("Best parameters: {}".format(grid_search.best_params_))
print("Best cross-validation score: {:.2f}".format(grid_search.best_score_))
```

```
Best parameters: {'C': 10, 'gamma': 100}
Best cross-validation score: 0.55
```



## Linear SVM after PCA

In [72]:

```
ln_svr = LinearSVR(C=10).fit(X_train_pca, y_train)
print('Train score on best parameters for LinearSVR - {}'.format(ln_svr.score(X_train_pca, y_train)))
print('Test score on best parameters for LinearSVR - {}'.format(ln_svr.score(X_test_pca, y_test)))
```

Train score on best parameters for LinearSVR - 0.4565816063539583  
Test score on best parameters for LinearSVR - 0.43280524491002903

In [73]:

```
#calculating the accuracy
ln_svr_accuracy = cross_val_score(estimator = ln_svr, X = X_train_pca, y = y_train, cv = 5)
print("Accuracy: {:.2f} %".format(ln_svr_accuracy.mean()*100))
```

Accuracy: 45.17 %

## Kernel(Linear)SVM after PCA

In [74]:

```
kl_svm = SVR(kernel='linear', C=10).fit(X_train_pca, y_train)
print('Train score on best parameters for SVR kernel - Linear {}'.format(kl_svm.score(X_train_pca, y_train)))
print('Test score on best parameters for SVR kernel - Linear {}'.format(kl_svm.score(X_test_pca, y_test)))
```

Train score on best parameters for SVR kernel - Linear 0.45574929544318354  
Test score on best parameters for SVR kernel - Linear 0.42599964884893393

In [75]:

```
#calculating the accuracy
kl_svm_accuracy = cross_val_score(estimator = kl_svm, X = X_train_pca, y = y_train, cv = 5)
print("Accuracy: {:.2f} %".format(kl_svm_accuracy.mean()*100))
```

Accuracy: 45.20 %

## Kernel(RBF) SVM after PCA

In [76]:

```
krbf_svm = SVR(kernel='rbf', gamma=100, C=10).fit(X_train_pca, y_train)
print('Train score on best parameters for SVR kernel - rbf {}'.format(krbf_svm.score(X_train_pca, y_train)))
print('Test score on best parameters for SVR kernel - rbf {}'.format(krbf_svm.score(X_test_pca, y_test)))
```

Train score on best parameters for SVR kernel - rbf 0.6782168760741613  
Test score on best parameters for SVR kernel - rbf 0.5401128539537601

In [77]:

```
#calculating the accuracy
krbf_svm_accuracy = cross_val_score(estimator = krbf_svm, X = X_train_pca, y = y_train, cv = 5)
print("Accuracy: {:.2f} %".format(krbf_svm_accuracy.mean()*100))
```

Accuracy: 55.19 %

Accuracy: 43.41 %

## Kernel SVM (Poly) after PCA

In [78]:

```
klp_svmm = SVR(kernel='poly', degree=3, C=10).fit(X_train_pca, y_train)
print('Train score on best parameters for SVR kernel - poly {}'.format(klp_svmm.score(X_train_pca,
y_train)))
print('Test score on best parameters for SVR kernel - poly {}'.format(klp_svmm.score(X_test_pca,y_
test)))
```

Train score on best parameters for SVR kernel - poly 0.4498346524343754  
Test score on best parameters for SVR kernel - poly 0.42378354017324316

In [79]:


```
#calculating the accuracy of Kernal SVM(Poly)
klp_svmm_accuracy = cross_val_score(estimator = klp_svmm, X = X_train_pca, y = y_train, cv = 5)
print("Accuracy: {:.2f} %".format(klp_svmm_accuracy.mean()*100))
```

Accuracy: 43.41 %

## Visualizing the train and test score of kernalized SVM regression models after applying PCA

In [80]:

```
fig, ax = plt.subplots(figsize=(8,4))
width = 0.3
plt.xlabel('SVR')
plt.ylabel('Accuracy')
outer_labels = ['Training LinearSVR','Test LinearSVR','Training SVR kernel - Linear','Test SVR ker
nel - Linear','Training SVR kernel - rbf',
                'Test SVR kernel - rbf','Training SVR kernel - poly','Test SVR kernel - poly']
inner_label = ['Training LinearSVR','Test LinearSVR','Training SVR kernel - Linear','Test SVR kern
el - Linear','Training SVR kernel - rbf',
               'Test SVR kernel - rbf','Training SVR kernel - poly','Test SVR kernel - poly']
list_1 = [0,.3,1,1.3,2,2.3,3,3.3]
ax.set_xticks(list_1)
for j in range(0,4,1) :
    ax.set_xticklabels(outer_labels,rotation=90)
    ax.set_xticklabels(inner_label,rotation=90)
train_accuracylist=[45.38,45.57,71.57,49.05]
test_accuracylist=[42.42,42.59,55.01,44.76]
for i in range(0,4,1) :
    ax.bar(i,train_accuracylist[i],width)
    ax.bar(i+width,test_accuracylist[i],width)
plt.grid()
```





## Compairing Model before applying PCA technique and after applying PCA technique

In [81]:

```
#Index Levels
outer_row= ['Before PCA','Before PCA','After PCA','After PCA']
inner_row = ['Training Accuracy','Test Accuracy','Training Accuracy','Test Accuracy']
grn_level = list(zip(outer_row,inner_row))
grn_level = pd.MultiIndex.from_tuples(grn_level)
```

In [82]:

```
#train and test accuracy score of the models as observed before and after PCA
data_subset = np.array([(0.5935,0.5900,0.6514,0.5935,0.5894,0.3634,0.8042,0.5825,0.5612),
                        (0.5582,0.5581,0.5782,0.5585,0.5567,0.3875,0.5161,0.5425,0.5212),
                        (0.4594,0.5509,0.6375,0.4594,0.4253,0.4481,0.6782,0.4557,0.4498),
                        (0.4365,0.5029,0.5246,0.4365,0.4066,0.4144,0.5401,0.4259,0.4237)])
```

In [83]:

```
data_df = pd.DataFrame(data_subset,index=grn_level,columns=['Linear Regression','Polynomial
Regressor','KNN Regressor',
                                                         'Ridge Regressor','Lasso Regressor','Linear
R','SVM - RBF Kernel','SVM - Linear Kernel','SVM - Poly Kernel'])
```

## Comparison Table for our Model

In [84]:

```
import seaborn as sns

cm = sns.light_palette("Purple", as_cmap=True)

s = data_df.style.background_gradient(cmap='tab20b')
s
```

Out[84]:

		Linear Regression	Polynomial Regressor	KNN Regressor	Ridge Regressor	Lasso Regressor	Linear SVR	SVM - RBF Kernel	SVM - Linear Kernel	SVM - Poly Kernel
Before PCA	Training Accuracy	0.5935	0.59	0.6514	0.5935	0.5894	0.3634	0.8042	0.5825	0.5612
	Test Accuracy	0.5582	0.5581	0.5782	0.5585	0.5567	0.3875	0.5161	0.5425	0.5212
After PCA	Training Accuracy	0.4594	0.5509	0.6375	0.4594	0.4253	0.4481	0.6782	0.4557	0.4498
	Test Accuracy	0.4365	0.5029	0.5246	0.4365	0.4066	0.4144	0.5401	0.4259	0.4237

## Observations

It is observed from the above steps PCA reduce the dimenstionality from 25 to 8 . Primary reason for using PCA is that it decrease time complexity of our model.Furthermore,Computation time decrease but as a penalty accuracy of the model decreases.If we refer out comparison table above before PCA model accuracy and after PCA accuracy we can find out that majority of our

model performed well before PCA and few of the model accuracy are well in after PCA. I heretore we will go ahead with before PCA model though it has higher dimensionality but accuracy is better.

## Deep Learning Model - Regression: Neural Networks

In [86]:

```
#install the packages
import numpy as np
import tensorflow as tf
from tensorflow import keras
from sklearn.model_selection import GridSearchCV
from keras.models import Sequential
from keras.layers import Dense
from keras.wrappers.scikit_learn import KerasClassifier
```

Using TensorFlow backend.

In [87]:

```
'''
Steps for creating Neural Network using Keras Classifier
'''

# Step 1: Create model
model = Sequential()
# Defining Input layer
model.add(Dense(24, input_dim = 25, activation = 'relu'))
# Defining Hidden layer
model.add(Dense(50, activation = 'relu'))
model.add(Dense(25, activation = 'relu'))
#Output layer
model.add(Dense(1, kernel_initializer='normal', activation = 'sigmoid'))

# Step 2: Build the computational graph - compile
model.compile(loss = 'mean_absolute_error', optimizer = 'adam', metrics = ['mean_absolute_error'] )

# Step 3: Train the model
model.fit(X_train, y_train, epochs = 30, batch_size = 50)
```

```
Epoch 1/30
3591/3591 [=====] - 1s 207us/step - loss: 3.9610 - mean_absolute_error: 3.9610
Epoch 2/30
3591/3591 [=====] - 0s 70us/step - loss: 3.6922 - mean_absolute_error: 3.6922
Epoch 3/30
3591/3591 [=====] - 0s 68us/step - loss: 3.6915 - mean_absolute_error: 3.6915
Epoch 4/30
3591/3591 [=====] - 0s 70us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 5/30
3591/3591 [=====] - 0s 69us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 6/30
3591/3591 [=====] - 0s 78us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 7/30
3591/3591 [=====] - 0s 78us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 8/30
3591/3591 [=====] - 0s 79us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 9/30
3591/3591 [=====] - 0s 81us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 10/30
3591/3591 [=====] - 0s 76us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 11/30
```

```

3591/3591 [=====] - 0s 67us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 12/30
3591/3591 [=====] - 0s 69us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 13/30
3591/3591 [=====] - 0s 70us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 14/30
3591/3591 [=====] - 0s 68us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 15/30
3591/3591 [=====] - 0s 67us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 16/30
3591/3591 [=====] - 0s 68us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 17/30
3591/3591 [=====] - 0s 70us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 18/30
3591/3591 [=====] - 0s 69us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 19/30
3591/3591 [=====] - 0s 69us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 20/30
3591/3591 [=====] - 0s 68us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 21/30
3591/3591 [=====] - 0s 69us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 22/30
3591/3591 [=====] - 0s 69us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 23/30
3591/3591 [=====] - 0s 79us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 24/30
3591/3591 [=====] - 0s 68us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 25/30
3591/3591 [=====] - 0s 68us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 26/30
3591/3591 [=====] - 0s 69us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 27/30
3591/3591 [=====] - 0s 71us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 28/30
3591/3591 [=====] - 0s 70us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 29/30
3591/3591 [=====] - 0s 77us/step - loss: 3.6914 - mean_absolute_error: 3.6914
Epoch 30/30
3591/3591 [=====] - 0s 69us/step - loss: 3.6914 - mean_absolute_error: 3.6914

```

Out[87]:

```
<keras.callbacks.callbacks.History at 0x294c3901288>
```

In [88]:

```

# Summarize Result
loss_and_metrics = model.evaluate(X_test, y_test)

print("Test Loss", loss_and_metrics[0])
print("Test Accuracy", loss_and_metrics[1])

```

```

1198/1198 [=====] - 0s 67us/step
Test Loss 3.749836247433008
Test Accuracy 3.7498362064361572

```



# Classification- Bank Marketing

Bank Marketing csv The data is related with direct marketing campaigns (phone calls) of a Portuguese banking institution. The classification goal is to predict if the client will subscribe a term deposit (variable y).

Data Set Information: The data is related with direct marketing campaigns of a Portuguese banking institution. The marketing campaigns were based on phone calls. Often, more than one contact to the same client was required, in order to access if the product (bank term deposit) would be ('yes') or not ('no') subscribed.

There are 21 columns and 41188 rows in this dataset. We have imported all the necessary files and libraries. We also filled the null values with mean and did the visualization using seaborn, pyplot, matplotlib.

Variables - age, job, marital, education, default, housing, loan, contact, month, day\_of\_week, duration, campaign, pdays , previous, poutcome, emp.var.rate, cons.price.idx, cons.conf.idx, euribor3m, nr.employed, y.

We will perform Classification on this dataset.

## Importing Libraries

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier

%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

## Read csv file

In [2]:

```
bnk = pd.read_csv(r"C:\Users\SST190000\Downloads\Applied ML\archive\bank_marketing.csv", sep = ';')
```

In [3]:

bnk

Out[3]:

	age	job	marital	education	default	housing	loan	contact	month	day_of_week	...	campaign	pdays	pre
0	56	housemaid	married	basic.4y	no	no	no	telephone	may	mon	...	1	999	
1	57	services	married	high.school	unknown	no	no	telephone	may	mon	...	1	999	
2	37	services	married	high.school	no	yes	no	telephone	may	mon	...	1	999	
3	40	admin.	married	basic.6y	no	no	no	telephone	may	mon	...	1	999	
4	56	services	married	high.school	no	no	yes	telephone	may	mon	...	1	999	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	
41183	73	retired	married	professional.course	no	yes	no	cellular	nov	fri	...	1	999	
41184	46	blue-collar	married	professional.course	no	no	no	cellular	nov	fri	...	1	999	
41185	56	retired	married	university.degree	no	yes	no	cellular	nov	fri	...	2	999	
41186	44	technician	married	professional.course	no	no	no	cellular	nov	fri	...	1	999	
41187	74	retired	married	professional.course	no	yes	no	cellular	nov	fri	...	3	999	

age job marital education default housing loan contact month day\_of\_week ... campaign pdays previous

41188 rows × 21 columns

## Getting information about the dataset

In [4]:

```
bnk.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41188 entries, 0 to 41187
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   41188 non-null  int64
1   job                   41188 non-null  object
2   marital               41188 non-null  object
3   education              41188 non-null  object
4   default                41188 non-null  object
5   housing                41188 non-null  object
6   loan                   41188 non-null  object
7   contact                41188 non-null  object
8   month                 41188 non-null  object
9   day_of_week           41188 non-null  object
10  duration               41188 non-null  int64
11  campaign               41188 non-null  int64
12  pdays                  41188 non-null  int64
13  previous               41188 non-null  int64
14  poutcome               41188 non-null  object
15  emp.var.rate           41188 non-null  float64
16  cons.price.idx          41188 non-null  float64
17  cons.conf.idx           41188 non-null  float64
18  euribor3m              41188 non-null  float64
19  nr.employed             41188 non-null  float64
20  y                       41188 non-null  object
dtypes: float64(5), int64(5), object(11)
memory usage: 6.6+ MB
```

In [5]:

```
bnk.describe()
```

Out[5]:

	age	duration	campaign	pdays	previous	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m
count	41188.00000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000
mean	40.02406	258.285010	2.567593	962.475454	0.172963	0.081886	93.575664	-40.502600	3.621297
std	10.42125	259.279249	2.770014	186.910907	0.494901	1.570960	0.578840	4.628198	1.734441
min	17.00000	0.000000	1.000000	0.000000	0.000000	-3.400000	92.201000	-50.800000	0.634000
25%	32.00000	102.000000	1.000000	999.000000	0.000000	-1.800000	93.075000	-42.700000	1.344000
50%	38.00000	180.000000	2.000000	999.000000	0.000000	1.100000	93.749000	-41.800000	4.857000
75%	47.00000	319.000000	3.000000	999.000000	0.000000	1.400000	93.994000	-36.400000	4.961000
max	98.00000	4918.000000	56.000000	999.000000	7.000000	1.400000	94.767000	-26.900000	5.045000

In [6]:

```
bnk.columns
```

Out[6]:

```
Index(['age', 'job', 'marital', 'education', 'default', 'housing', 'loan',
      'contact', 'month', 'day_of_week', 'duration', 'campaign', 'pdays',
      'previous', 'poutcome', 'emp.var.rate', 'cons.price.idx',
      'cons.conf.idx', 'euribor3m', 'nr.employed', 'y'],
      dtype='object')
```



In [7]:

```
bnk.shape
```

Out[7]:

```
(41188, 21)
```

## Getting information about the dataset

In [8]:

```
bnk.isnull().sum()
```

Out[8]:

```
age          0
job          0
marital      0
education    0
default      0
housing      0
loan         0
contact      0
month        0
day_of_week  0
duration     0
campaign     0
pdays       0
previous     0
poutcome     0
emp.var.rate 0
cons.price.idx 0
cons.conf.idx 0
euribor3m    0
nr.employed  0
y            0
dtype: int64
```

## Check for Duplicates and Null values

In [9]:

```
bnk.duplicated().sum()
```

Out[9]:

```
12
```

In [10]:

```
duplicates = bnk[bnk.duplicated()]
print("Duplicate rows : ")
duplicates
```

Duplicate rows :

Out[10]:

	age	job	marital	education	default	housing	loan	contact	month	day_of_week	...	campaign	pdays	pre
1266	39	blue-collar	married	basic.6y	no	no	no	telephone	may	thu ...		1	999	
12261	36	retired	married	unknown	no	no	no	telephone	jul	thu ...		1	999	
14234	27	technician	single	professional.course	no	no	no	cellular	jul	mon ...		2	999	
14050	47	technician	divorced	high.school	no	no	no	cellular	jul	thu ...		2	999	

109956	4/	technician	divorced	nigh.school	no	yes	no	cellular	jul	day_of_week	tru	...	...	campaign	3	999	pre
age	job	marital	education	default	housing	loan	contact	month	day_of_week	...	...	...	...	...	...	...	...
18465	32	technician	single	professional.course	no	yes	no	cellular	jul	thu	...	...	...	1	999	...	...
20216	55	services	married	high.school	unknown	no	no	cellular	aug	mon	...	...	...	1	999	...	...
20534	41	technician	married	professional.course	no	yes	no	cellular	aug	tue	...	...	...	1	999	...	...
25217	39	admin.	married	university.degree	no	no	no	cellular	nov	tue	...	...	...	2	999	...	...
28477	24	services	single	high.school	no	yes	no	cellular	apr	tue	...	...	...	1	999	...	...
32516	35	admin.	married	university.degree	no	yes	no	cellular	may	fri	...	...	...	4	999	...	...
36951	45	admin.	married	university.degree	no	no	no	cellular	jul	thu	...	...	...	1	999	...	...
38281	71	retired	single	university.degree	no	no	no	telephone	oct	tue	...	...	...	1	999	...	...

12 rows × 21 columns

## Drop duplicates

In [11]:

```
bnk.drop_duplicates(keep=False, inplace=True)
bnk.duplicated().sum()
```

Out[11]:

0

In [12]:

```
bnk.isna().sum()
```

Out[12]:

```
age          0
job          0
marital      0
education    0
default      0
housing      0
loan         0
contact      0
month        0
day_of_week  0
duration     0
campaign     0
pdays       0
previous     0
poutcome     0
emp.var.rate 0
cons.price.idx 0
cons.conf.idx 0
euribor3m    0
nr.employed  0
y            0
dtype: int64
```

## Imputing 5%-10% null values

In [13]:

```
bnk['emp.var.rate'] = bnk['emp.var.rate'].mask(np.random.random(bnk['emp.var.rate'].shape) < .1)
bnk['duration'] = bnk['duration'].mask(np.random.random(bnk['duration'].shape) < .05)
bnk['campaign'] = bnk['campaign'].mask(np.random.random(bnk['campaign'].shape) < .1)
bnk['pdays'] = bnk['pdays'].mask(np.random.random(bnk['pdays'].shape) < .1)
bnk['previous'] = bnk['previous'].mask(np.random.random(bnk['previous'].shape) < .05)
```

In [14]:

```
bnk.isna().sum()
```

```
Out[14]:
```

```
age                0
job                0
marital            0
education          0
default            0
housing            0
loan              0
contact            0
month              0
day_of_week        0
duration           2138
campaign           4085
pdays             4088
previous           1980
poutcome           0
emp.var.rate       4192
cons.price.idx      0
cons.conf.idx       0
euribor3m           0
nr.employed         0
y                  0
dtype: int64
```

```
In [15]:
```

```
bnk.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 41164 entries, 0 to 41187
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   age                   41164 non-null  int64
 1   job                   41164 non-null  object
 2   marital               41164 non-null  object
 3   education             41164 non-null  object
 4   default               41164 non-null  object
 5   housing               41164 non-null  object
 6   loan                  41164 non-null  object
 7   contact               41164 non-null  object
 8   month                 41164 non-null  object
 9   day_of_week           41164 non-null  object
10   duration              39026 non-null  float64
11   campaign              37079 non-null  float64
12   pdays                 37076 non-null  float64
13   previous              39184 non-null  float64
14   poutcome              41164 non-null  object
15   emp.var.rate          36972 non-null  float64
16   cons.price.idx         41164 non-null  float64
17   cons.conf.idx          41164 non-null  float64
18   euribor3m             41164 non-null  float64
19   nr.employed           41164 non-null  float64
20   y                     41164 non-null  object
dtypes: float64(9), int64(1), object(11)
memory usage: 6.9+ MB
```

## Replace null values with mean

```
In [16]:
```

```
bnk['emp.var.rate'].fillna(bnk['emp.var.rate'].mean(), inplace=True)
bnk['duration'].fillna(bnk['duration'].mean(), inplace=True)
bnk['campaign'].fillna(bnk['campaign'].mean(), inplace=True)
bnk['pdays'].fillna(bnk['pdays'].mean(), inplace=True)
bnk['previous'].fillna(bnk['previous'].mean(), inplace=True)
```

```
In [17]:
```

```
bnk.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 41164 entries, 0 to 41187
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   41164 non-null  int64
1   job                   41164 non-null  object
2   marital               41164 non-null  object
3   education             41164 non-null  object
4   default              41164 non-null  object
5   housing              41164 non-null  object
6   loan                 41164 non-null  object
7   contact              41164 non-null  object
8   month                41164 non-null  object
9   day_of_week          41164 non-null  object
10  duration              41164 non-null  float64
11  campaign              41164 non-null  float64
12  pdays                41164 non-null  float64
13  previous              41164 non-null  float64
14  poutcome              41164 non-null  object
15  emp.var.rate          41164 non-null  float64
16  cons.price.idx         41164 non-null  float64
17  cons.conf.idx          41164 non-null  float64
18  euribor3m             41164 non-null  float64
19  nr.employed            41164 non-null  float64
20  y                     41164 non-null  object
dtypes: float64(9), int64(1), object(11)
memory usage: 6.9+ MB
```

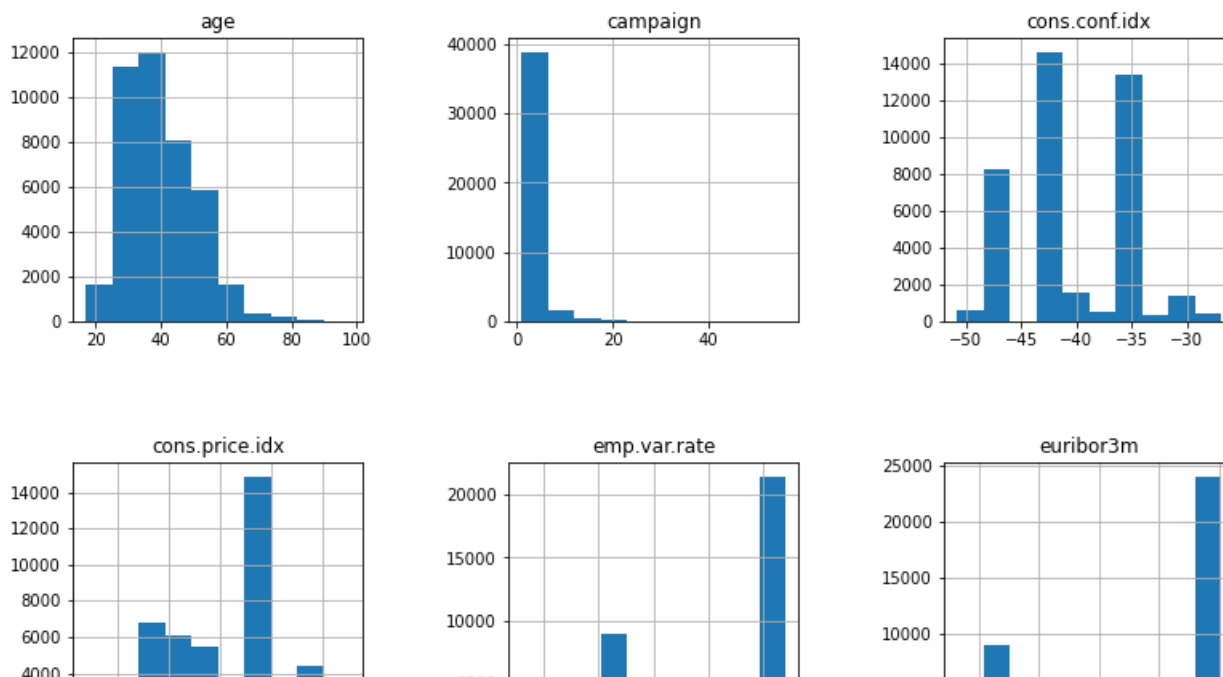
## Exploratory Data Analysis

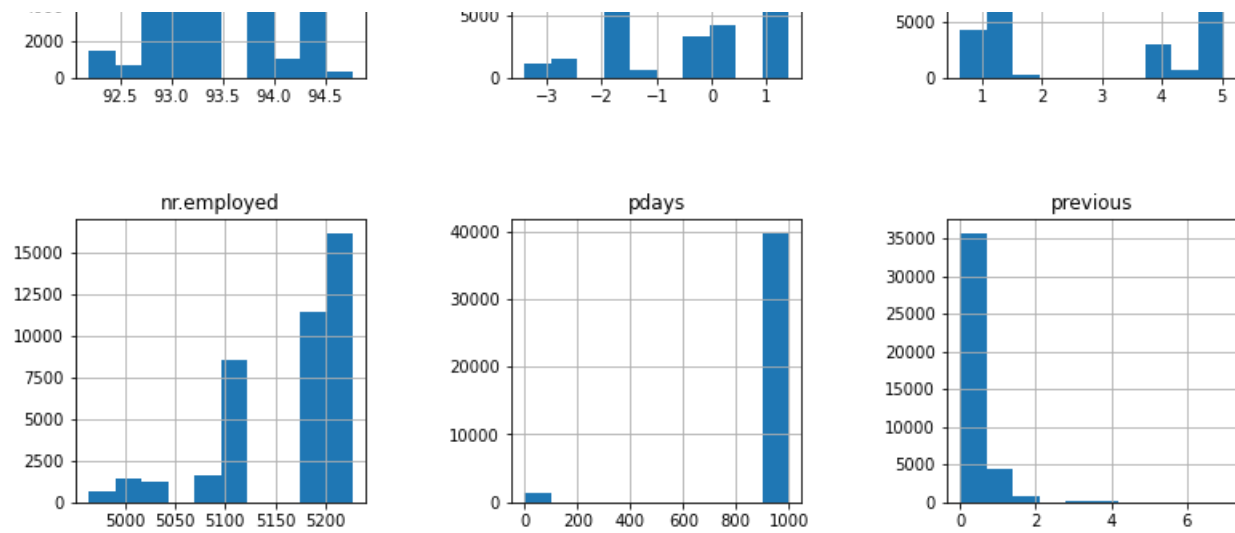
### Exploring numerical variables in 'bnk'

#### Histogram Subplots

```
In [18]:
```

```
col = ['age', 'campaign', 'pdays', 'previous', 'emp.var.rate', 'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed']
bnk.hist(column=col, figsize=(13,13))
plt.subplots_adjust(wspace = 0.5, hspace = 0.5)
plt.show()
```






## Distribution of age variable

In [19]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
sns.distplot(bnk['age'], hist=True, kde=True,
             bins=int(150/10), color = 'purple',
             hist_kws={'edgecolor':'black'})
bca.set_title('Age Distribution', fontsize=15)
```

Out[19]:

Text(0.5, 1.0, 'Age Distribution')




## Count of Duration

In [20]:

```
sns.countplot(x='duration', data=bnk)
bca.set_title('Count of Duration', fontsize=15)
```

Out[20]:

Text(0.5, 1.0, 'Count of Duration')



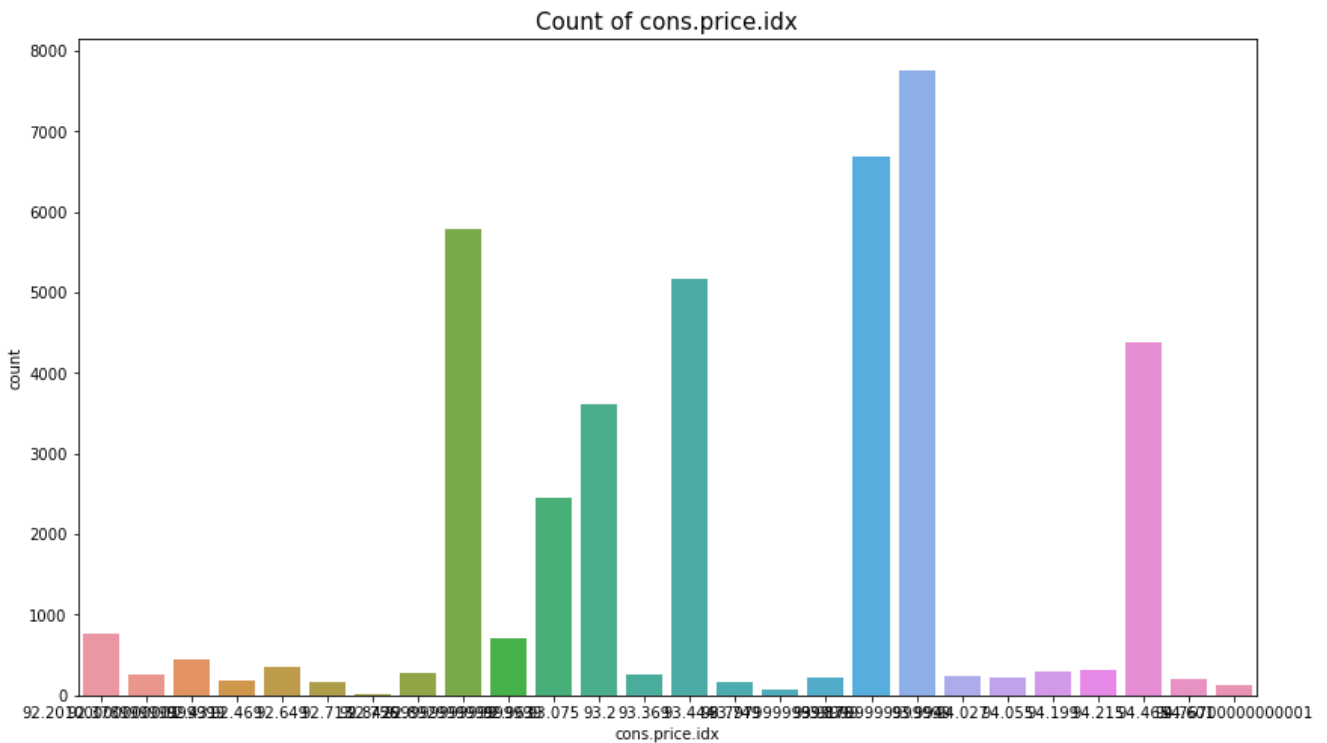
## Count of cons.price.idx

In [21]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
sns.countplot(x='cons.price.idx', data=bnk)
bca.set_title('Count of cons.price.idx', fontsize=15)
```

Out[21]:

Text(0.5, 1.0, 'Count of cons.price.idx')



## Count of emp.var.rate


In [22]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
sns.countplot(x='emp.var.rate', data=bnk)
bca.set_title('Count of emp.var.rate', fontsize=15)
```

```
bca.set_title('Count of emp.var.rate', fontsize=15)
```

Out[22]:

```
Text(0.5, 1.0, 'Count of emp.var.rate')
```



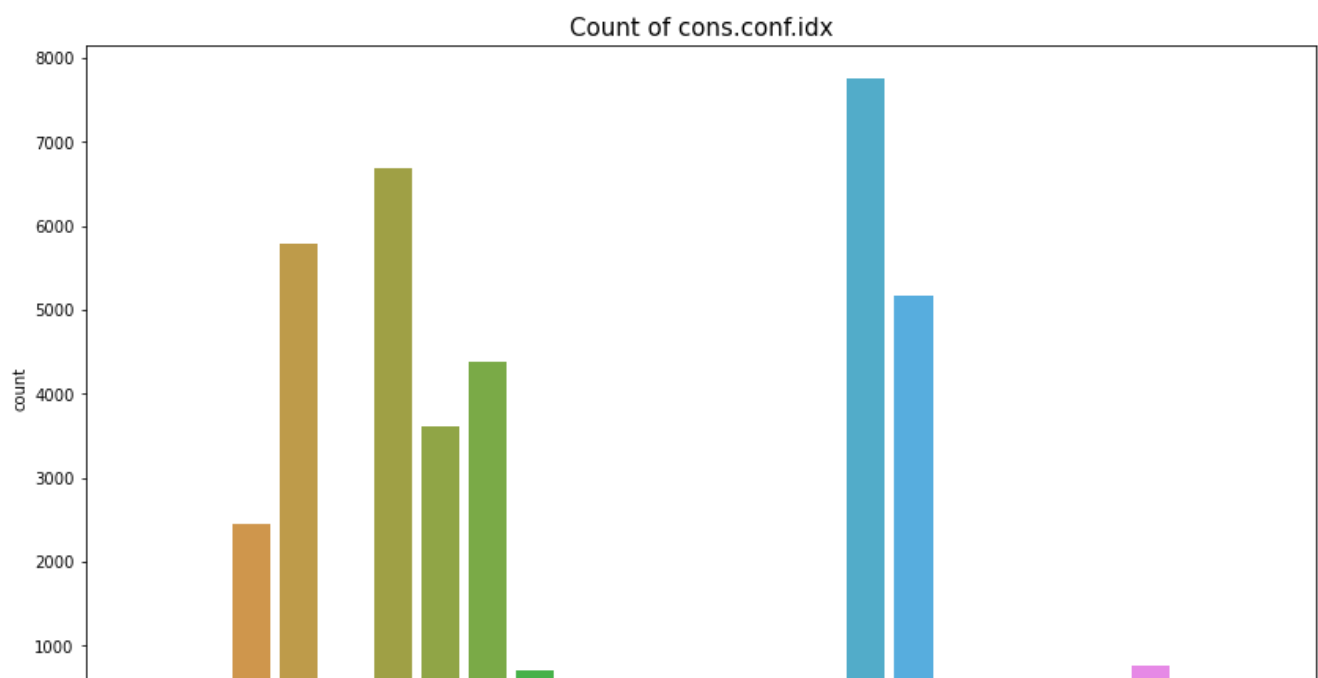
## Count of cons.conf.idx


In [23]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
sns.countplot(x='cons.conf.idx', data=bnk)
bca.set_title('Count of cons.conf.idx', fontsize=15)
```

Out[23]:

```
Text(0.5, 1.0, 'Count of cons.conf.idx')
```





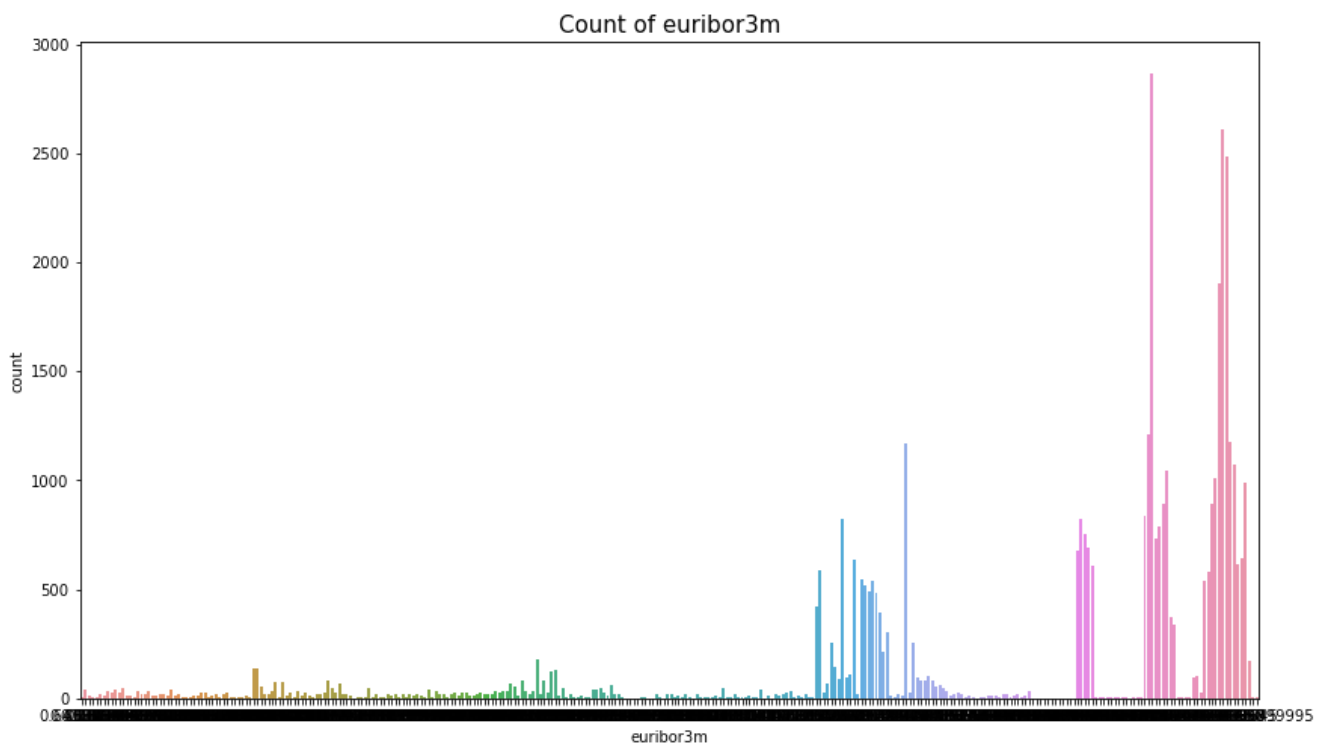
## Count of euribor3m

In [24]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
sns.countplot(x='euribor3m', data=bnk)
bca.set_title('Count of euribor3m', fontsize=15)
```

Out[24]:

Text(0.5, 1.0, 'Count of euribor3m')




## Exploring variation of numerical variables w.r.t target variable y


### Variation in age

In [25]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot(x='y', y='age', data=bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Age', fontsize=15)
bca1.set_title('Variation in age', fontsize=15)
bca1.tick_params(labelsize=15)
```








## Variation in duration

In [26]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot( x='y', y= 'duration', data =bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Duration', fontsize=15)
bca1.set_title('Variation in duration', fontsize=15)
bca1.tick_params(labelsize=15)
```

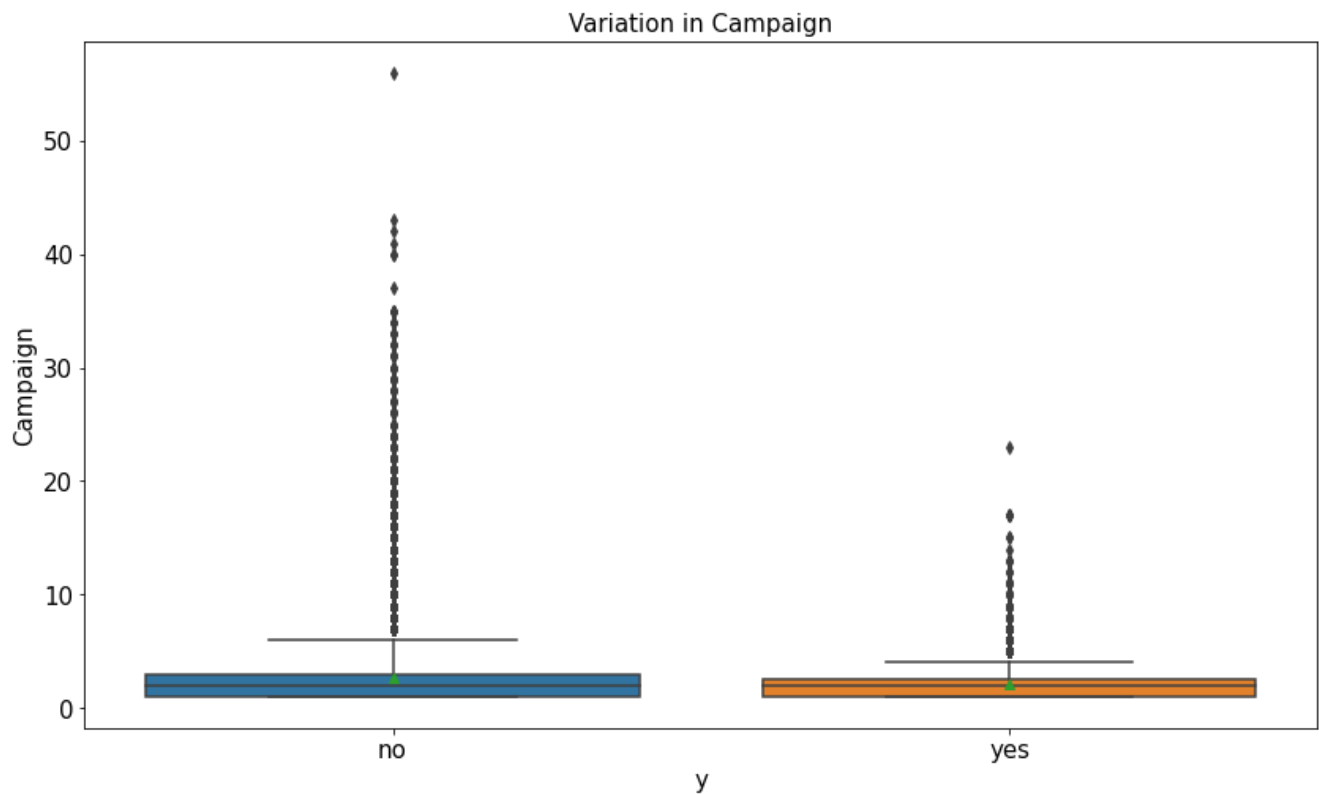


## Variation in Campaign

In [27]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot( x='y', y= 'campaign', data =bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Campaign', fontsize=15)
```


```
bca1.set_title('Variation in Campaign', fontsize=15)
bca1.tick_params(labelsize=15)
```



## Variation in Pdays

In [28]:


```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot(x='y', y='pdays', data=bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Pdays', fontsize=15)
bca1.set_title('Variation in Pdays', fontsize=15)
bca1.tick_params(labelsize=15)
```



## Variation in Previous

In [29]:


```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot(x='y', y='previous', data=bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Previous', fontsize=15)
bca1.set_title('Variation in Previous', fontsize=15)
bca1.tick_params(labelsize=15)
```




## Variation in Emp.Var.Rate

In [30]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot(x='y', y='emp.var.rate', data=bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Emp.Var.Rate', fontsize=15)
bca1.set_title('Variation in Emp.Var.Rate', fontsize=15)
bca1.tick_params(labelsize=15)
```






## Variation in Cons.Price.Idx

In [31]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot( x='y', y= 'cons.price.idx', data =bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Cons.Price.Idx', fontsize=15)
bca1.set_title('Variation in Cons.Price.Idx', fontsize=15)
bca1.tick_params(labelsize=15)
```




## Variation in Cons.Conf.Idx

In [32]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot( x='y', y= 'cons.conf.idx', data =bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Cons.Conf.Idx', fontsize=15)
bca1.set_title('Variation in Cons.Conf.Idx', fontsize=15)
bca1.tick_params(labelsize=15)
```


Variation in Cons.Conf.Idx



## Variation in Euribor3m

In [33]:

```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot(x='y', y='euribor3m', data=bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Euribor3m', fontsize=15)
bca1.set_title('Variation in Euribor3m', fontsize=15)
bca1.tick_params(labelsize=15)
```




## Variation in Nr Employed

## Variation in Nr.Employed

In [34]:


```
fig, bca = plt.subplots()
fig.set_size_inches(14, 8)
bca1 = sns.boxplot( x='y', y= 'nr.employed', data =bnk, showmeans=True)
bca1.set_xlabel('y', fontsize=15)
bca1.set_ylabel('Nr.Employed', fontsize=15)
bca1.set_title('Variation in Nr.Employed', fontsize=15)
bca1.tick_params(labels=15)
```




## Heatmap depicting correlation between all numerical variables

In [35]:

```
plt.subplots(figsize=(14,8))
sns.heatmap(bnk.corr(), annot=True)
plt.show()
```





## Encoding and storing target variable 'y'

We perform one-hot-encoding on target variable 'y' in bnk dataframe as it is categorical data. We store the result in a new variable 'y'.

In [36]:

```
y = pd.get_dummies(bnk['y'], columns = ['y'], prefix = ['y'], drop_first = True)
bnk.head()
```

Out[36]:

	age	job	marital	education	default	housing	loan	contact	month	day_of_week	...	campaign	pdays	previous	po
0	56	housemaid	married	basic.4y	no	no	no	telephone	may	mon	...	2.570404	999.0	0.172596	non
1	57	services	married	high.school	unknown	no	no	telephone	may	mon	...	1.000000	999.0	0.000000	non
2	37	services	married	high.school	no	yes	no	telephone	may	mon	...	1.000000	999.0	0.000000	non
3	40	admin.	married	basic.6y	no	no	no	telephone	may	mon	...	1.000000	999.0	0.000000	non
4	56	services	married	high.school	no	no	yes	telephone	may	mon	...	1.000000	999.0	0.000000	non

5 rows × 21 columns

## Creating a new dataframe 'bank\_client'

We are creating the bank\_client dataset to store information of bank clients. The attributes included are namely - age, job, marital, education, default, housing, loan

In [37]:

```
bank_client = bnk.iloc[:, 0:7]
bank_client.head()
```

Out[37]:

	age	job	marital	education	default	housing	loan
0	56	housemaid	married	basic.4y	no	no	no
1	57	services	married	high.school	unknown	no	no
2	37	services	married	high.school	no	yes	no
3	40	admin.	married	basic.6y	no	no	no
4	56	services	married	high.school	no	no	yes

## Exploring variables in bank\_client

### Age Count distribution


In [38]:

```
fig, bca = plt.subplots()
fig.set_size_inches(30, 15)
sns.countplot(x = 'age', data = bank_client)
bca.set_xlabel('Age', fontsize=25)
```

```
bca.set_ylabel('Count', fontsize=25)
bca.set_title('Age Count Distribution', fontsize=25)
```

Out[38]:


Text(0.5, 1.0, 'Age Count Distribution')



## Age Distribution

In [39]:

```
bca1 = sns.boxplot( y=bank_client["age"] )
bca1.set_xlabel('People Age', fontsize=15)
bca1.set_ylabel('Age', fontsize=15)
bca1.set_title('Age Distribution', fontsize=15)
bca1.tick_params(labels=15)
```



## Jobs Count Distribution


In [40]:

```
fig, bca = plt.subplots()
fig.set_size_inches(30, 10)
sns.countplot(x = 'job', data = bank_client)
bca.set_xlabel('Jobs', fontsize=25)
bca.set_ylabel('Count', fontsize=25)
bca.set_title('Jobs Count Distribution', fontsize=25)
```



Out[40]:

```
Text(0.5, 1.0, 'Jobs Count Distribution')
```




## Marital Status Count Distribution

In [41]:

```
fig, bca = plt.subplots()
fig.set_size_inches(10, 5)
sns.countplot(x = 'marital', data = bank_client)
bca.set_xlabel('Marital Status', fontsize=15)
bca.set_ylabel('Count', fontsize=15)
bca.set_title('Marital Status Count Distribution', fontsize=25)
```

Out[41]:

```
Text(0.5, 1.0, 'Marital Status Count Distribution')
```




## Education Count Distribution

In [42]:

```
fig, bca = plt.subplots()
fig.set_size_inches(20, 5)
sns.countplot(x = 'education', data = bank_client)
bca.set_xlabel('Education', fontsize=15)
bca.set_ylabel('Count', fontsize=15)
bca.set_title('Education Count Distribution', fontsize=25)
```

Out[42]:

Text(0.5, 1.0, 'Education Count Distribution')




## Housing Loan Count Distribution

In [43]:

```
fig, bca = plt.subplots()
fig.set_size_inches(10, 5)
sns.countplot(x = 'housing', data = bank_client)
bca.set_xlabel('Housing Loan', fontsize=15)
bca.set_ylabel('Count', fontsize=15)
bca.set_title('Housing Loan Count Distribution', fontsize=15)
```

Out[43]:

Text(0.5, 1.0, 'Housing Loan Count Distribution')




## Personal Loan Count Distribution

In [44]:

```
fig, bca = plt.subplots()
fig.set_size_inches(10, 5)
sns.countplot(x = 'loan', data = bank_client)
bca.set_xlabel('Personal Loan', fontsize=15)
bca.set_ylabel('Count', fontsize=15)
bca.set_title('Personal Loan Count Distribution', fontsize=15)
```

Out[44]:

Text(0.5, 1.0, 'Personal Loan Count Distribution')




## Default Credit Count Distribution

In [45]:

```
fig, bca = plt.subplots()
fig.set_size_inches(10, 5)
sns.countplot(x = 'default', data = bank_client)
bca.set_xlabel('Default Credit', fontsize=15)
bca.set_ylabel('Count', fontsize=15)
bca.set_title('Default Credit Count Distribution', fontsize=15)
```

Out[45]:

Text(0.5, 1.0, 'Default Credit Count Distribution')



## Treating categorical variables

In [46]:

```
bank_client['job'].unique()
```

Out[46]:

```
array(['housemaid', 'services', 'admin.', 'blue-collar', 'technician',
       'retired', 'management', 'unemployed', 'self-employed', 'unknown',
       'entrepreneur', 'student'], dtype=object)
```

In [47]:

```
dummy = pd.get_dummies(bank_client['job'],prefix = 'Job_N')
print(dummy)
```

```

      Job_N_admin.  Job_N_blue-collar  Job_N_entrepreneur  Job_N_housemaid  \
0                0                0                0                1
1                0                0                0                0
2                0                0                0                0
3                1                0                0                0
4                0                0                0                0
...
41183            0                0                0                0
41184            0                1                0                0
41185            0                0                0                0
41186            0                0                0                0
41187            0                0                0                0

      Job_N_management  Job_N_retired  Job_N_self-employed  Job_N_services  \
0                    0                0                0                0
1                    0                0                0                1
2                    0                0                0                1
3                    0                0                0                0
4                    0                0                0                1
...
41183            0                1                0                0
41184            0                0                0                0
41185            0                1                0                0
41186            0                0                0                0
41187            0                1                0                0

      Job_N_student  Job_N_technician  Job_N_unemployed  Job_N_unknown
0                0                0                0                0
1                0                0                0                0
2                0                0                0                0
3                0                0                0                0
4                0                0                0                0
...
41183            0                0                0                0
41184            0                0                0                0
41185            0                0                0                0
41186            0                1                0                0
41187            0                0                0                0

```

[41164 rows x 12 columns]

In [48]:

```
bank_client = bank_client.join(dummy)
bank_client
```

Out[48]:

	age	job	marital	education	default	housing	loan	Job_N_admin.	Job_N_blue-collar	Job_N_entrepreneur	Job_N_I
0	56	housemaid	married	basic.4y	no	no	no	0	0	0	
1	57	services	married	high.school	unknown	no	no	0	0	0	
2	37	services	married	high.school	no	yes	no	0	0	0	
3	40	admin.	married	basic.6y	no	no	no	1	0	0	
4	56	services	married	high.school	no	no	yes	0	0	0	
...	...	...	...	...	...	...	...	...	...	...	
41183	73	retired	married	professional.course	no	yes	no	0	0	0	
41184	46	blue-collar	married	professional.course	no	no	no	0	1	0	
41185	56	retired	married	university.degree	no	yes	no	0	0	0	
41186	44	technician	married	professional.course	no	no	no	0	0	0	
41187	74	retired	married	professional.course	no	yes	no	0	0	0	

41164 rows x 19 columns

In [49]:

```
bank_client['marital'].unique()
```

Out[49]:

```
array(['married', 'single', 'divorced', 'unknown'], dtype=object)
```

In [50]:

```
lc=LabelEncoder()  
bank_client['Marital_N']=lc.fit_transform(bank_client['marital'])  
bank_client
```

Out[50]:

	age	job	marital	education	default	housing	loan	Job_N_admin.	Job_N_blue-collar	Job_N_entrepreneur	Job_N_I
0	56	housemaid	married	basic.4y	no	no	no	0	0	0	
1	57	services	married	high.school	unknown	no	no	0	0	0	
2	37	services	married	high.school	no	yes	no	0	0	0	
3	40	admin.	married	basic.6y	no	no	no	1	0	0	
4	56	services	married	high.school	no	no	yes	0	0	0	
...	...	...	...	...	...	...	...	...	...	...	
41183	73	retired	married	professional.course	no	yes	no	0	0	0	
41184	46	blue-collar	married	professional.course	no	no	no	0	1	0	
41185	56	retired	married	university.degree	no	yes	no	0	0	0	
41186	44	technician	married	professional.course	no	no	no	0	0	0	
41187	74	retired	married	professional.course	no	yes	no	0	0	0	

41164 rows × 20 columns

In [51]:

```
bank_client['education'].unique()
```

Out[51]:

```
array(['basic.4y', 'high.school', 'basic.6y', 'basic.9y',  
      'professional.course', 'unknown', 'university.degree',  
      'illiterate'], dtype=object)
```

In [52]:

```
bank_client=pd.concat((bank_client,pd.get_dummies(bank_client['education'])),axis=1)
```

In [53]:

```
bank_client['default'].unique()
```

Out[53]:

```
array(['no', 'unknown', 'yes'], dtype=object)
```

In [54]:

```
bank_client['housing'].unique()
```

Out[54]:

```
array(['no', 'yes', 'unknown'], dtype=object)
```

In [55]:

```
bank_client['loan'].unique()
```

Out[55]:

```
array(['no', 'yes', 'unknown'], dtype=object)
```

In [56]:

```
lc=LabelEncoder()  
bank_client['Default_N']=lc.fit_transform(bank_client['default'])
```

In [57]:

```
lc=LabelEncoder()  
bank_client['Housing_N']=lc.fit_transform(bank_client['housing'])
```

In [58]:

```
lc=LabelEncoder()  
bank_client['Loan_N']=lc.fit_transform(bank_client['loan'])
```

In [59]:

```
bank_client.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 41164 entries, 0 to 41187  
Data columns (total 31 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   age                   41164 non-null  int64  
1   job                   41164 non-null  object  
2   marital               41164 non-null  object  
3   education             41164 non-null  object  
4   default               41164 non-null  object  
5   housing               41164 non-null  object  
6   loan                  41164 non-null  object  
7   Job_N_admin.          41164 non-null  uint8  
8   Job_N_blue-collar     41164 non-null  uint8  
9   Job_N_entrepreneur    41164 non-null  uint8  
10  Job_N_housemaid       41164 non-null  uint8  
11  Job_N_management      41164 non-null  uint8  
12  Job_N_retired         41164 non-null  uint8  
13  Job_N_self-employed   41164 non-null  uint8  
14  Job_N_services        41164 non-null  uint8  
15  Job_N_student         41164 non-null  uint8  
16  Job_N_technician      41164 non-null  uint8  
17  Job_N_unemployed      41164 non-null  uint8  
18  Job_N_unknown         41164 non-null  uint8  
19  Marital_N             41164 non-null  int32  
20  basic.4y              41164 non-null  uint8  
21  basic.6y              41164 non-null  uint8  
22  basic.9y              41164 non-null  uint8  
23  high.school           41164 non-null  uint8  
24  illiterate            41164 non-null  uint8  
25  professional.course   41164 non-null  uint8  
26  university.degree     41164 non-null  uint8  
27  unknown               41164 non-null  uint8  
28  Default_N             41164 non-null  int32  
29  Housing_N             41164 non-null  int32  
30  Loan_N                41164 non-null  int32  
dtypes: int32(4), int64(1), object(6), uint8(20)  
memory usage: 5.2+ MB
```

In [60]:

```
bank_client = bank_client.drop(['job', 'marital', 'education', 'housing', 'default', 'loan'], axis  
= 1)
```

In [61]:

```
bank_client.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 41164 entries, 0 to 41187
Data columns (total 25 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   age                   41164 non-null  int64
 1   Job_N_admin.          41164 non-null  uint8
 2   Job_N_blue-collar    41164 non-null  uint8
 3   Job_N_entrepreneur    41164 non-null  uint8
 4   Job_N_housemaid       41164 non-null  uint8
 5   Job_N_management      41164 non-null  uint8
 6   Job_N_retired         41164 non-null  uint8
 7   Job_N_self-employed   41164 non-null  uint8
 8   Job_N_services        41164 non-null  uint8
 9   Job_N_student         41164 non-null  uint8
10   Job_N_technician      41164 non-null  uint8
11   Job_N_unemployed      41164 non-null  uint8
12   Job_N_unknown         41164 non-null  uint8
13   Marital_N            41164 non-null  int32
14   basic.4y             41164 non-null  uint8
15   basic.6y             41164 non-null  uint8
16   basic.9y             41164 non-null  uint8
17   high.school          41164 non-null  uint8
18   illiterate           41164 non-null  uint8
19   professional.course   41164 non-null  uint8
20   university.degree     41164 non-null  uint8
21   unknown              41164 non-null  uint8
22   Default_N            41164 non-null  int32
23   Housing_N            41164 non-null  int32
24   Loan_N               41164 non-null  int32
dtypes: int32(4), int64(1), uint8(20)
memory usage: 3.3 MB
```

In [62]:

```
bank_client['age'] = bank_client['age'].astype(int)
```

In [63]:

```
bank_client.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 41164 entries, 0 to 41187
Data columns (total 25 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   age                   41164 non-null  int32
 1   Job_N_admin.          41164 non-null  uint8
 2   Job_N_blue-collar    41164 non-null  uint8
 3   Job_N_entrepreneur    41164 non-null  uint8
 4   Job_N_housemaid       41164 non-null  uint8
 5   Job_N_management      41164 non-null  uint8
 6   Job_N_retired         41164 non-null  uint8
 7   Job_N_self-employed   41164 non-null  uint8
 8   Job_N_services        41164 non-null  uint8
 9   Job_N_student         41164 non-null  uint8
10   Job_N_technician      41164 non-null  uint8
11   Job_N_unemployed      41164 non-null  uint8
12   Job_N_unknown         41164 non-null  uint8
13   Marital_N            41164 non-null  int32
14   basic.4y             41164 non-null  uint8
15   basic.6y             41164 non-null  uint8
16   basic.9y             41164 non-null  uint8
17   high.school          41164 non-null  uint8
18   illiterate           41164 non-null  uint8
19   professional.course   41164 non-null  uint8
20   university.degree     41164 non-null  uint8
21   unknown              41164 non-null  uint8
22   Default_N            41164 non-null  int32
--  --
--  --
```

```
23 Housing_N          41164 non-null  int32
24 Loan_N              41164 non-null  int32
dtypes: int32(5), uint8(20)
memory usage: 3.1 MB
```

In [64]:

```
def age(dataframe):
    dataframe.loc[dataframe['age'] <= 32, 'age'] = 1
    dataframe.loc[(dataframe['age'] > 32) & (dataframe['age'] <= 47), 'age'] = 2
    dataframe.loc[(dataframe['age'] > 47) & (dataframe['age'] <= 70), 'age'] = 3
    dataframe.loc[(dataframe['age'] > 70) & (dataframe['age'] <= 98), 'age'] = 4

    return dataframe

age(bank_client) ;
```

In [65]:

```
bank_client.head()
```

Out [65]:

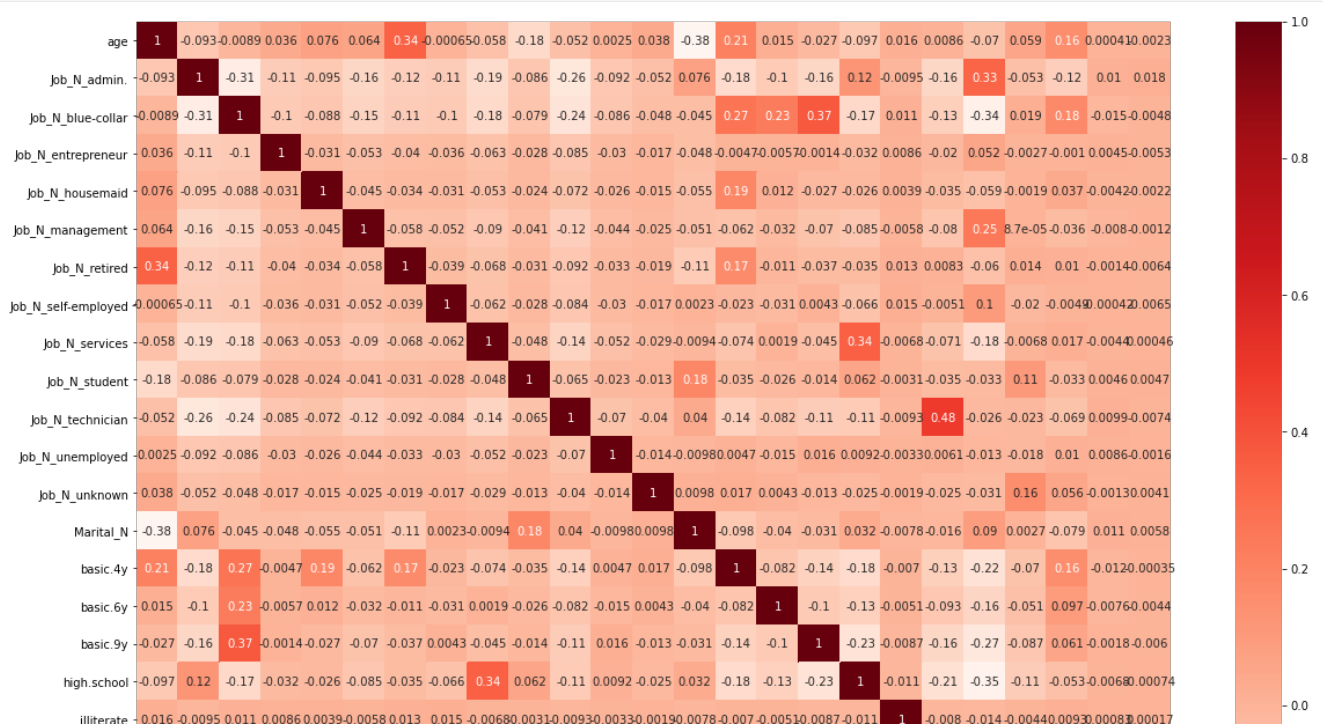
	age	Job_N_admin.	Job_N_blue-collar	Job_N_entrepreneur	Job_N_housemaid	Job_N_management	Job_N_retired	Job_N_self-employed	Job_N_s
0	3	0	0	0	1	0	0	0	
1	3	0	0	0	0	0	0	0	
2	2	0	0	0	0	0	0	0	
3	2	1	0	0	0	0	0	0	
4	3	0	0	0	0	0	0	0	

5 rows × 25 columns

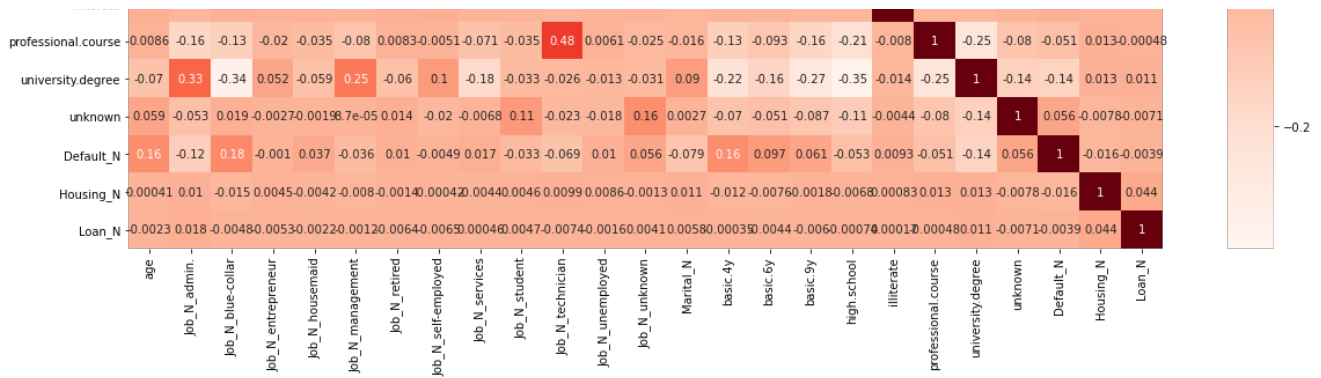
## Correlation between variables in bank\_client dataset

In [66]:

```
plt.figure(figsize=(20,15))
cor = bank_client.corr()
sns.heatmap(cor, annot=True, cmap=plt.cm.Reds)
plt.show()
```







## Creating new dataset 'other\_attr'

We are now creating a dataset to store the attributes - contact, month, day\_of\_week and duration.

In [67]:

```
other_attr = bnk.iloc[:, 7:11]
other_attr.head()
```

Out [67]:

	contact	month	day_of_week	duration
0	telephone	may	mon	261.0
1	telephone	may	mon	149.0
2	telephone	may	mon	226.0
3	telephone	may	mon	151.0
4	telephone	may	mon	307.0

## Check for null values

In [68]:

```
other_attr.isnull().sum()
```

Out [68]:

```
contact      0
month        0
day_of_week  0
duration     0
dtype: int64
```


## Exploring the attributes in 'other\_attr'


### Calls duration

In [69]:

```
dur = sns.boxplot(x = 'duration', data = other_attr)
dur.set_xlabel('Calls', fontsize=15)
dur.set_ylabel('Duration', fontsize=15)
dur.set_title('Calls Distribution', fontsize=15)
dur.tick_params(labelsize=20)
```

Calls Distribution






## Contacts Count

In [70]:

```
fig, bca = plt.subplots()
fig.set_size_inches(10, 5)
sns.countplot(x = 'contact', data = other_attr)
bca.set_xlabel('Contacts', fontsize=15)
bca.set_ylabel('Count', fontsize=15)
bca.set_title('Contacts Count', fontsize=15)
```

Out[70]:

Text(0.5, 1.0, 'Contacts Count')



## Months Count


In [71]:

```
fig, bca = plt.subplots()
fig.set_size_inches(10, 5)
sns.countplot(x = 'month', data = other_attr, order = ['mar', 'apr', 'may', 'jun', 'jul', 'aug', 'sep', 'oct', 'nov', 'dec'])
bca.set_xlabel('Months', fontsize=15)
bca.set_ylabel('Count', fontsize=15)
bca.set_title('Months Count', fontsize=15)
```

Out[71]:

Text(0.5, 1.0, 'Months Count')






## Days Of Week Count

In [72]:

```
fig, bca = plt.subplots()
fig.set_size_inches(10, 5)
sns.countplot(x = 'day_of_week', data = other_attr)
bca.set_xlabel('Days', fontsize=15)
bca.set_ylabel('Count', fontsize=15)
bca.set_title('DaysOfWeek Count', fontsize=15)
```

Out[72]:

Text(0.5, 1.0, 'DaysOfWeek Count')



## Treating categorical variables

In [73]:

```
other_attr.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 41164 entries, 0 to 41187
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   contact     41164 non-null  object
1   month       41164 non-null  object
2   day_of_week 41164 non-null  object
3   duration    41164 non-null  float64
dtypes: float64(1), object(3)
memory usage: 2.8+ MB
```

In [74]:

```
other_attr['contact'].unique()
```

Out[74]:

```
array(['telephone', 'cellular'], dtype=object)
```

In [75]:

```
other_attr['contact'] = other_attr['contact'].map({'telephone':1, 'cellular':2}).astype(int)
```

In [76]:

```
other_attr['month'].unique()
```

Out[76]:

```
array(['may', 'jun', 'jul', 'aug', 'oct', 'nov', 'dec', 'mar', 'apr',  
      'sep'], dtype=object)
```

In [77]:

```
other_attr["month"] = other_attr["month"].str.capitalize()
```

In [78]:

```
other_attr["month"] = pd.to_datetime(other_attr.month, format='%b', errors='coerce').dt.month  
other_attr = other_attr.sort_values(by="month")
```

In [79]:

```
other_attr['month'].unique()
```

Out[79]:

```
array([ 3,  4,  5,  6,  7,  8,  9, 10, 11, 12], dtype=int64)
```

In [80]:

```
other_attr['day_of_week'].unique()
```

Out[80]:

```
array(['tue', 'mon', 'thu', 'wed', 'fri'], dtype=object)
```

In [81]:

```
lc=LabelEncoder()  
other_attr['day_of_week']=lc.fit_transform(other_attr['day_of_week'])
```

In [82]:

```
other_attr['day_of_week'].unique()
```

Out[82]:

```
array([3, 1, 2, 4, 0])
```

In [83]:

```
other_attr['duration'] = other_attr['duration'].astype(int)
```

## Creating new dataset 'cont\_attr'

Here we are creating a new dataset for the social and economic context attributes, which are - emp.var.rate, cons.price.idx, cons.conf.idx, euribor3m, nr.employed.

In [84]:

```
cont_attr= bnk.loc[ , ['emp.var.rate', 'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed']]
cont_attr.head()
```

Out[84]:

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	nr.employed
0	1.100000	93.994	-36.4	4.857	5191.0
1	0.079252	93.994	-36.4	4.857	5191.0
2	1.100000	93.994	-36.4	4.857	5191.0
3	1.100000	93.994	-36.4	4.857	5191.0
4	1.100000	93.994	-36.4	4.857	5191.0

## Creating a dataset 'remain\_attr'

This dataset contains all the remaining attributes ( y excluded). These are - campaign, pdays, previous, poutcome

In [85]:

```
remain_attr = bnk.loc[ , ['campaign', 'pdays', 'previous', 'poutcome']]
remain_attr.head()
```

Out[85]:

	campaign	pdays	previous	poutcome
0	2.570404	999.0	0.172596	nonexistent
1	1.000000	999.0	0.000000	nonexistent
2	1.000000	999.0	0.000000	nonexistent
3	1.000000	999.0	0.000000	nonexistent
4	1.000000	999.0	0.000000	nonexistent

## Treating the categorical attributes

In [86]:

```
remain_attr['poutcome'].unique()
```

Out[86]:

```
array(['nonexistent', 'failure', 'success'], dtype=object)
```

In [87]:

```
remain_attr['poutcome'] = remain_attr['poutcome'].map({'nonexistent':1, 'failure':2, 'success':3}).astype(int)
```

## Creating the 'final\_bank' dataset

We will now merge/concat all the above datasets that we created and curated as per need into one final dataset for our analysis

analysis.

In [88]:

```
final_bank= pd.concat([bank_client, other_attr, cont_attr, remain_attr], axis = 1)
```

## Exploring the dataset

In [89]:

```
final_bank.shape
```

Out[89]:

```
(41164, 38)
```

In [90]:

```
final_bank.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 41164 entries, 0 to 41187
Data columns (total 38 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   41164 non-null  int32
1   Job_N_admin.          41164 non-null  uint8
2   Job_N_blue-collar     41164 non-null  uint8
3   Job_N_entrepreneur    41164 non-null  uint8
4   Job_N_housemaid       41164 non-null  uint8
5   Job_N_management      41164 non-null  uint8
6   Job_N_retired         41164 non-null  uint8
7   Job_N_self-employed   41164 non-null  uint8
8   Job_N_services        41164 non-null  uint8
9   Job_N_student         41164 non-null  uint8
10  Job_N_technician      41164 non-null  uint8
11  Job_N_unemployed      41164 non-null  uint8
12  Job_N_unknown         41164 non-null  uint8
13  Marital_N             41164 non-null  int32
14  basic.4y              41164 non-null  uint8
15  basic.6y              41164 non-null  uint8
16  basic.9y              41164 non-null  uint8
17  high.school           41164 non-null  uint8
18  illiterate            41164 non-null  uint8
19  professional.course   41164 non-null  uint8
20  university.degree     41164 non-null  uint8
21  unknown               41164 non-null  uint8
22  Default_N             41164 non-null  int32
23  Housing_N             41164 non-null  int32
24  Loan_N                41164 non-null  int32
25  contact               41164 non-null  int32
26  month                 41164 non-null  int64
27  day_of_week           41164 non-null  int32
28  duration              41164 non-null  int32
29  emp.var.rate          41164 non-null  float64
30  cons.price.idx        41164 non-null  float64
31  cons.conf.idx         41164 non-null  float64
32  euribor3m             41164 non-null  float64
33  nr.employed           41164 non-null  float64
34  campaign              41164 non-null  float64
35  pdays                41164 non-null  float64
36  previous              41164 non-null  float64
37  poutcome              41164 non-null  int32
dtypes: float64(8), int32(9), int64(1), uint8(20)
memory usage: 5.3 MB
```

In [91]:

```
final_bank.isna().sum()
```

Out[91]:

```

age                                0
Job_N_admin.                      0
Job_N_blue-collar                 0
Job_N_entrepreneur               0
Job_N_housemaid                  0
Job_N_management                  0
Job_N_retired                    0
Job_N_self-employed              0
Job_N_services                   0
Job_N_student                    0
Job_N_technician                 0
Job_N_unemployed                 0
Job_N_unknown                    0
Marital_N                        0
basic.4y                         0
basic.6y                         0
basic.9y                         0
high.school                      0
illiterate                       0
professional.course              0
university.degree                0
unknown                          0
Default_N                        0
Housing_N                        0
Loan_N                           0
contact                          0
month                            0
day_of_week                      0
duration                         0
emp.var.rate                     0
cons.price.idx                   0
cons.conf.idx                    0
euribor3m                       0
nr.employed                      0
campaign                         0
pdays                            0
previous                         0
poutcome                         0
dtype: int64

```

In [92]:

```
final bank['campaign'].unique()
```

Out[92]:

```
array([ 2.57040373,  1.          ,  2.          ,  3.          ,  4.          ,
        5.          ,  6.          ,  7.          ,  8.          ,  9.          ,
        10.         , 11.         , 12.         , 13.         , 19.         ,
        18.         , 23.         , 14.         , 22.         , 25.         ,
        16.         , 17.         , 15.         , 20.         , 56.         ,
        42.         , 28.         , 26.         , 27.         , 32.         ,
        21.         , 24.         , 29.         , 31.         , 30.         ,
        35.         , 41.         , 37.         , 40.         , 33.         ,
        34.         , 43.         ])
```

In [93]:

```
final_bank['campaign'].fillna(final_bank['campaign'].mean(), inplace=True)
```

In [94]:

```
final bank.isna().sum()
```

Out[94]:

age	0
Job_N_admin.	0
Job_N_blue-collar	0
Job_N_entrepreneur	0
Job_N_housemaid	0
Job_N_management	0
Job_N_retired	0
Job_N_self-employed	0

```

Job_N_self-employed      0
Job_N_services           0
Job_N_student            0
Job_N_technician         0
Job_N_unemployed         0
Job_N_unknown            0
Marital_N                0
basic.4y                 0
basic.6y                 0
basic.9y                 0
high.school              0
illiterate               0
professional.course      0
university.degree        0
unknown                  0
Default_N                0
Housing_N                0
Loan_N                   0
contact                  0
month                    0
day_of_week              0
duration                 0
emp.var.rate             0
cons.price.idx           0
cons.conf.idx            0
euribor3m                0
nr.employed              0
campaign                 0
pdays                   0
previous                 0
poutcome                 0
dtype: int64

```

In [95]:

```
final_bank.describe()
```

Out[95]:

	age	Job_N_admin.	Job_N_blue-collar	Job_N_entrepreneur	Job_N_housemaid	Job_N_management	Job_N_retired	Job_N_self-employed
count	41164.000000	41164.000000	41164.000000	41164.000000	41164.000000	41164.000000	41164.000000	41164.000000
mean	1.978598	0.253037	0.224759	0.035371	0.025751	0.071033	0.041687	0.000000
std	0.735708	0.434757	0.417429	0.184717	0.158392	0.256883	0.199875	0.000000
min	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	2.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
75%	2.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
max	4.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

8 rows × 38 columns

## Splitting the data

We already have our target variable stored in 'y' from the beginning. Also, we have separately curated our final\_bank dataset. So, it does not contain our target variable y from our original dataset.

In [96]:

```

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(final_bank, y, test_size = 0.2, random_state = 0
)

```

In [97]:

```
from sklearn.model_selection import KFold
```



```

from sklearn.model_selection import cross_val_score
from sklearn.metrics import confusion_matrix, accuracy_score
k_fold = KFold(n_splits=10, shuffle=True, random_state=0)

```

In [98]:

```
X_train.head()
```

Out[98]:

	age	Job_N_admin.	Job_N_blue-collar	Job_N_entrepreneur	Job_N_housemaid	Job_N_management	Job_N_retired	Job_N_self-employed	Job
20018	2	0	0	0	0	0	0	0	
39695	1	1	0	0	0	0	0	0	
17238	3	0	0	0	1	0	0	0	
5924	3	0	1	0	0	0	0	0	
34656	2	1	0	0	0	0	0	0	

5 rows × 38 columns

## Scaling the data

In our `final_bank` data, we can see that the minimum and maximum value ranges from quite high to quite low values. For this reason, we are scaling our data with `StandardScaler`. We do so to scale our features centred around the zero and have unit variance.

In [99]:

```

#train-test split
from sklearn.preprocessing import StandardScaler
sc_X = StandardScaler()
X_train = sc_X.fit_transform(X_train)
X_test = sc_X.transform(X_test)

```

## Voting Classifier

In [100]:

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC

from sklearn.ensemble import VotingClassifier

```

### Hard voting

In [101]:

```

log_reg_clf = LogisticRegression(random_state= 0, C = 100, max_iter = 1000)
log_reg_clf.fit(X_train, y_train)

dtree_clf = DecisionTreeClassifier(max_depth = 1, random_state = 0)
dtree_clf.fit(X_train, y_train)

svc_clf = SVC(C = 0.1, gamma = 0.01, probability = True, random_state= 0)
svc_clf.fit(X_train, y_train)

hard_voting_clf = VotingClassifier(estimators=[('lr', log_reg_clf), ('dt', dtree_clf), ('svc',
svc_clf)], voting='hard')
hard_voting_clf.fit(X_train, y_train)

from sklearn.metrics import accuracy_score
for clf in (log_reg_clf, dtree_clf, svc_clf, hard_voting_clf):
    clf.fit(X_train, y_train)

```

```
y_pred = clf.predict(X_test)
print(clf.__class__.__name__, '%.4f'%accuracy_score(y_test, y_pred))
```

```
LogisticRegression 0.9093
DecisionTreeClassifier 0.8876
SVC 0.8987
VotingClassifier 0.9014
```

## Soft Voting

In [102]:

```
log_reg_clf = LogisticRegression(random_state= 0, C = 100, max_iter = 1000)
log_reg_clf.fit(X_train, y_train)

dtree_clf = DecisionTreeClassifier(max_depth = 1, random_state = 0)
dtree_clf.fit(X_train, y_train)

svc_clf = SVC(C = 0.1, gamma = 0.01, probability = True, random_state= 0)
svc_clf.fit(X_train, y_train)

soft_voting_clf = VotingClassifier(estimators=[('lr', log_reg_clf), ('dt', dtree_clf), ('svc',
svc_clf)], voting='soft')
soft_voting_clf.fit(X_train, y_train)

from sklearn.metrics import accuracy_score
for clf in (log_reg_clf, dtree_clf, svc_clf, soft_voting_clf):
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)
    print(clf.__class__.__name__, '%.4f'%accuracy_score(y_test, y_pred))
```

```
LogisticRegression 0.9093
DecisionTreeClassifier 0.8876
SVC 0.8987
VotingClassifier 0.9030
```

## Bagging

### Bagging for Decision Tree Classifier

In [104]:

```
from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier

dtree_clf = DecisionTreeClassifier(random_state=0)
dtree_bag_clf = BaggingClassifier(dtree_clf, n_estimators=500, max_samples=100, bootstrap=True, n_j
obs=-1, random_state=0)

dtree_bag_clf.fit(X_train, y_train)
y_pred = dtree_bag_clf.predict(X_test)
```

In [105]:

```
dtree_bag_clf.fit(X_train, y_train)

# train and test scores
print('Train score: %.2f'%dtree_bag_clf.score(X_train, y_train))
print('Test score: %.2f'%dtree_bag_clf.score(X_test, y_test))
```

```
Train score: 0.91
Test score: 0.91
```

In [106]:

```
print(confusion_matrix(y_test, dtree_bag_clf.predict(X_test) ))
```

```
from sklearn.metrics import classification_report
print(classification_report(y_train, dtree_bag_clf.predict(X_train)))
```

```
[[7128 180]
 [ 557 368]]
      precision    recall  f1-score   support

     0       0.93      0.98      0.95      29218
     1       0.67      0.39      0.49       3713

 accuracy          0.91      32931
 macro avg       0.80      0.68      0.72      32931
 weighted avg    0.90      0.91      0.90      32931
```

## Random Forest Classifier

### GridSearch

In [107]:

```
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier

param_grid = {'n_estimators': [200, 300, 400, 500],
              'max_depth': np.arange(1, 10)}

rf_gridsearch = GridSearchCV(RandomForestClassifier(random_state=0), param_grid, cv=10,
return_train_score=True)
rf_gridsearch.fit(X_train, y_train)
print("Best parameters for RandomForest Clf: {}".format(rf_gridsearch.best_params_))
print("Best cross-validation score: {:.2f}".format(rf_gridsearch.best_score_))
```

Best parameters for RandomForest Clf: {'max\_depth': 9, 'n\_estimators': 500}  
Best cross-validation score: 0.91

## Random Forest Classifier

In [108]:

```
rf_clf = RandomForestClassifier(n_estimators=400, max_depth = 9, bootstrap=True, n_jobs=-1, random_
state=0)
rf_clf.fit(X_train, y_train)

pred_rf = rf_clf.predict(X_test)

#train and test scores
print('Train score: {:.2f}'.format(rf_clf.score(X_train, y_train)))
print('Test score: {:.2f}'.format(rf_clf.score(X_test, y_test)))
```

Train score: 0.92  
Test score: 0.91

In [109]:

```
print(confusion_matrix(y_test, rf_clf.predict(X_test)))
from sklearn.metrics import classification_report
print(classification_report(y_train, rf_clf.predict(X_train)))
```

```
[[7189 119]
 [ 641 284]]
      precision    recall  f1-score   support

     0       0.93      0.99      0.96      29218
     1       0.86      0.39      0.54       3713

 accuracy          0.92      32931
 macro avg       0.89      0.69      0.75      32931
```

macro avg	0.89	0.89	0.70	32931
weighted avg	0.92	0.92	0.91	32931

## Pasting

### Decision Tree Classifier

In [110]:

```
from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier

# pasting: bootstrap = False

dtree_clf = DecisionTreeClassifier(criterion = 'entropy', random_state=0)
dtree_bag_clf = BaggingClassifier(dtree_clf, n_estimators=500, max_samples=100, bootstrap=False, random_state=0)

dtree_bag_clf.fit(X_train, y_train)
y_pred = dtree_bag_clf.predict(X_test)

from sklearn.metrics import accuracy_score

# train and test scores
print('Train score: %.2f'%dtree_bag_clf.score(X_train, y_train))
print('Test score: %.2f'%dtree_bag_clf.score(X_test, y_test))
```

Train score: 0.91  
Test score: 0.91

In [111]:

```
print(confusion_matrix(y_test, dtree_bag_clf.predict(X_test)))
from sklearn.metrics import classification_report
print(classification_report(y_train, dtree_bag_clf.predict(X_train)))
```

```
[[7172  136]
 [ 621  304]]
```

	precision	recall	f1-score	support
0	0.92	0.98	0.95	29218
1	0.70	0.32	0.44	3713
accuracy			0.91	32931
macro avg	0.81	0.65	0.69	32931
weighted avg	0.89	0.91	0.89	32931

### SVC Classifier

In [112]:

```
from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier

# pasting: bootstrap = False

svc_clf = SVC(C = 0.1, gamma = 0.01, probability = True, random_state= 0)
svc_bag_clf = BaggingClassifier(svc_clf, n_estimators=500, max_samples=100, bootstrap=False, random_state=0)

svc_bag_clf.fit(X_train, y_train)
y_pred = svc_bag_clf.predict(X_test)

from sklearn.metrics import accuracy_score

#train and test scores
print('Train score: %.2f'%svc_bag_clf.score(X_train, y_train))
print('Test score: %.2f'%svc_bag_clf.score(X_test, y_test))
```

```
print('Train score: %.2f'%svc_bag_clf.score(X_train, y_train))
print('Test score: %.2f'%svc_bag_clf.score(X_test, y_test))
```

Train score: 0.89  
Test score: 0.89

In [113]:

```
print(confusion_matrix(y_test, svc_bag_clf.predict(X_test)))
from sklearn.metrics import classification_report
print(classification_report(y_train, svc_bag_clf.predict(X_train)))
```

```
[[7289  19]
 [ 875  50]]

              precision    recall  f1-score   support

         0       0.89        1.00        0.94        29218
         1       0.86        0.06        0.11         3713

 accuracy          0.89          0.89          0.89          32931
 macro avg          0.87          0.53          0.53          32931
 weighted avg          0.89          0.89          0.85          32931
```

## ADA Boost Classifier

### Decision Tree

In [114]:

```
from sklearn.ensemble import AdaBoostClassifier

dtree_ada_clf = AdaBoostClassifier(DecisionTreeClassifier(max_depth=1), n_estimators=200,
                                   algorithm="SAMME.R", learning_rate=0.5, random_state=0)

dtree_ada_clf.fit(X_train, y_train)
predictions = dtree_ada_clf.predict(X_test)

#train and test scores
print('Train score: %.2f'%dtree_ada_clf.score(X_train, y_train))
print('Test score: %.2f'%dtree_ada_clf.score(X_test, y_test))
```

Train score: 0.91  
Test score: 0.91

In [115]:

```
confusion_matrix(y_test, predictions)
```

Out[115]:

```
array([[7108,  200],
       [ 560,  365]], dtype=int64)
```

In [116]:

```
print(confusion_matrix(y_test, dtree_ada_clf.predict(X_test)))
from sklearn.metrics import classification_report
print(classification_report(y_train, dtree_ada_clf.predict(X_train)))
```

```
[[7108  200]
 [ 560  365]]

              precision    recall  f1-score   support

         0       0.93        0.97        0.95        29218
         1       0.66        0.40        0.50         3713
```

accuracy			0.91	32931
macro avg	0.79	0.69	0.72	32931
weighted avg	0.90	0.91	0.90	32931

## Logistic Regression

In [117]:

```
from sklearn.ensemble import AdaBoostClassifier

log_reg_ada_clf = AdaBoostClassifier(LogisticRegression(solver='liblinear'), n_estimators=500,
                                     algorithm="SAMME.R", learning_rate=0.1, random_state=0)

log_reg_ada_clf.fit(X_train, y_train)
predictions = log_reg_ada_clf.predict(X_test)

#train and test scores
print('Train score: %.2f'%log_reg_ada_clf.score(X_train, y_train))
print('Test score: %.2f'%log_reg_ada_clf.score(X_test, y_test))
```

Train score: 0.91  
Test score: 0.91

In [118]:

```
print(confusion_matrix(y_test, log_reg_ada_clf.predict(X_test)))
from sklearn.metrics import classification_report
print(classification_report(y_train, log_reg_ada_clf.predict(X_train)))
```

```
[[7194  114]
 [ 663 262]]
```

	precision	recall	f1-score	support
0	0.92	0.98	0.95	29218
1	0.70	0.28	0.40	3713

accuracy			0.91	32931
macro avg	0.81	0.63	0.68	32931
weighted avg	0.89	0.91	0.89	32931

## Gradient Boosting Classifier

In [119]:

```
from sklearn.ensemble import GradientBoostingClassifier

gbrt = GradientBoostingClassifier(random_state=0, max_depth=5, learning_rate=0.01)
gbrt.fit(X_train, y_train)

#train and test scores
print("Accuracy on training set: {:.3f}".format(gbrt.score(X_train, y_train)))
print("Accuracy on test set: {:.3f}".format(gbrt.score(X_test, y_test)))
```

Accuracy on training set: 0.904  
Accuracy on test set: 0.901

In [120]:

```
print(confusion_matrix(y_test, gbrt.predict(X_test)))
from sklearn.metrics import classification_report
print(classification_report(y_train, gbrt.predict(X_train)))
```

```
[[7276  32]
 [ 779 146]]
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

	0	0.90	1.00	0.95	29218
	1	0.87	0.17	0.28	3713
accuracy				0.90	32931
macro avg		0.89	0.58	0.62	32931
weighted avg		0.90	0.90	0.87	32931

## PCA

In [121]:

```
from sklearn.decomposition import PCA

pca = PCA()
pca.fit(X_train)

cumsum = np.cumsum(pca.explained_variance_ratio_)
d = np.argmax(cumsum >= 0.95) + 1
```

In [122]:

```
d
```

Out[122]:

```
29
```

In [123]:

```
pca = PCA(n_components=0.95)
X_reduced = pca.fit_transform(X_train)
```

In [124]:

```
pca.n_components_
```

Out[124]:

```
29
```

In [125]:

```
np.sum(pca.explained_variance_ratio_)
```

Out[125]:

```
0.9538315288981813
```

In [126]:

```
pca = PCA(n_components = 29)
X_reduced = pca.fit_transform(X_train)
X_recovered = pca.inverse_transform(X_reduced)
```

In [127]:

```
X_reduced_pca = X_reduced
```

In [128]:

```
from sklearn.decomposition import IncrementalPCA

n_batches = 100

inc_pca = IncrementalPCA(n_components=29)
```

```

inc_pca = IncrementalPCA(n_components=29)
for X_batch in np.array_split(X_train, n_batches):
    print(".", end="")
    inc_pca.partial_fit(X_batch)

X_train_reduced = inc_pca.transform(X_train)

```

In [129]:

```

from sklearn.decomposition import IncrementalPCA

n_batches = 100

inc_pca = IncrementalPCA(n_components=29)
for X_batch in np.array_split(X_test, n_batches):
    print(".", end="")
    inc_pca.partial_fit(X_batch)

X_test_reduced = inc_pca.transform(X_test)

```

In [130]:

```
X_train_reduced.shape
```

Out[130]:

```
(32931, 29)
```

In [131]:

```
X_test_reduced.shape
```

Out[131]:

```
(8233, 29)
```

## Models on PCA data

### Logistic Regression

In [132]:

```

from sklearn.linear_model import LogisticRegression
log_model = LogisticRegression()
log_model.fit(X_train_reduced, y_train)
log_pred = log_model.predict(X_test_reduced)

```

In [133]:

```

print(confusion_matrix(y_test, log_pred))
from sklearn.metrics import classification_report
print(classification_report(y_train, log_model.predict(X_train_reduced)))

```

```

[[7086  222]
 [ 711 214]]

```

	precision	recall	f1-score	support
0	0.92	0.98	0.95	29218
1	0.65	0.36	0.46	3713
accuracy			0.91	32931
macro avg	0.79	0.67	0.70	32931
weighted avg	0.89	0.91	0.89	32931



In [134]:

```
lr_score_train = log_model.score(X_train_reduced,y_train)
lr_score_train
```

Out[134]:

0.9056208435820352

In [135]:

```
lr_score_test = log_model.score(X_test_reduced,y_test)
lr_score_test
```

Out[135]:

0.886675573909875

## KNN Classifier

In [136]:

```
from sklearn import model_selection
from sklearn.neighbors import KNeighborsClassifier
```

```
neighbors = np.arange(0,25)
```

```
cv_scores = []
```

In [137]:

```
# To determine best k-value
for k in neighbors:
    k_val = k+1
    knn_clf = KNeighborsClassifier(n_neighbors = k_val, weights='uniform', p=2, metric='euclidean')
    k_fold = model_selection.KFold(n_splits=10, random_state=123)
    cross_val_scores = model_selection.cross_val_score(knn_clf, X_train_reduced, y_train, cv=k_fold,
    , scoring='accuracy')
    cv_scores.append(cross_val_scores.mean()*100)
    print("k=%d %0.2f (+/- %0.2f)" % (k_val, cross_val_scores.mean()*100, cross_val_scores.std()*100))
```

```
optimal_kval = neighbors[cv_scores.index(max(cv_scores))]
print ("The optimal number of neighbors is %d with %0.1f%%" % (optimal_kval,
cv_scores[optimal_kval]))
```


```
plt.plot(neighbors, cv_scores)
plt.xlabel('Number of Neighbors K')
plt.ylabel('Train Accuracy')
plt.show()
```

```
k=1 87.67 (+/- 0.61)
k=2 89.31 (+/- 0.62)
k=3 89.17 (+/- 0.43)
k=4 89.55 (+/- 0.53)
k=5 89.57 (+/- 0.44)
k=6 89.71 (+/- 0.47)
k=7 89.66 (+/- 0.35)
k=8 89.78 (+/- 0.45)
k=9 89.79 (+/- 0.43)
k=10 89.91 (+/- 0.37)
k=11 89.87 (+/- 0.35)
k=12 89.90 (+/- 0.45)
k=13 89.92 (+/- 0.45)
k=14 89.93 (+/- 0.46)
k=15 89.96 (+/- 0.40)
k=16 90.00 (+/- 0.46)
k=17 90.03 (+/- 0.39)
```

```

k=17 90.00 (+/- 0.39)
k=18 90.01 (+/- 0.46)
k=19 90.04 (+/- 0.46)
k=20 90.01 (+/- 0.53)
k=21 90.08 (+/- 0.48)
k=22 90.00 (+/- 0.50)
k=23 90.14 (+/- 0.46)
k=24 90.06 (+/- 0.44)
k=25 90.10 (+/- 0.47)
The optimal number of neighbors is 22 with 90.1%

```



In [138]:

```

from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=24)
knn.fit(X_train_reduced, y_train)
knn_pred = knn.predict(X_test_reduced)

```

In [139]:

```

# for train data
from sklearn.metrics import classification_report
print(classification_report(y_train, knn.predict(X_train_reduced)))

```

	precision	recall	f1-score	support
0	0.91	0.99	0.95	29218
1	0.71	0.25	0.37	3713
accuracy			0.90	32931
macro avg	0.81	0.62	0.66	32931
weighted avg	0.89	0.90	0.88	32931

In [140]:

```

# for test data
confusion_matrix(y_test, knn.predict(X_test_reduced))

```

Out[140]:

```

array([[7249, 59],
       [ 875, 50]], dtype=int64)

```

In [141]:

```

print(classification_report(y_test, knn.predict(X_test_reduced)))

```

	precision	recall	f1-score	support
0	0.89	0.99	0.94	7308
1	0.46	0.05	0.10	925
accuracy			0.89	8233
macro avg	0.68	0.52	0.52	8233

weighted avg      0.84      0.89      0.84      8233

In [142]:

```
knn_score_test = knn.score(X_test_reduced,y_test)
knn_score_test
```

Out[142]:

0.88655411150249

## Linear SVM

In [143]:

```
from sklearn.svm import LinearSVC

svm = LinearSVC()
svm.fit(X_train_reduced,y_train)
svc_pred = svm.predict(X_test_reduced)
```

In [144]:

```
# for train set
from sklearn.metrics import classification_report
print(classification_report(y_train, svm.predict(X_train_reduced)))
```

	precision	recall	f1-score	support
0	0.92	0.98	0.95	29218
1	0.67	0.31	0.42	3713
accuracy			0.90	32931
macro avg	0.79	0.64	0.68	32931
weighted avg	0.89	0.90	0.89	32931

In [145]:

```
#for test set
from sklearn.metrics import classification_report
print(classification_report(y_test, svm.predict(X_test_reduced)))
```

	precision	recall	f1-score	support
0	0.90	0.98	0.94	7308
1	0.48	0.17	0.25	925
accuracy			0.89	8233
macro avg	0.69	0.57	0.59	8233
weighted avg	0.86	0.89	0.86	8233

In [146]:

```
svm_score_train = svm.score(X_train_reduced,y_train)
svm_score_train
```

Out[146]:

0.904588381767939

In [147]:

```
svm_score_test = svm.score(X_test_reduced,y_test)
svm_score_test
```

```
Out[147]:  
0.8860682618729503
```

## Decision Tree Classifier

```
In [148]:
```

```
d_tree = DecisionTreeClassifier(criterion='entropy', random_state=0)  
d_tree.fit(X_train_reduced, y_train)  
d_tree_pred = d_tree.predict(X_test_reduced)
```

```
In [149]:
```

```
print(classification_report(y_test, d_tree.predict(X_test_reduced)))
```

	precision	recall	f1-score	support
0	0.91	0.88	0.89	7308
1	0.23	0.28	0.26	925
accuracy			0.81	8233
macro avg	0.57	0.58	0.57	8233
weighted avg	0.83	0.81	0.82	8233

```
In [150]:
```

```
print(classification_report(y_train, d_tree.predict(X_train_reduced)))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	29218
1	1.00	1.00	1.00	3713
accuracy			1.00	32931
macro avg	1.00	1.00	1.00	32931
weighted avg	1.00	1.00	1.00	32931

```
In [151]:
```

```
d_tree_score_train = d_tree.score(X_train_reduced,y_train)  
d_tree_score_train
```

```
Out[151]:  
0.9999696334760559
```

```
In [152]:
```

```
d_tree_score_test = d_tree.score(X_test_reduced,y_test)  
d_tree_score_test
```

```
Out[152]:  
0.814283979108466
```

## Kernalized SVM(linear, rbf, poly)

```
In [153]:
```

```
from sklearn.svm import SVC  
from sklearn.model_selection import GridSearchCV  
  
tuned_parameters = [{'kernel': ['rbf'], 'gamma': [0.1],
```

```

'C': [1]],
{'kernel': ['linear'], 'C': [1]},
{'kernel': ['poly'], 'degree': [3], 'C': [10]}

```

In [154]:

```

clf = GridSearchCV(SVC(), tuned_parameters, cv=5, scoring='precision')
clf.fit(X_train, y_train)

print(clf.cv_results_)

{'mean_fit_time': array([91.10583048, 70.27198591, 80.00121274]), 'std_fit_time':
array([1.04248985, 7.38563372, 1.80327034]), 'mean_score_time': array([5.95953856, 2.66312752,
2.90309567]), 'std_score_time': array([0.01897421, 0.03878049, 0.03692636]), 'param_C':
masked_array(data=[1, 1, 10],
              mask=[False, False, False],
              fill_value='?',
              dtype=object), 'param_gamma': masked_array(data=[0.1, --, --],
              mask=[False, True, True],
              fill_value='?',
              dtype=object), 'param_kernel': masked_array(data=['rbf', 'linear', 'poly'],
              mask=[False, False, False],
              fill_value='?',
              dtype=object), 'param_degree': masked_array(data=[--, --, 3],
              mask=[ True, True, False],
              fill_value='?',
              dtype=object), 'params': [{'C': 1, 'gamma': 0.1, 'kernel': 'rbf'}, {'C': 1, 'kernel': '
linear'}, {'C': 10, 'degree': 3, 'kernel': 'poly'}], 'split0_test_score': array([0.67412141,
0.63099631, 0.60408163]), 'split1_test_score': array([0.63467492, 0.62939297, 0.53682171]),
'split2_test_score': array([0.62170088, 0.65116279, 0.56097561]), 'split3_test_score':
array([0.63276836, 0.64312268, 0.59578544]), 'split4_test_score': array([0.59744409, 0.62666667, 0
.5462963 ]), 'mean_test_score': array([0.63214193, 0.63626828, 0.56879214]), 'std_test_score':
array([0.02482757, 0.00933899, 0.02669553]), 'rank_test_score': array([2, 1, 3])}

```

In [155]:

```

print('The best model is: ', clf.best_params_)
print('This model produces a mean cross-validated score (precision) of', clf.best_score_)

```

The best model is: {'C': 1, 'kernel': 'linear'}  
This model produces a mean cross-validated score (precision) of 0.6362682830306745

In [156]:

```

svm_ker_lin = SVC(kernel='linear', C=1)
svm_ker_rbf = SVC(kernel='rbf', gamma=0.1, C=1)
svm_ker_poly = SVC(kernel='poly', degree=3, C=10)

```

In [157]:

```

svm_ker_lin.fit(X_train_reduced, y_train)
svm_ker_rbf.fit(X_train_reduced, y_train)
svm_ker_poly.fit(X_train_reduced, y_train)

ker_lin_pred = svm.predict(X_test_reduced)
ker_rbf_pred = svm.predict(X_test_reduced)
ker_poly_pred = svm.predict(X_test_reduced)

```

In [158]:

```

print(classification_report(y_train, svm_ker_lin.predict(X_train_reduced)))

```

	precision	recall	f1-score	support
0	0.92	0.98	0.95	29218
1	0.65	0.29	0.41	3713
accuracy			0.90	32931
macro avg	0.79	0.64	0.68	32931
weighted avg	0.89	0.90	0.89	32931

In [159]:

```
print(classification_report(y_train, svm_ker_rbf.predict(X_train_reduced)))
```

	precision	recall	f1-score	support
0	0.94	0.99	0.97	29218
1	0.89	0.52	0.66	3713
accuracy			0.94	32931
macro avg	0.92	0.76	0.81	32931
weighted avg	0.94	0.94	0.93	32931

In [160]:

```
print(classification_report(y_train, svm_ker_poly.predict(X_train_reduced)))
```

	precision	recall	f1-score	support
0	0.94	0.98	0.96	29218
1	0.80	0.50	0.62	3713
accuracy			0.93	32931
macro avg	0.87	0.74	0.79	32931
weighted avg	0.92	0.93	0.92	32931

In [161]:

```
print(classification_report(y_test, svm_ker_poly.predict(X_test_reduced)))
```

	precision	recall	f1-score	support
0	0.91	0.91	0.91	7308
1	0.30	0.32	0.31	925
accuracy			0.84	8233
macro avg	0.61	0.61	0.61	8233
weighted avg	0.84	0.84	0.84	8233

In [162]:

```
print(classification_report(y_test, svm_ker_lin.predict(X_test_reduced)))
```

	precision	recall	f1-score	support
0	0.90	0.98	0.94	7308
1	0.53	0.15	0.24	925
accuracy			0.89	8233
macro avg	0.72	0.57	0.59	8233
weighted avg	0.86	0.89	0.86	8233

In [163]:

```
print(classification_report(y_test, svm_ker_rbf.predict(X_test_reduced)))
```

	precision	recall	f1-score	support
0	0.89	0.99	0.94	7308
1	0.31	0.02	0.04	925
accuracy			0.88	8233
macro avg	0.60	0.51	0.49	8233

weighted avg	0.82	0.88	0.84	8233
--------------	------	------	------	------

In [164]:

```
# Since linear kernel is our best model, we will consider it's train and test scores.
#test score
svm_ker_lin_score_test = svm_ker_lin.score(X_test_reduced,y_test)
svm_ker_lin_score_test
```

Out[164]:

0.8895906716871128

In [165]:

```
#train score
svm_ker_lin_score_train = svm_ker_lin.score(X_train_reduced,y_train)
svm_ker_lin_score_train
```

Out[165]:

0.9028574899031307

**Results from Project 1 for all models :**

***Train scores for our models are as follows :***

***Logistic Regression : 0.9092951929792596***

***KNN Classification : 0.9031915216665148***

***Linear SVM : 0.9085056633567156***

***Kernalized SVM : 0.8997904709847864***

***Decision Tree Classifier : 0.9999696334760559***

***Test scores for our models are as follows :***

***Logistic Regression : 0.9086602696465444***

***KNN Classification : 0.9001579011296004***

***Linear SVM : 0.9069597959431556***

***Kernalized SVM : 0.9086602696465444***

***Decision Tree Classifier : 0.8933560063160452***

***As we can see from our train and test scores from Project 1 and our train and test scores when using our PCA reduced dataset, the scores have dropped. Hence, we can say that the dimensionally reduced dataset results in a poorer score compared to the original dataset. Although, PCA did help in saving computational time by reducing the features to 29 from 38.***

In [166]:

```
import numpy
from keras.models import Sequential
from keras.layers import Dense
from keras.wrappers.scikit_learn import KerasClassifier

# create model
clf_model = Sequential()
clf_model.add(Dense(12, input_dim=45, activation='relu'))
clf_model.add(Dense(8, activation='relu'))
```

```
clf_model.add(Dense(1, activation='sigmoid'))

# Compile model
clf_model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
return model

# Fit the model
clf_model.fit(X_train, y_train, epochs=150, batch_size=10)

# evaluate the model
clf_model_scores = clf_model.evaluate(X_test, y_test)
print("\ns: %.2f%%" % (clf_model.metrics_names[1], scores[1]*100))
```

```
-----
ModuleNotFoundError                                Traceback (most recent call last)
<ipython-input-166-2a2fa6886b37> in <module>
      1 import numpy
----> 2 from keras.models import Sequential
      3 from keras.layers import Dense
      4 from keras.wrappers.scikit_learn import KerasClassifier
      5
```

**ModuleNotFoundError:** No module named 'keras'

In [ ]:

```
y_predict = clf_model.predict(X_test)
y_predict
```

In [ ]: